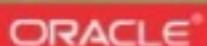




Microservices + Oracle | A Bright Future

Kelly Goetsch

Director, Product Management



January 28th 2016



Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.



Program Agenda

- 1 ➤ Introduction
- 2 ➤ History of Microservices
- 3 ➤ Non-technical Prerequisites
- 4 ➤ Architectural Prerequisites
- 5 ➤ Technical Prerequisites
- 6 ➤ Implementing Microservices
- 7 ➤ How Oracle Products Support Microservices
- 8 ➤ Summary



An aerial photograph of a massive aircraft fuselage, likely a Boeing 787, being assembled in a large industrial hangar. The fuselage is white with blue protective covers on the upper sections. Numerous workers in blue uniforms are scattered throughout the hangar, some standing on elevated walkways and others working directly on the aircraft. The floor is a light-colored concrete with yellow safety markings. In the foreground, the large, teal-colored vertical stabilizer is being lowered into place, its base visible at the bottom center of the frame.

Introduction

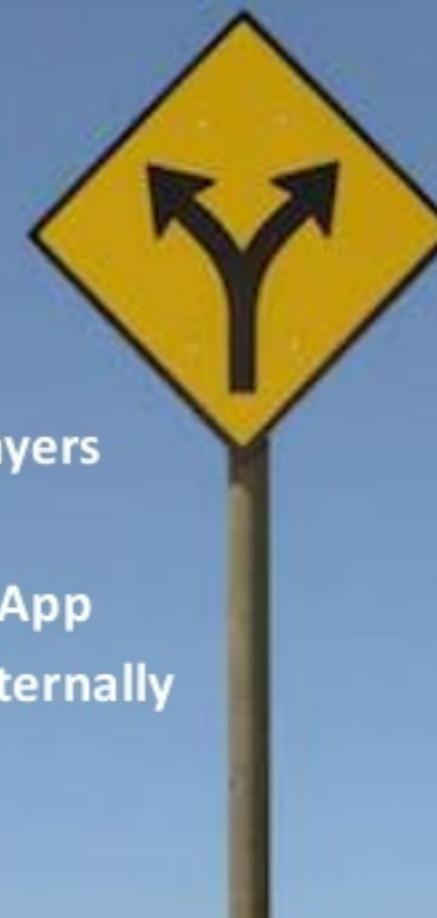


What Are Microservices?

Minimal function services that are deployed separately but can interact together to achieve a broader use-case

Monolithic Applications

- Single, Monolithic App
- Must Deploy Entire App
- One Database for Entire App
- Organized Around Technology Layers
- State In Each Runtime Instance
- One Technology Stack for Entire App
- In-process Calls Locally, SOAP Externally



Microservices

- Many, Smaller Minimal Function Microservices
- Can Deploy Each Microservice Independently
- Each Microservice Often Has Its Own Datastore
- Organized Around Business Capabilities
- State is Externalized
- Choice of Technology for Each Microservice
- REST Calls Over HTTP, Messaging, or Binary

Microservices Are Analogous to Unix Utilities

Same concept, different decade

"Write programs that do one thing and do it well. Write programs to work together. Write programs to handle text streams, because that is a universal interface."



Doug McIlroy

Inventor of the Unix Pipe

- Unix Executable: Does one thing and does it well
- Runs independent of other commands
- Produces text-based response

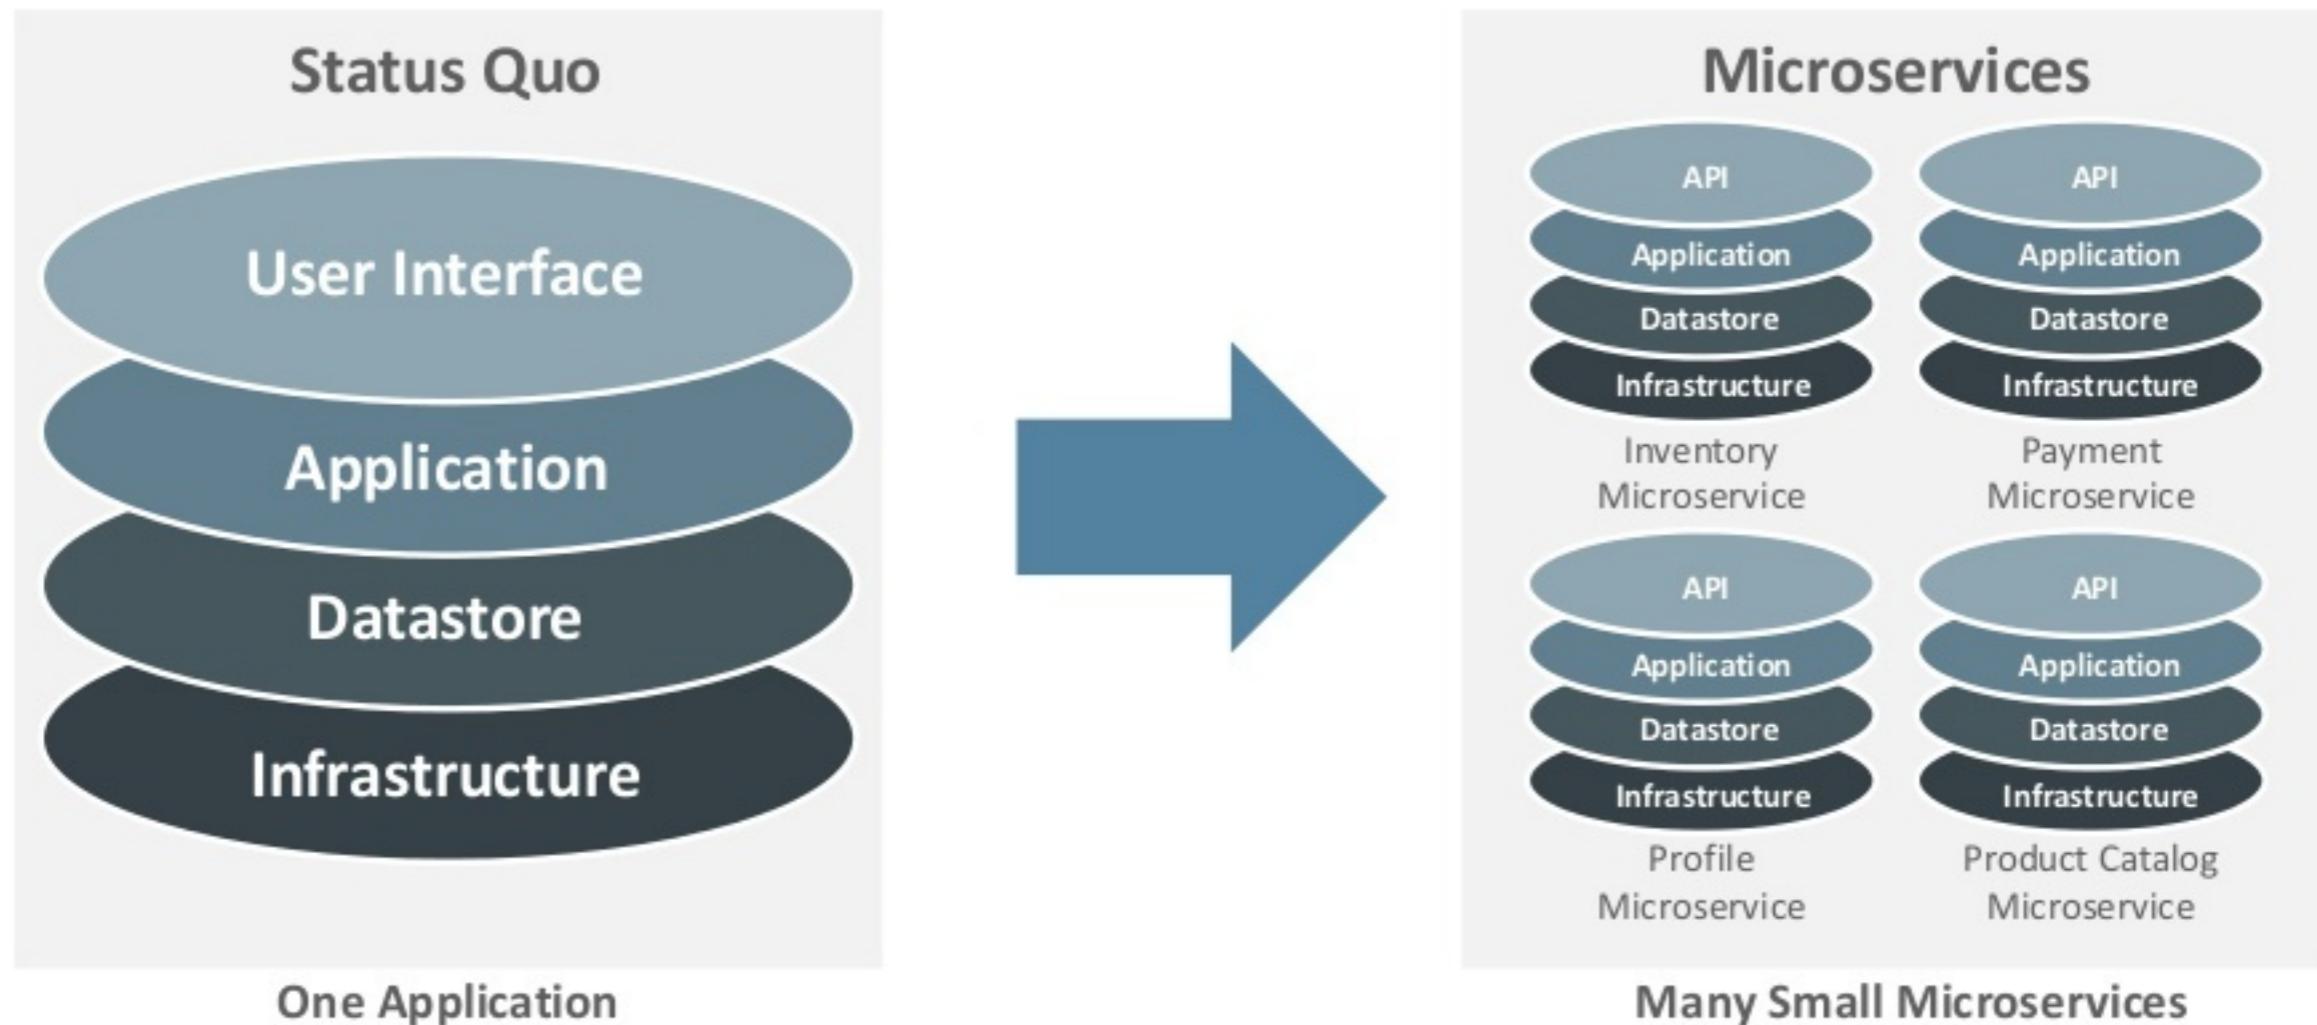
```
curl -v -H "Accept: application/json" -H "Content-type: application/json" -X POST  
-d '{"productId":645887,"quantity":"1"}'  
"http://localhost:8840/rest/ShoppingCart/"
```

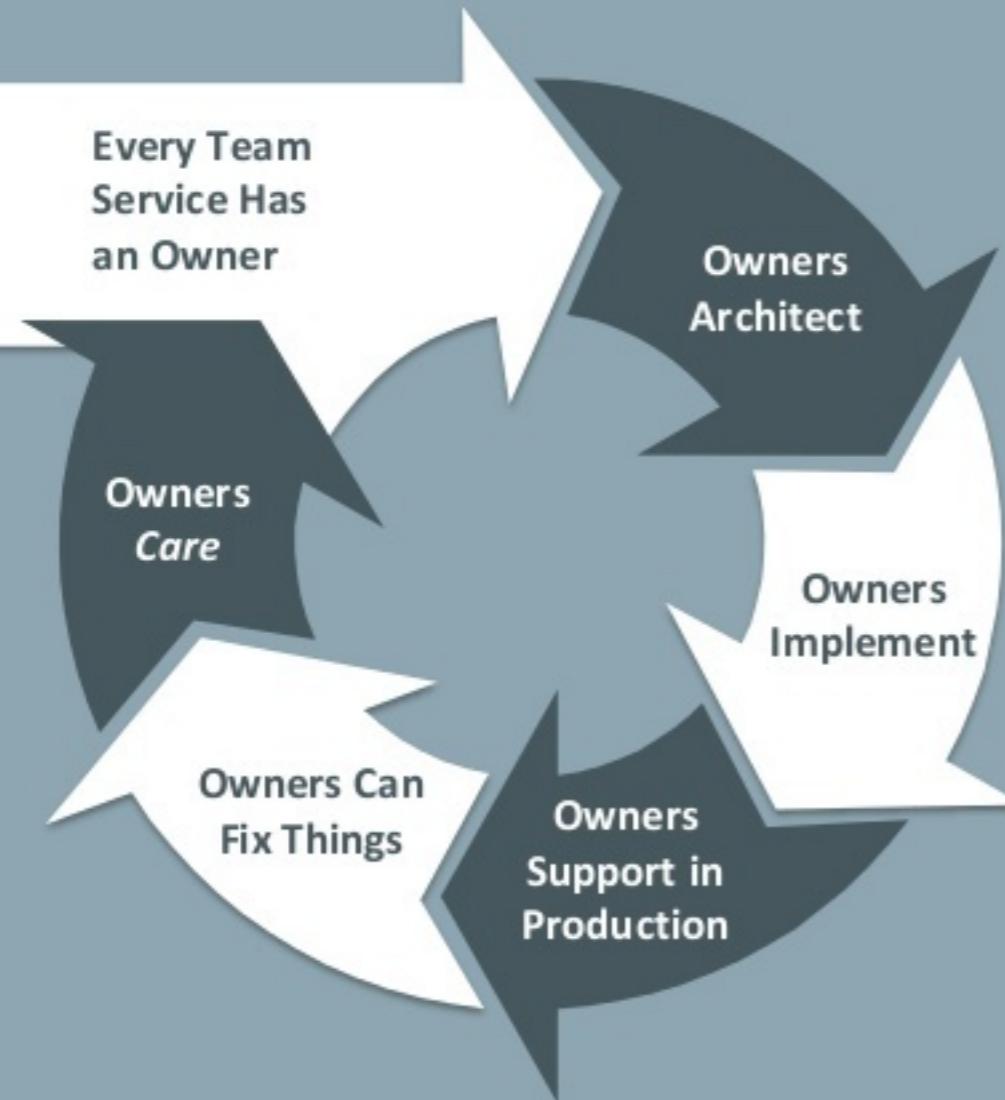


- Microservice: Does one thing and does it well
- Runs independent of other microservices
- Produces text-based response to clients



Microservices Apps Are Developed/Deployed Independently





Ownership is Key to the Success of Microservices

In traditional enterprises, anyone individual has very low ownership of anything. It's classic tragedy of the commons

Fundamentally, Microservices is a Tradeoff

Do you want...

Traditional App Development

Easier Deployment/Ops

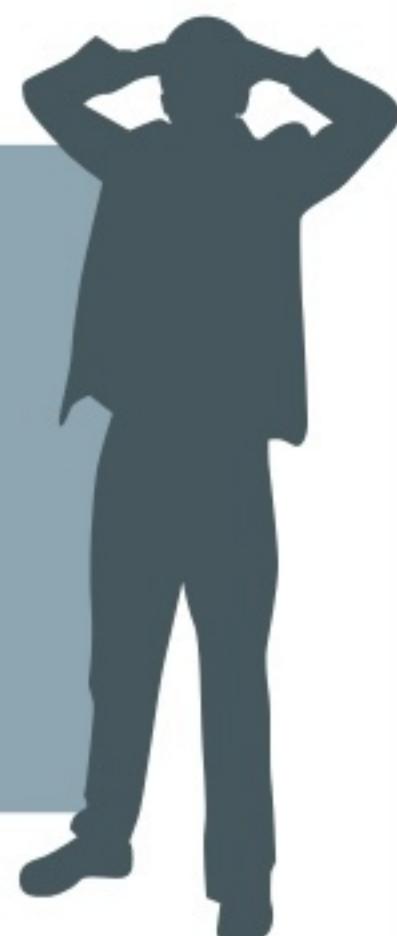
- One big block of code, sometimes broken into semi-porous modules
- Complexity handled inside the big block of code
- Each big block is hard to develop but easy to deploy



Microservices

Easier Development

- Many small blocks of code, each developed and deployed independently
- Complexity encapsulated in each microservice
- Each microservice is easy to develop but hard to deploy





Top 5 Signs It's Time To Look at Microservices

1. 100+ developers for an app
2. 5m lines of code for an app
3. Monthly or quarterly releases to production
4. > 1 quarter backlog of development work
5. > 20% developer turnover

Common Microservice Adoption Use Cases



I want to **extend** my existing monolithic application by adding microservices on the periphery.

I want to **decompose** an existing modular application into a microservices-style application

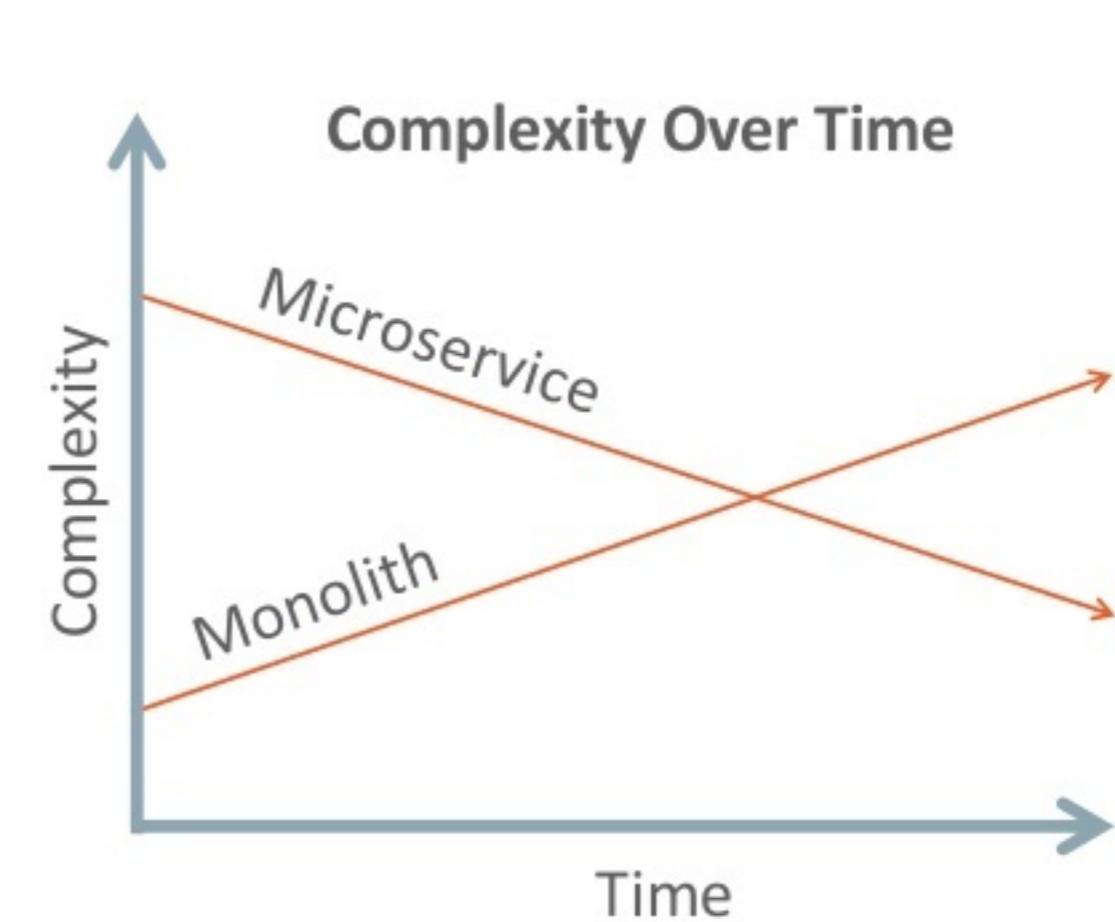
I want to build a **net new** microservices-style application from the ground up.



Sometimes Monolithic Applications Are Still a Good Fit

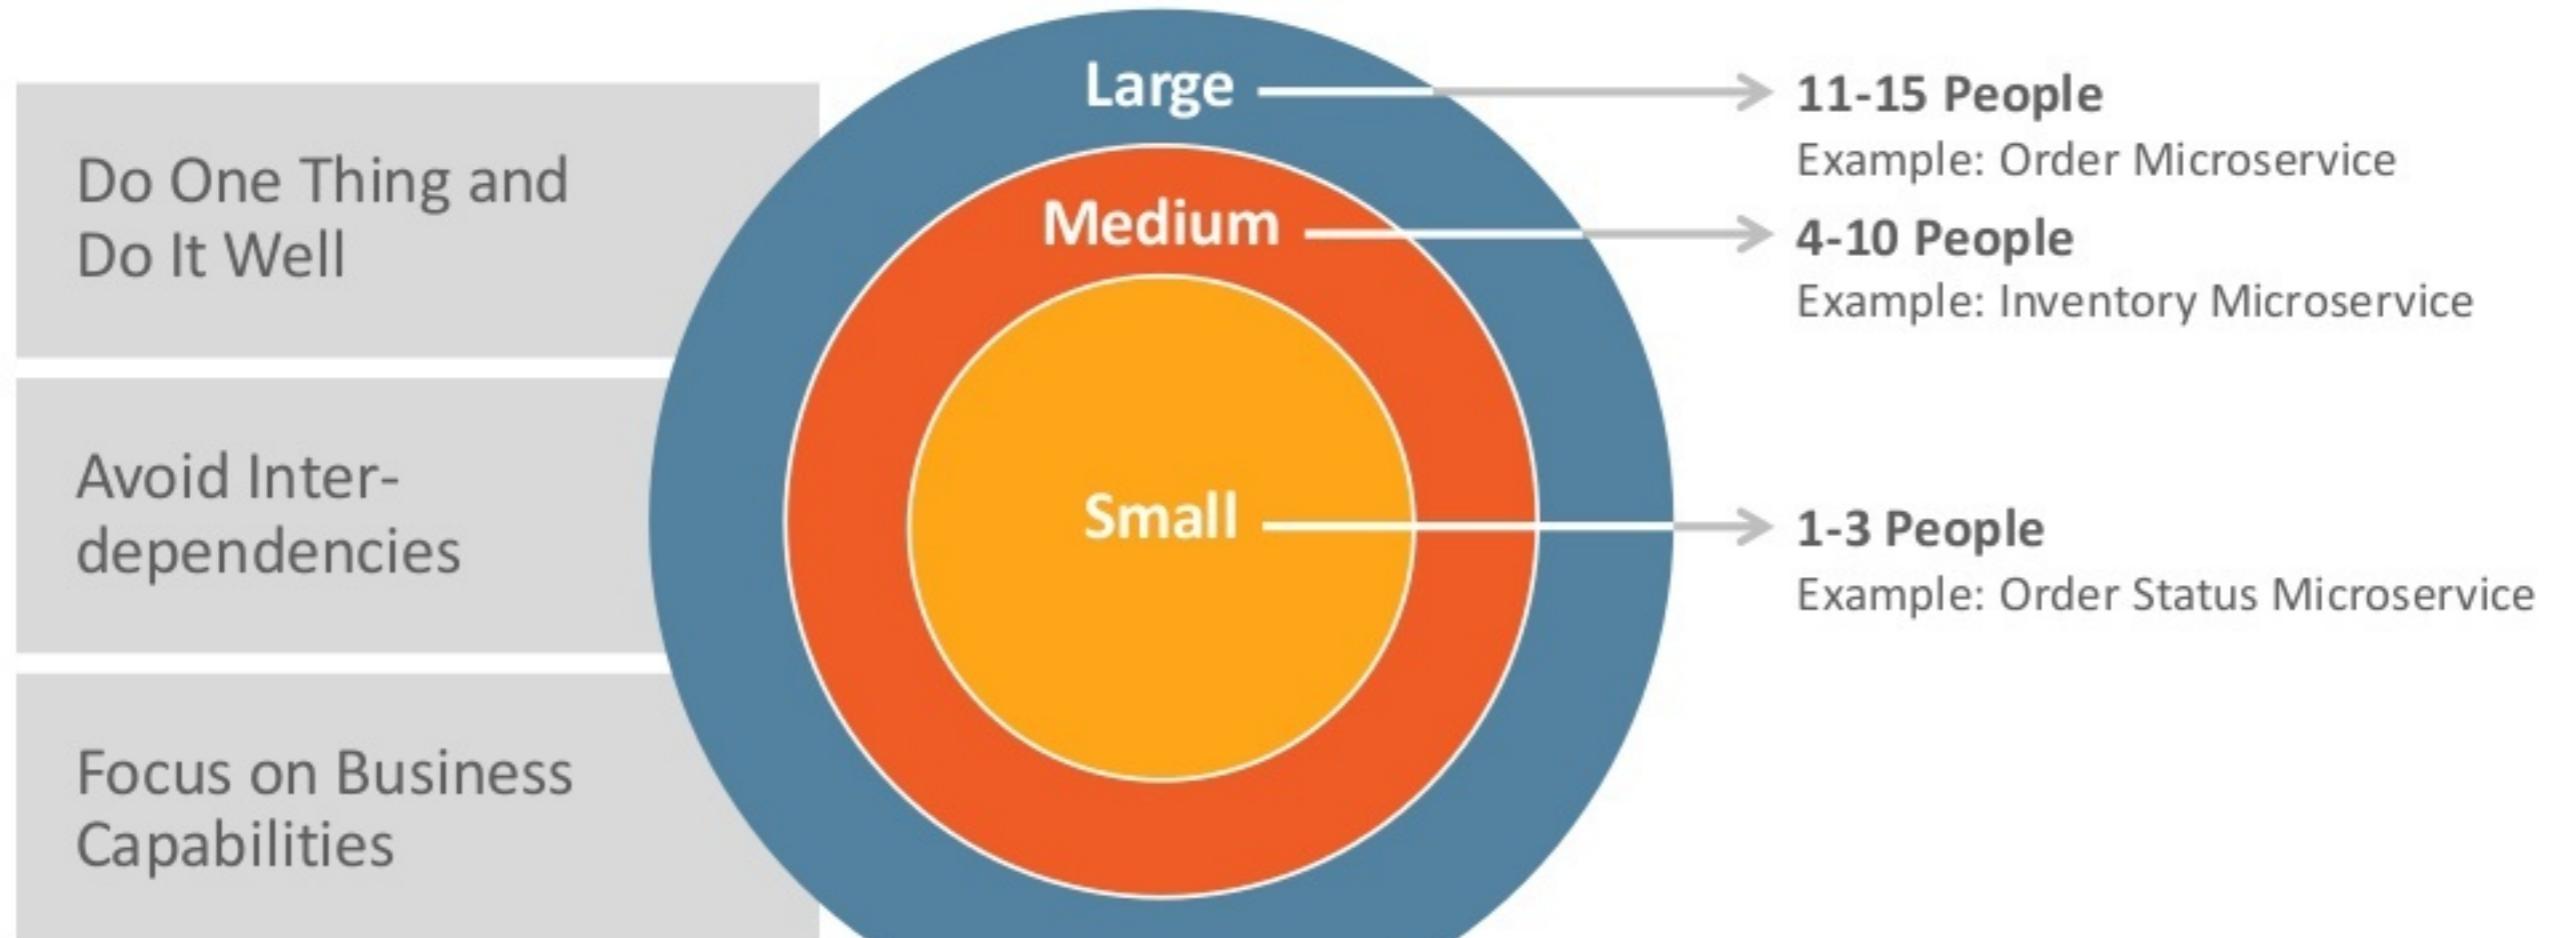
Microservices add complexity

- For less complex applications, monoliths are always better in both the long and short-run
- For moderately complex applications, monoliths are still probably better in both the long and short-run
- For complex apps, microservices may pay off over time but it takes a long time to offset the high up-front investment required to do it



How Big Should a Microservice Be?

Can have hundreds of microservices for a larger application

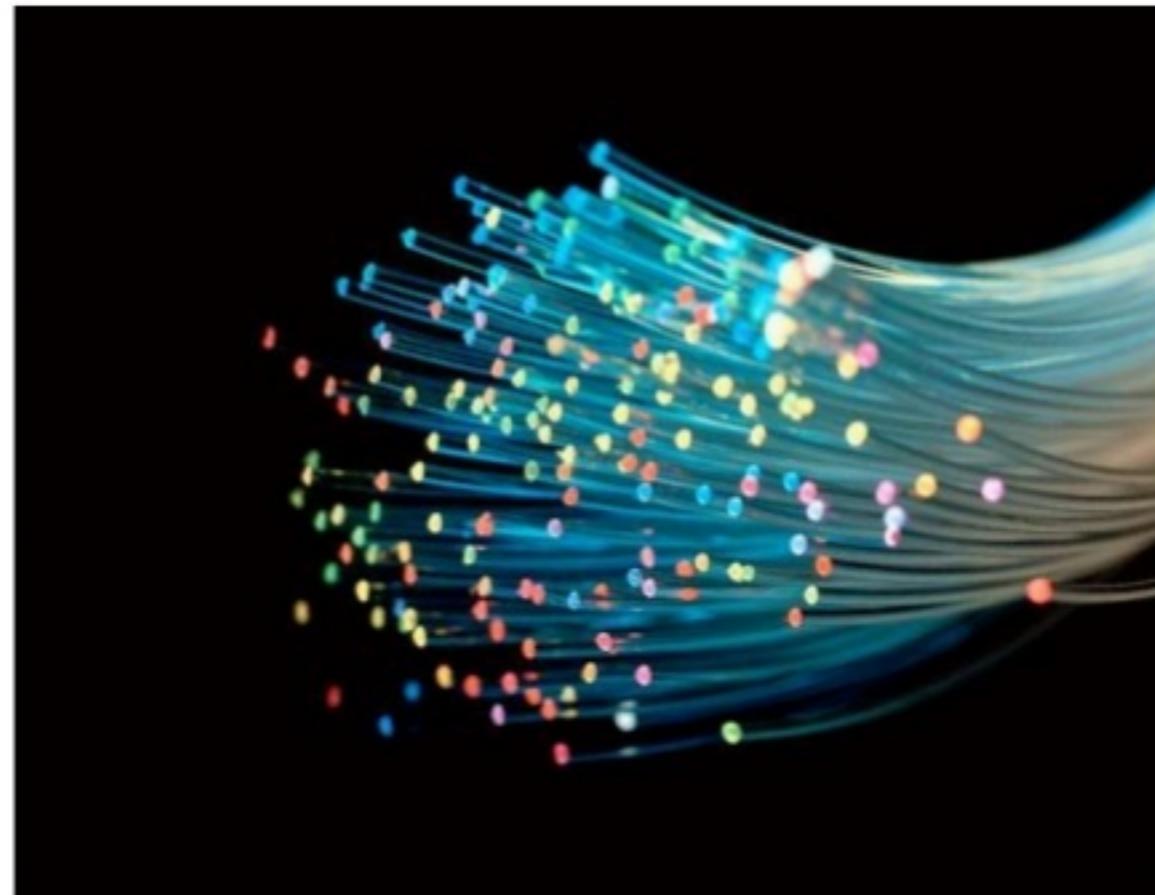


Small Teams = Much Better Communication

Low Latency/High Bandwidth Communication



Legacy

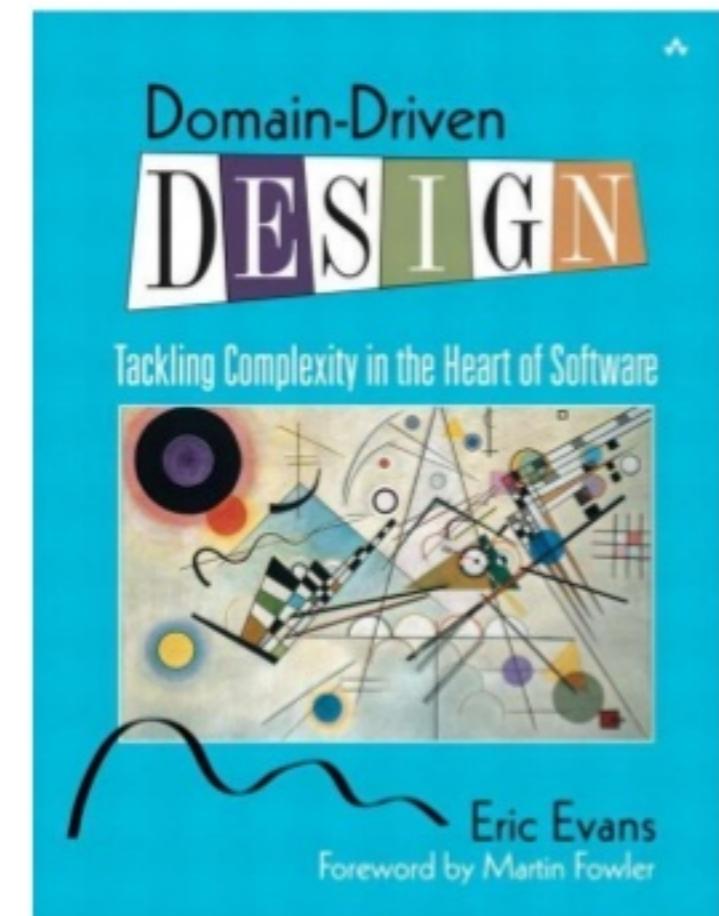


Microservices

Definitive Guide to Deciding Microservice Size

Domain-Driven Design: Tackling Complexity in the Heart of Software by Eric Evans

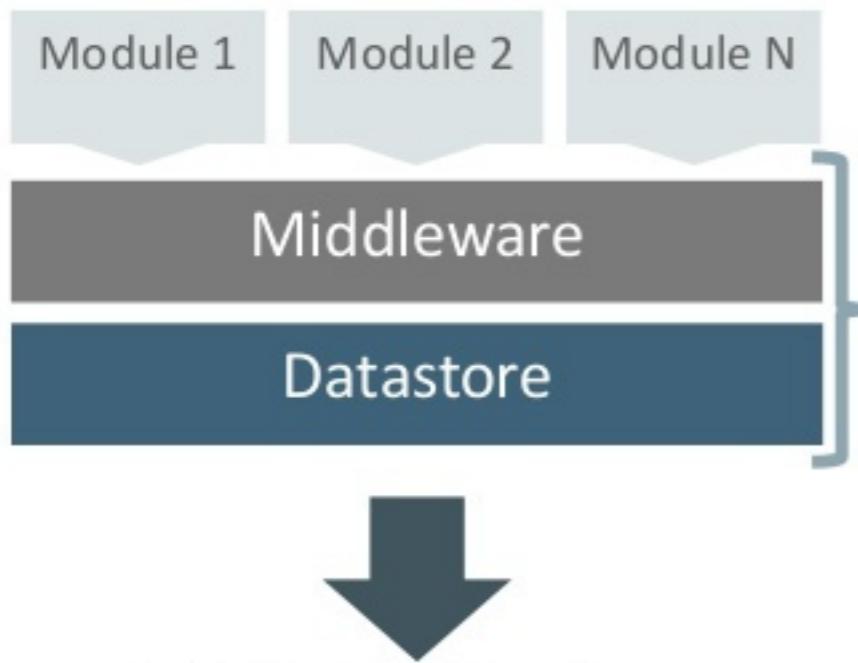
- Bounded Context is a central pattern in Domain-Driven Design
- Deals with designing software based on the underlying domain
- You can't build a big unified domain model for an entire system
- Divides a large system into Bounded Contexts, each of which can have a unified model



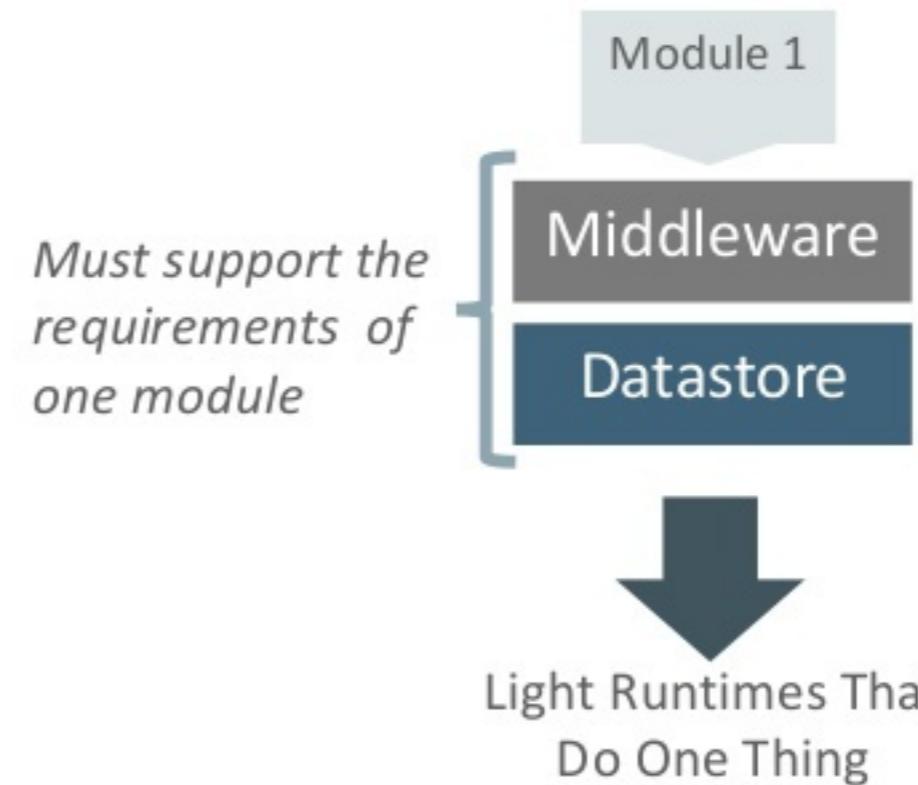
“Micro” in Microservices != Runtime Weight

Microservices *tend* to use smaller runtimes but you can use what you have today

Monoliths



Microservices



You're Doing Microservices If You...

Can build a
microservice
independently

Can release each
microservice
independently

Don't share a
datasource across
microservices



Microservices Are Not a Panacea

Technical Complexity

- Microservices doesn't eliminate complexity - it just moves it and often adds to it
- Monolithic applications allow you to deal with complexity in one body of code
- Forces move to distributed computing

Effort

- Testing, logging, monitoring, security, versioning, etc all become much harder
- Polyglot exponentially increases the number of lifecycles required
- A lot of duplicated effort since each team is independent and goal is to minimize dependencies

Organization

- Most organizations are organized around horizontal technology layers – need to build small product-focused teams
- Much higher skills required
- Many developers will not want to do production support



History of Microservices

Microservices go *way* back



Microservice Principles Have Been With Us For Decades

Reduce Complexity
Through
Modularization

Loose
Coupling

Do One Thing and
Do It Well

Focus on Business
Capabilities, Not
Technology Layers

*The principles behind microservices are
often just good architecture principles*



History of Microservices as a Term

A trend without a name, until “microservices” was used in 2013

- First Google search of "microservices" occurred in October of 2013
 - expanded exponentially from there
- Began in the mid to late 2000's as a reaction against monoliths and against traditional SOA
- No one person is credited with inventing the term or popularizing it. It was an idea whose time had come

2007

2009

2011

2013

2015



SOA vs. Microservices

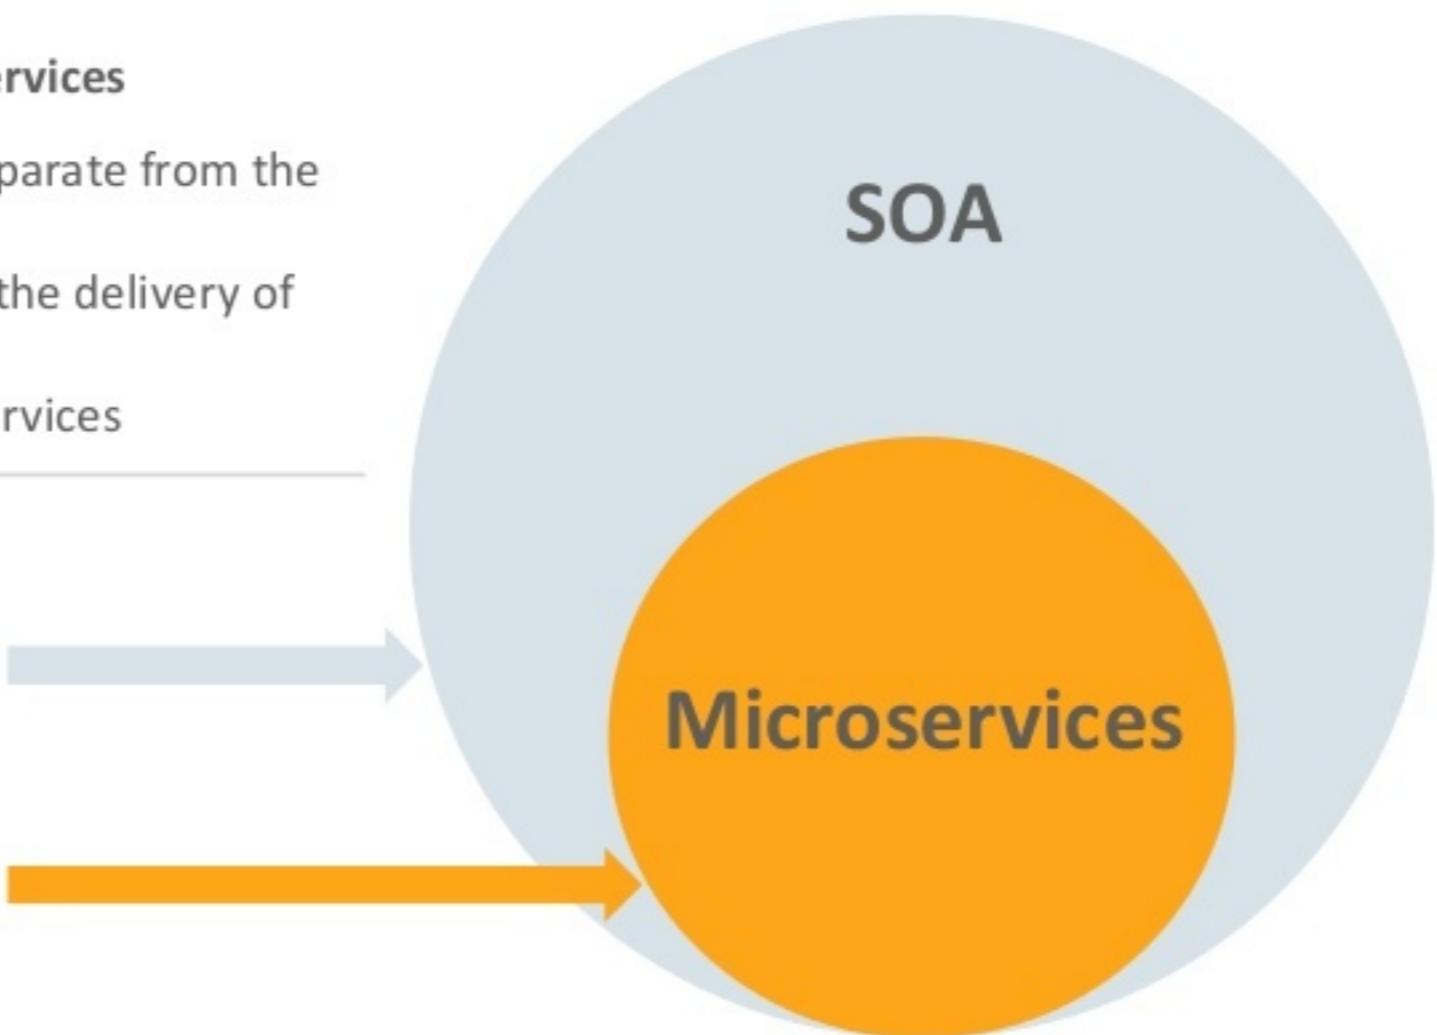
SOA is the general idea, where microservices are a very specific way of achieving it

All of the tenets of SOA also apply to microservices

1. Keeping consumption of services separate from the provisioning of services
2. Separating infra management from the delivery of application capability
3. Separating teams and decoupling services

Implementation Differences

- Favors centralized orchestration
- Needlessly complicated by SOAP
- “Dumb endpoints, smart pipes”



- Favors distributed choreography
- REST + HTTP/S = simple
- “Smart endpoints, dumb pipes”

SOA vs. Microservices Misconceptions



“Microservices removes the need for an Enterprise Service Bus”



Don't confuse the product with the pattern

“Microservices solves the problems of SOA”



Don't confuse improper SOA deployments as problems with SOA

“Companies like Netflix and LinkedIn use microservices, so we should too”



Netflix and LinkedIn are in the platform business. Is that your business too?

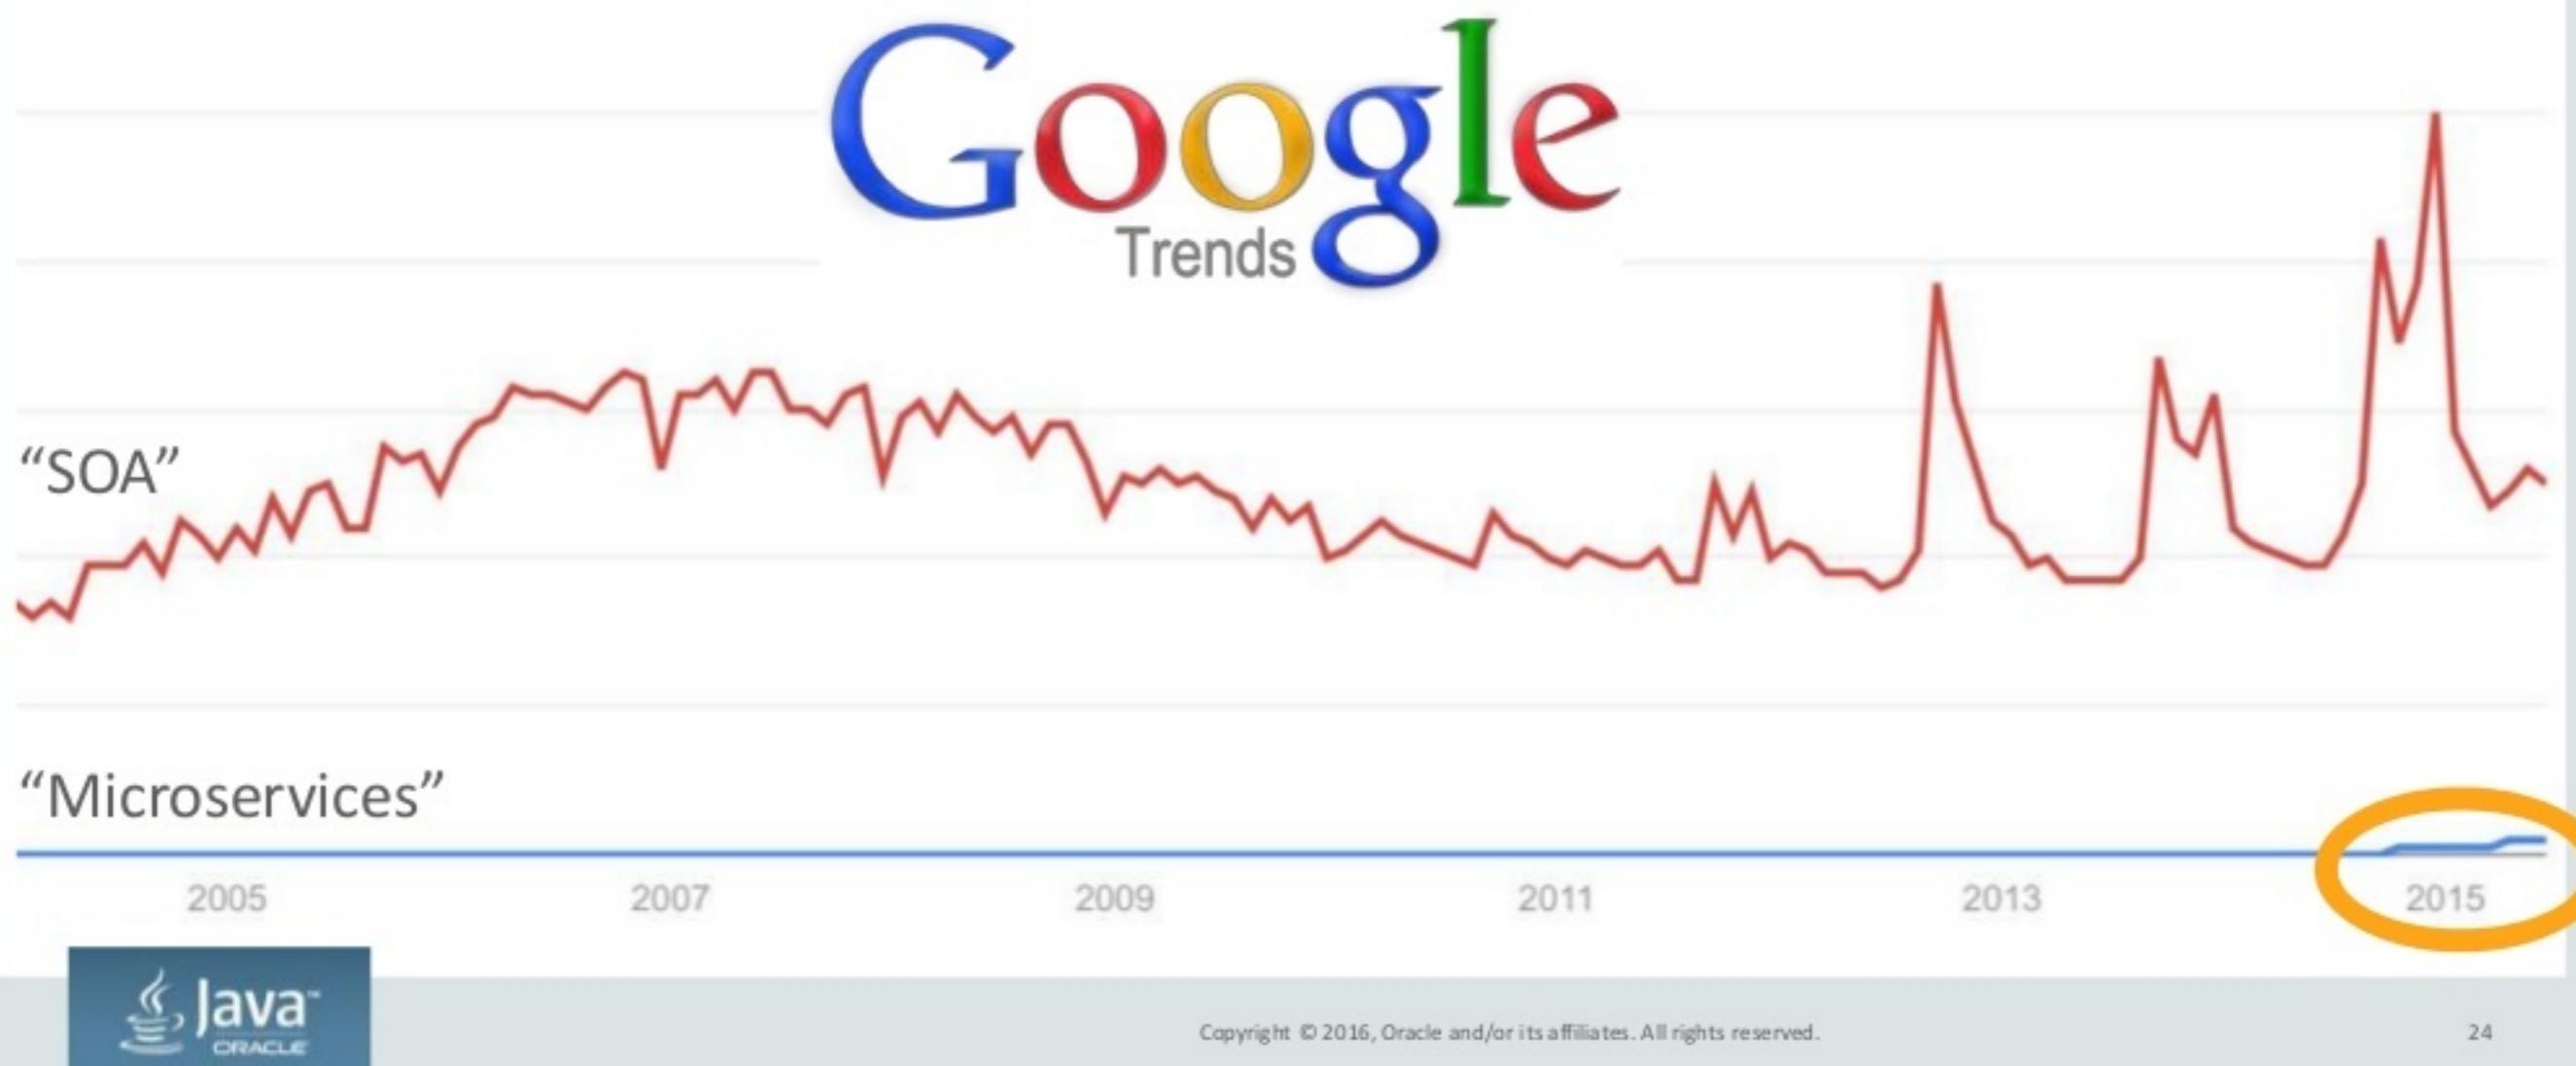
“We must choose microservices, or SOA”



Use both

SOA Still Reigns Supreme

SOA isn't going anywhere



Microservice Principles Are Old; Implementation is New

Microservices is not just rebranded SOA

- Teams independently architect, develop, deploy and maintain each microservice
- Each microservice often has its own datastore, which may not always be 100% up-to-date
- Microservices is fully decentralized - no ESBs, no single database, no top-down anything
- Responses from microservices are not manipulated by an intermediary, like an ESB
- Microservices favors simple transports - XML or JSON over HTTP/S. No SOAP
- Any given instance of a microservice is stateless - state, config, and data pushed externally
- Microservices support polyglot - each microservice team is free to pick the best technology
- DevOps principles - automated setup and developers owning production support
- Use of containers, which allow for simple app packaging and fast startup time
- Use of cloud, for the elastic infrastructure, platform and software services



Non-technical Prerequisites

Conway's Law

*“Any piece of software reflects
the organizational
structure that produced it”*

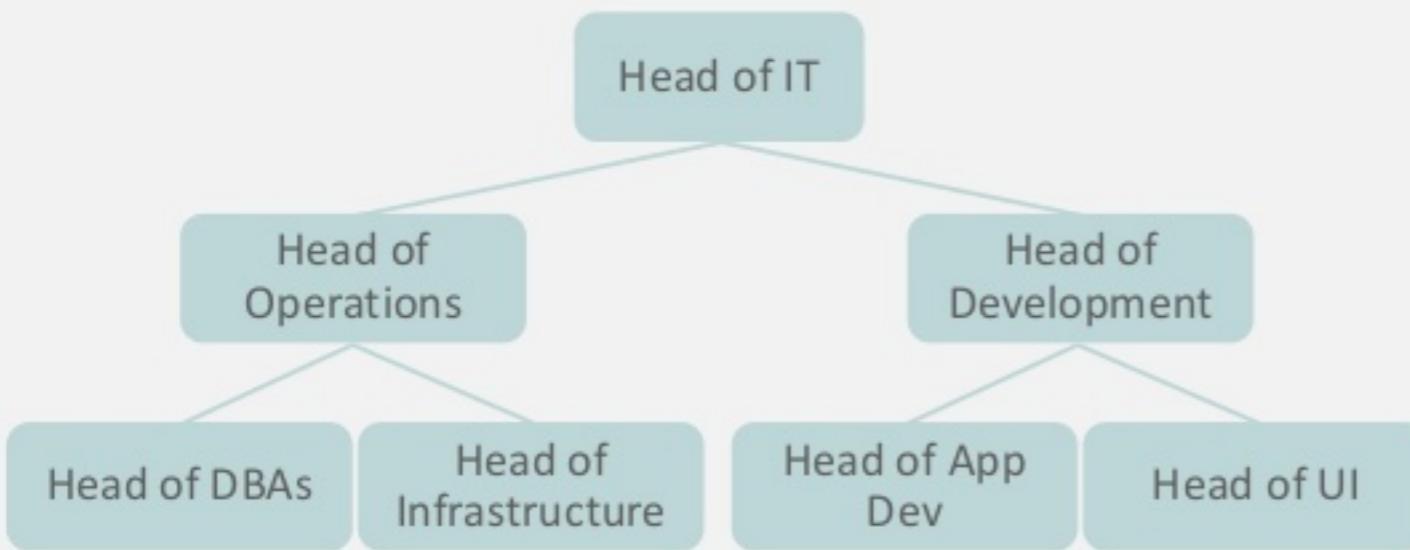


Melvin Conway
1968

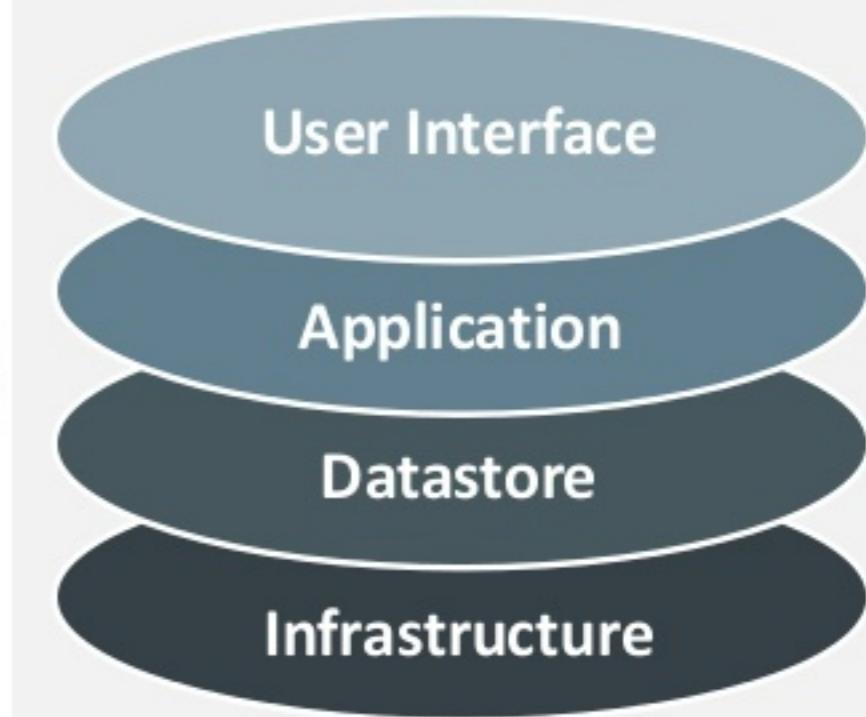
Conway's Law In Action

Any piece of software reflects the organizational structure that produced it

Typical Enterprise Organization Structure



Resulting Software



An Enormous Monolith

Characteristics of Different Organization Types

Centralized Organizations Focused on Technology Layers

- Produce monoliths
- Simple change requires extensive coordination across all of the different layers
- Business logic is spread everywhere because it's easier to bury it in the layer you own
- No real ownership

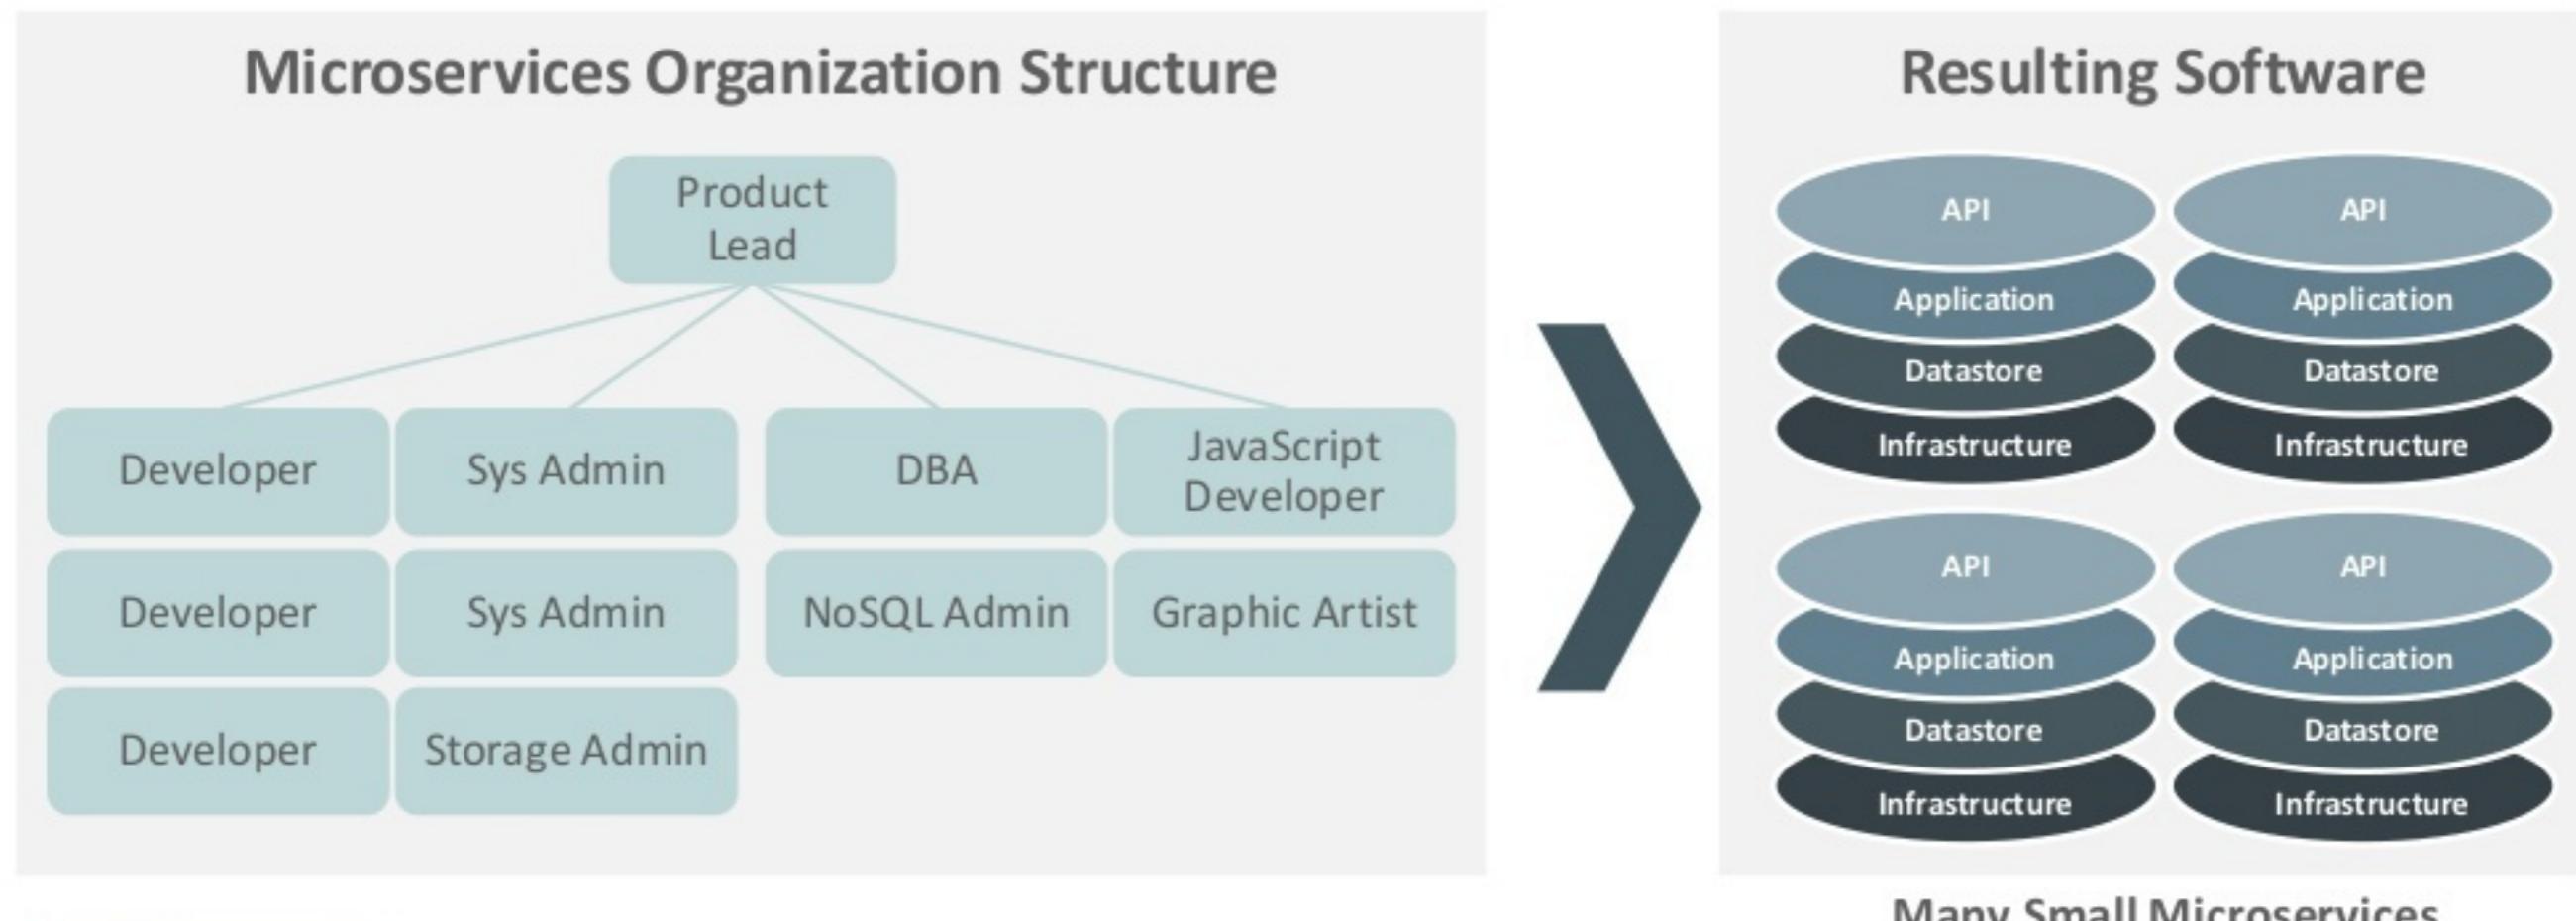
Distributed Organizations Focused on Products

- Produce microservices
- Small teams can make any change they want to an individual microservice
- Architecture clean because developers have 100% control
- True ownership

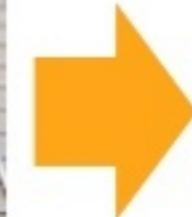


Successful Teams Structure Their Teams Around Products

Build small vertical teams

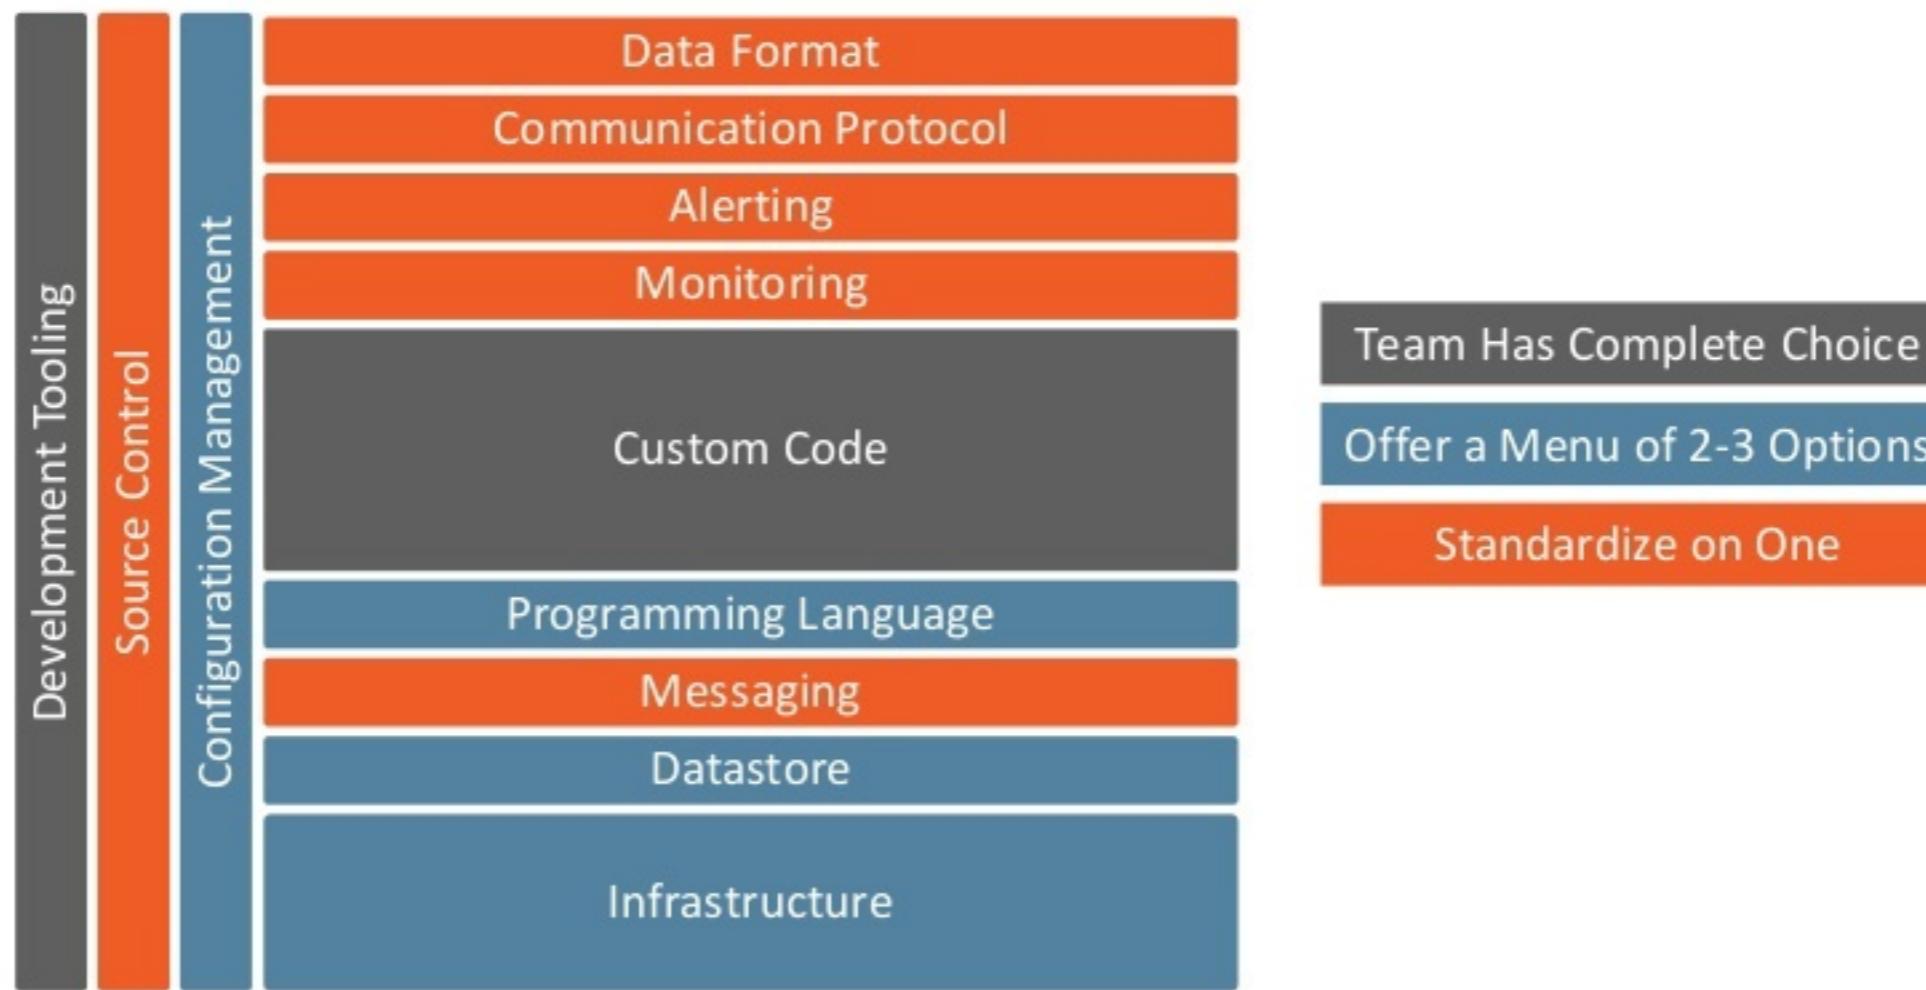


Horizontal -> Vertical Teams



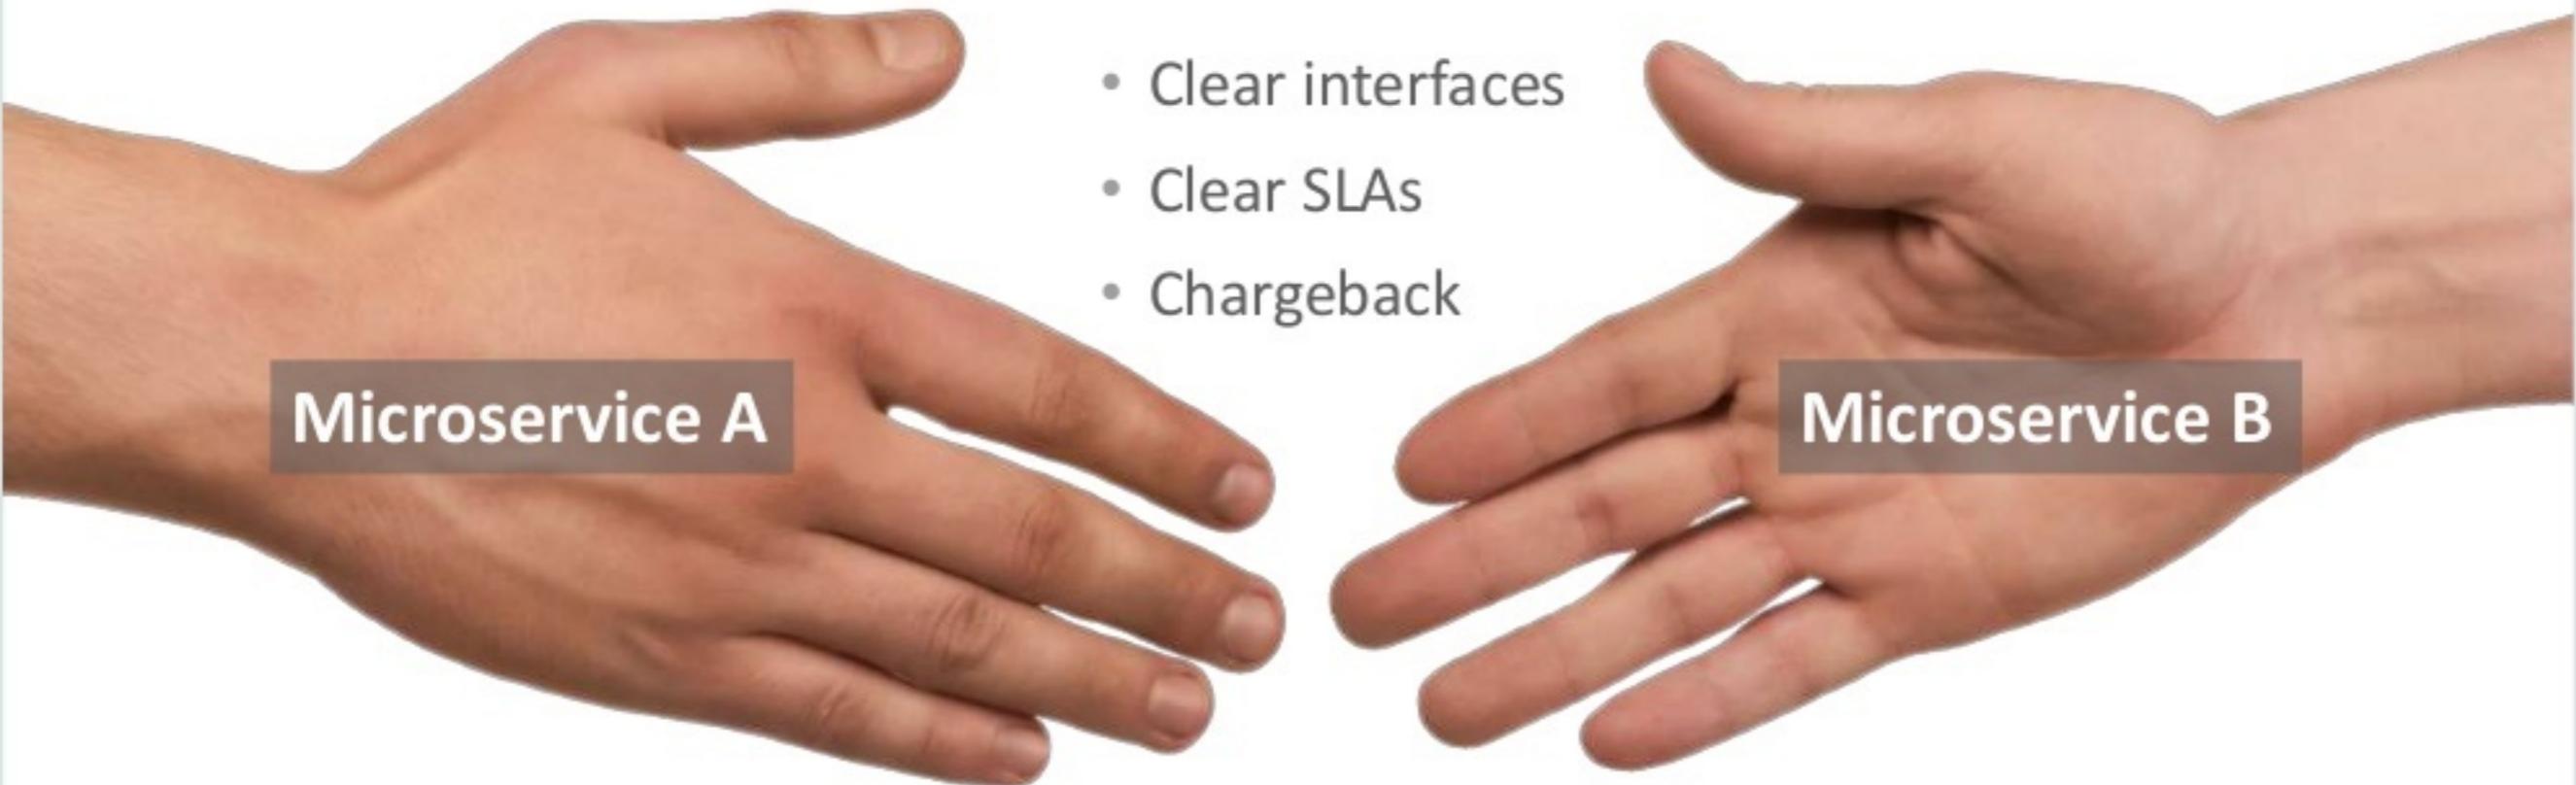
Standardize Where it Makes Sense. Be Pragmatic

Teams do not have 100% freedom



Each Team Should Treat Other Teams as External Vendors

Good fences make good neighbors



Microservice A

- Clear interfaces
- Clear SLAs
- Chargeback

Microservice B

DevOps Must Also Be Adopted Prior to Microservices

Mostly a function of culture

Culture

Technology

Respect

- Dev respect for ops
- Ops respect for dev

Infra as Code

- Don't build envs by hand
- Scripts checked in and managed as src

Discuss

- Ops should be in dev discussions
- Dev should be in ops discussions
- Shared runbooks

Shared Version Control

- Single system
- Ship trunk
- Enable features through flags

Avoid Blaming

- No fingerpointing!
- Everyone should have some culpability

One Step Build/Deploy

- One button build/deploy
- If verification fails, stop and alert or take action

"Done" Means Released

- Dev's responsibility shouldn't ever end – production support required
- "Throwing it over the wall" is dead

Don't Fix Anything

- If something breaks, re-deploy.
- Don't fix
- Fix environment setup scripts



More About DevOps

<http://www.slideshare.net/KellyGoetsch/mastering-devops-with-oracle>



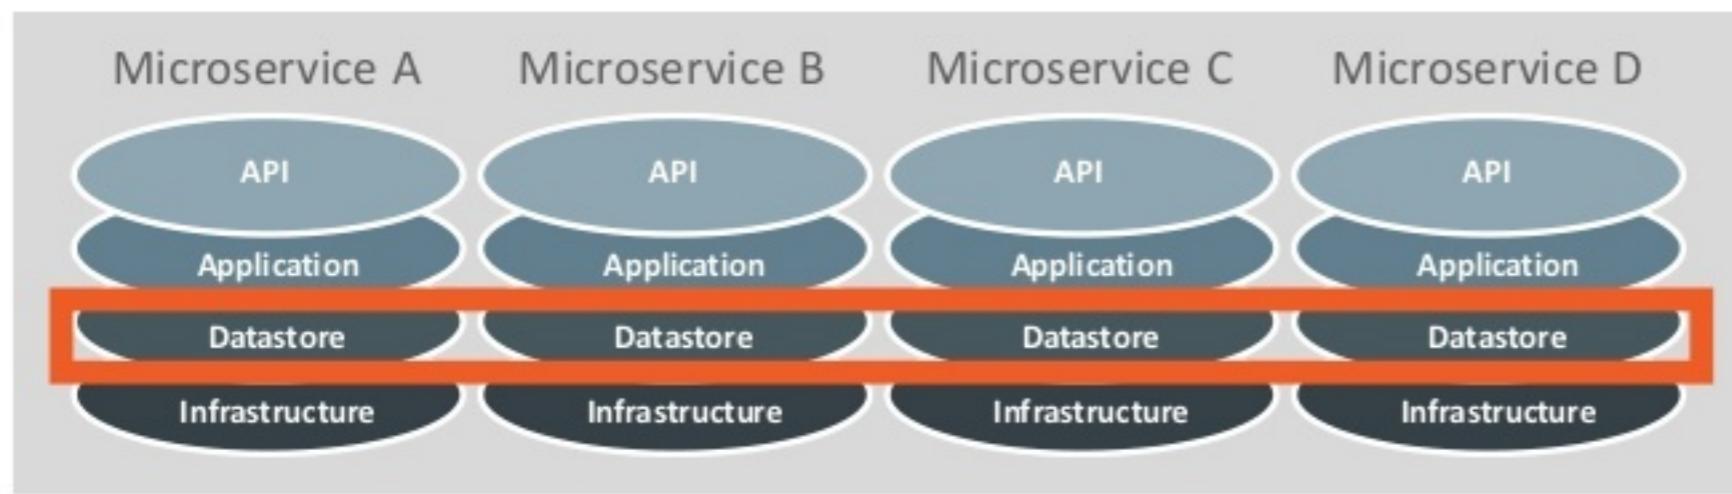


Architectural Prerequisites



Microservices Forces Move To Distributed Computing

Introduces enormous complexity – monoliths don't suffer from this



- Distributed computing is a natural consequence of microservices because each microservice has its own datastore
- Sharing datastores across microservices introduces coupling – very bad!
- There will always be latency between microservices
- Latency = eventual consistency

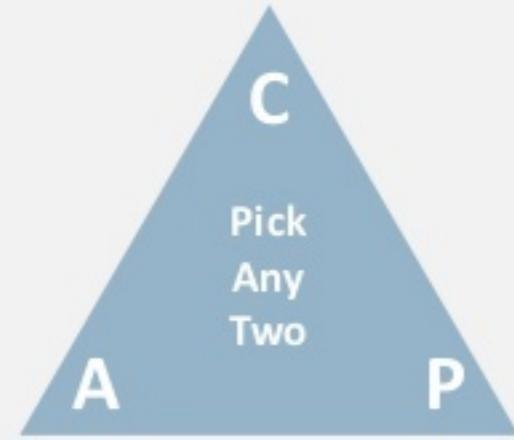
- All data exchange between microservices must be through API layer or messaging – no accessing datastores cross-microservices
- Must implement high-speed messaging between microservices. REST + HTTP probably isn't fast enough
- May end up duplicating data across datastores – e.g. a customer's profile

Rules of Distributed Computing

Computer science is about tradeoffs

Theory

CAP Theorem – Pick Any Two



Consistency

Each node shows the same data at all times

Availability

Each node is available for writes at all times

Partition Tolerance

Able to handle network outages



More Information

Practice

Pick One

Partition Tolerance is non-negotiable because we have networks that can always fail

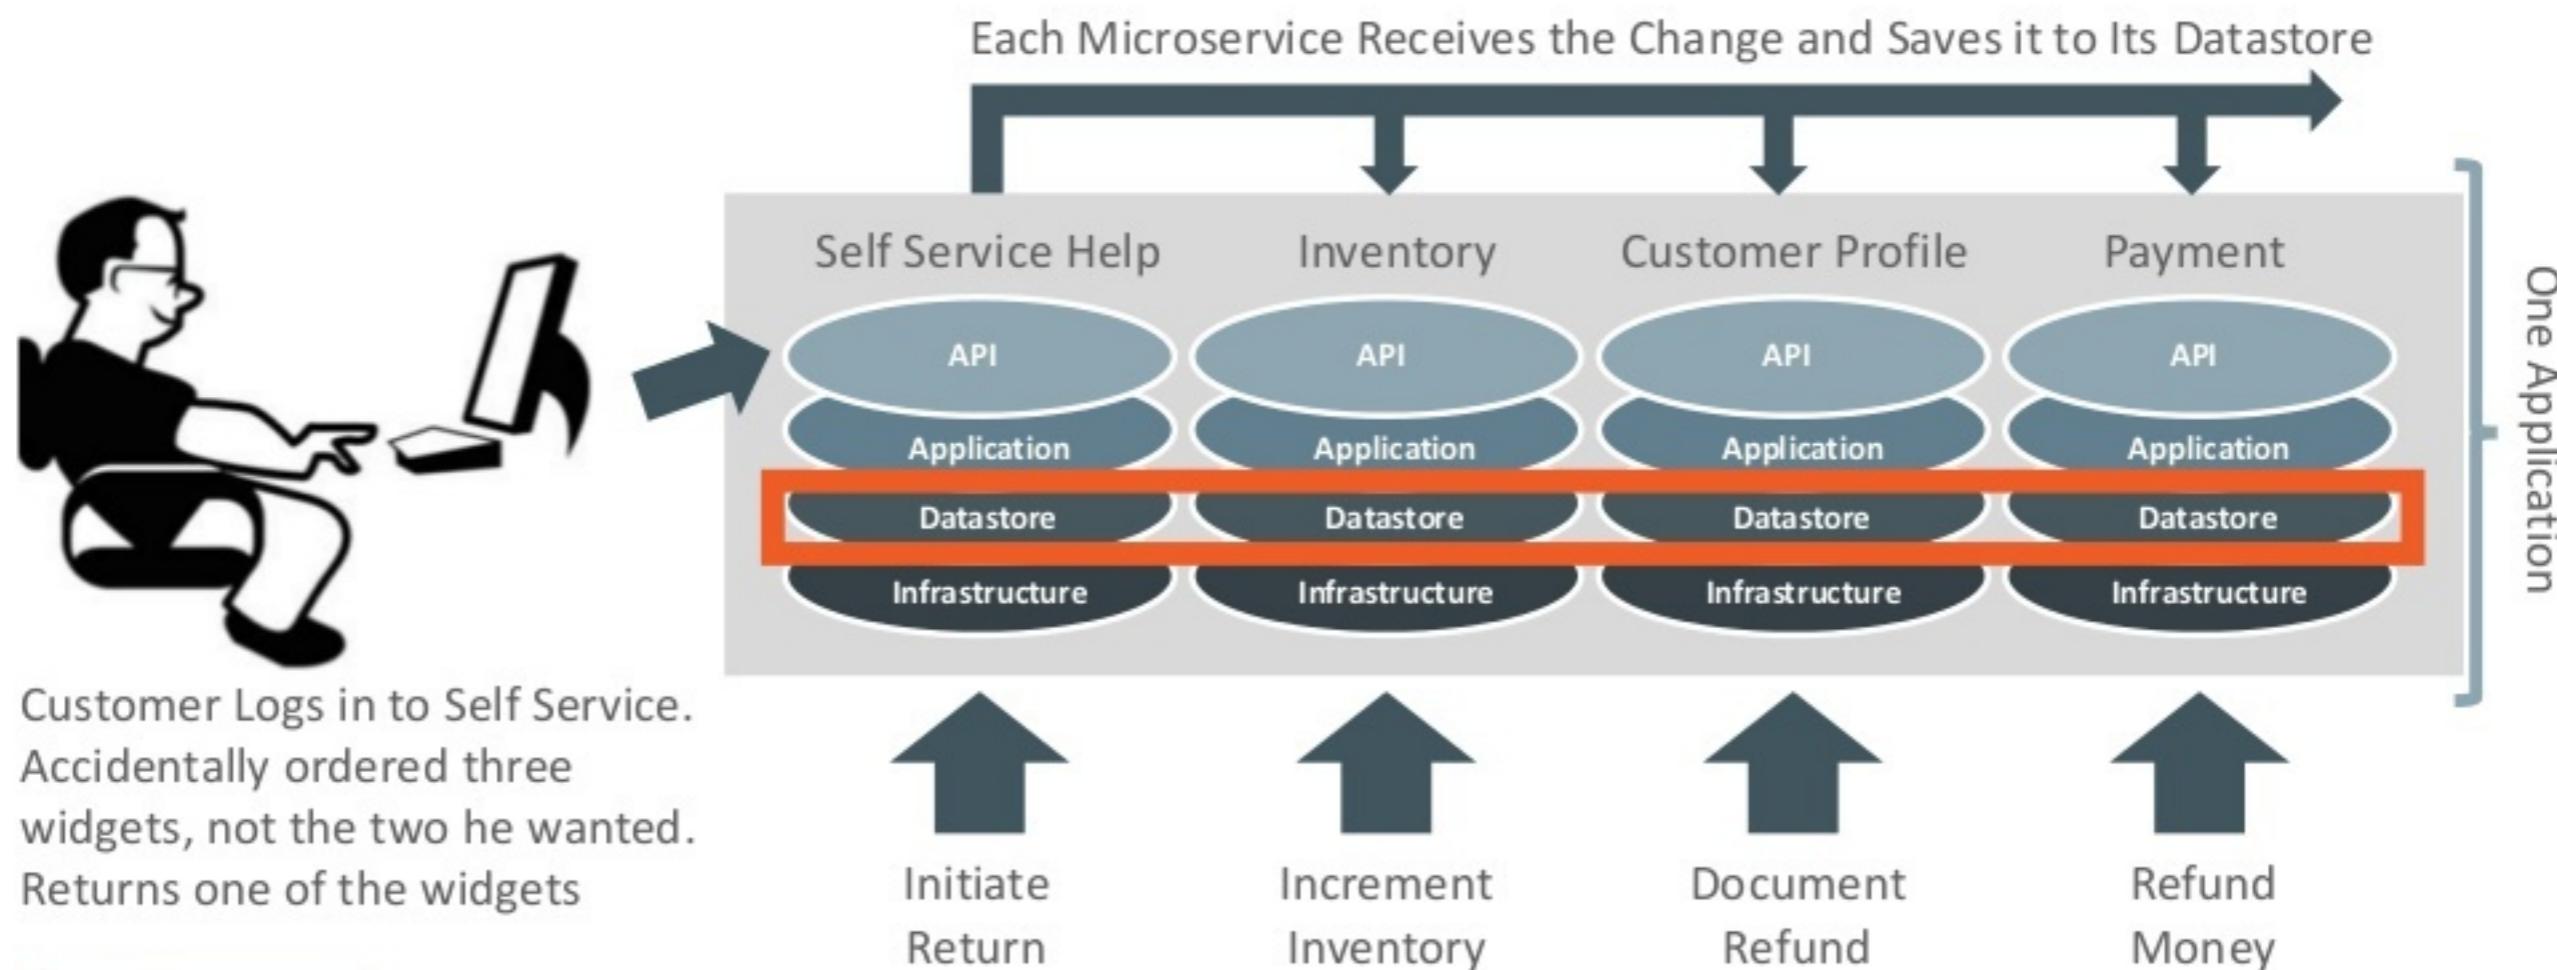
Enterprise IT Systems: Often CP

Microservice Systems: Often AP

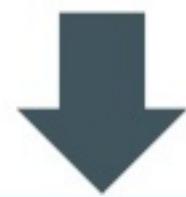
Each microservice can be CP, AP or CA but the system as a whole is always CP or AP

Microservices Usually Forces Eventual Consistency

Synchronous communication leads to availability and performance problems



Synchronous Calls With Microservices Are Very Bad



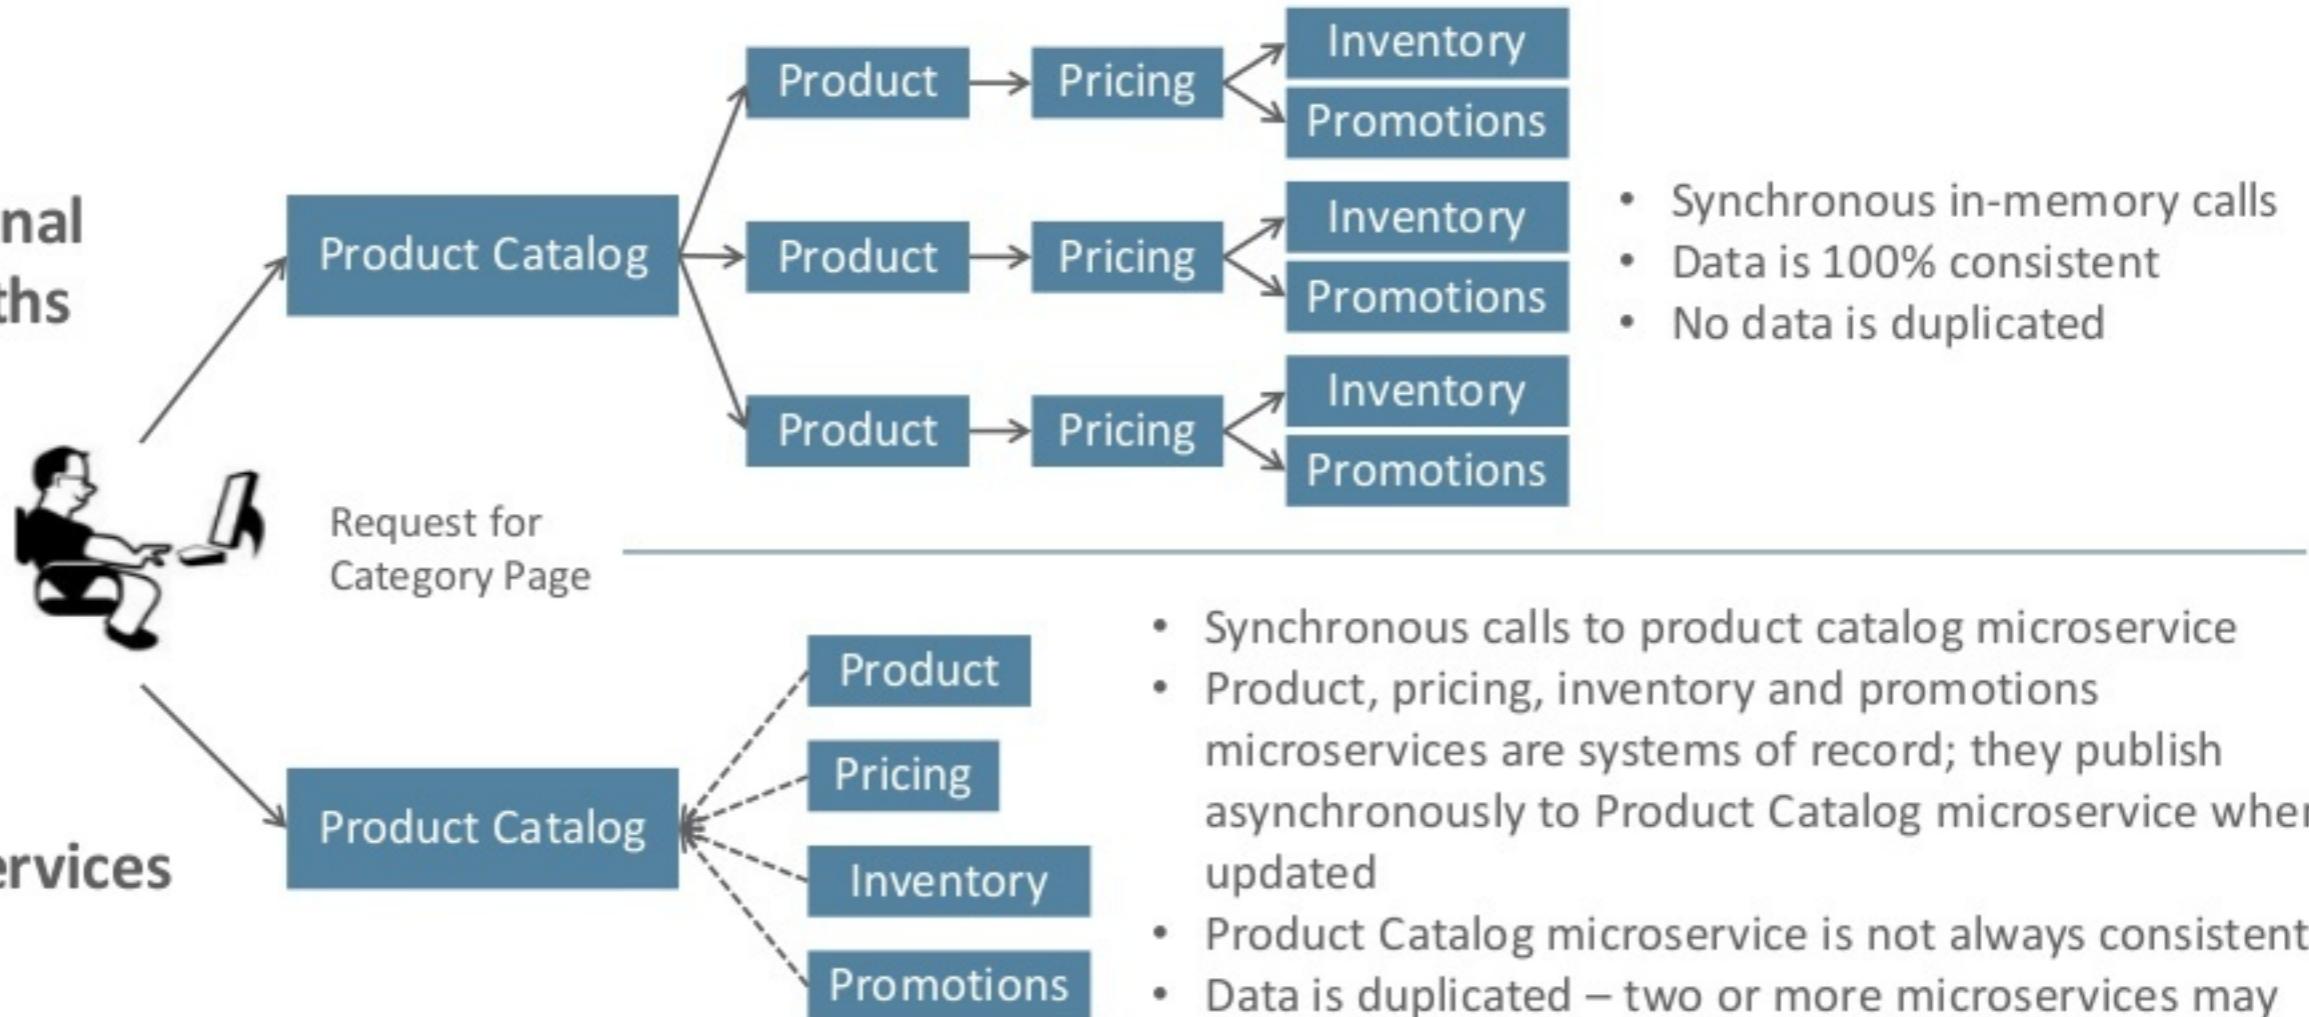
Chaining == coupling == downtime

The availability of microservice A depends on B, B depends on C, etc



Avoid Synchronous Calls for Read-only Data By Copying

Traditional Monoliths



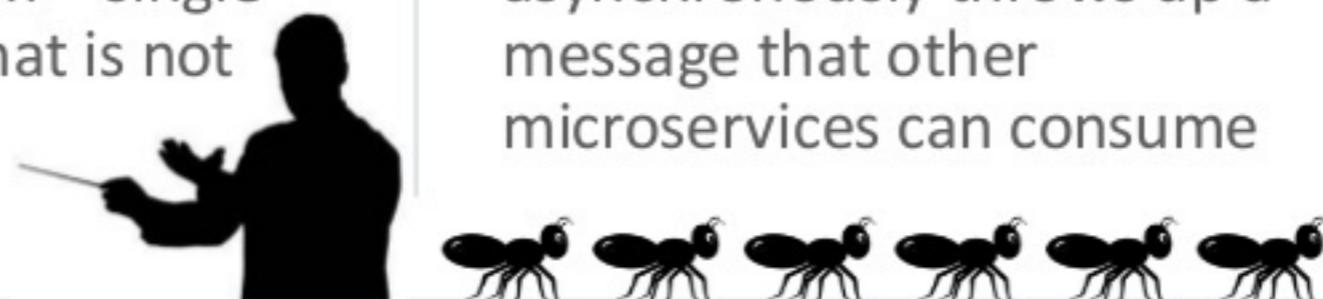
Choreography Tends to Be Better Than Orchestration

Orchestration

- Top-down coordination of discrete actions
- Used in centralized, monolithic applications
- Brittle – centralized by nature
- Each “action” registers with centralized system – single point of failure that is not very flexible

Choreography

- Bottom-up coordination of discrete actions
- Used in distributed, microservice applications
- Resilient – distributed by nature
- Each microservice asynchronously throws up a message that other microservices can consume



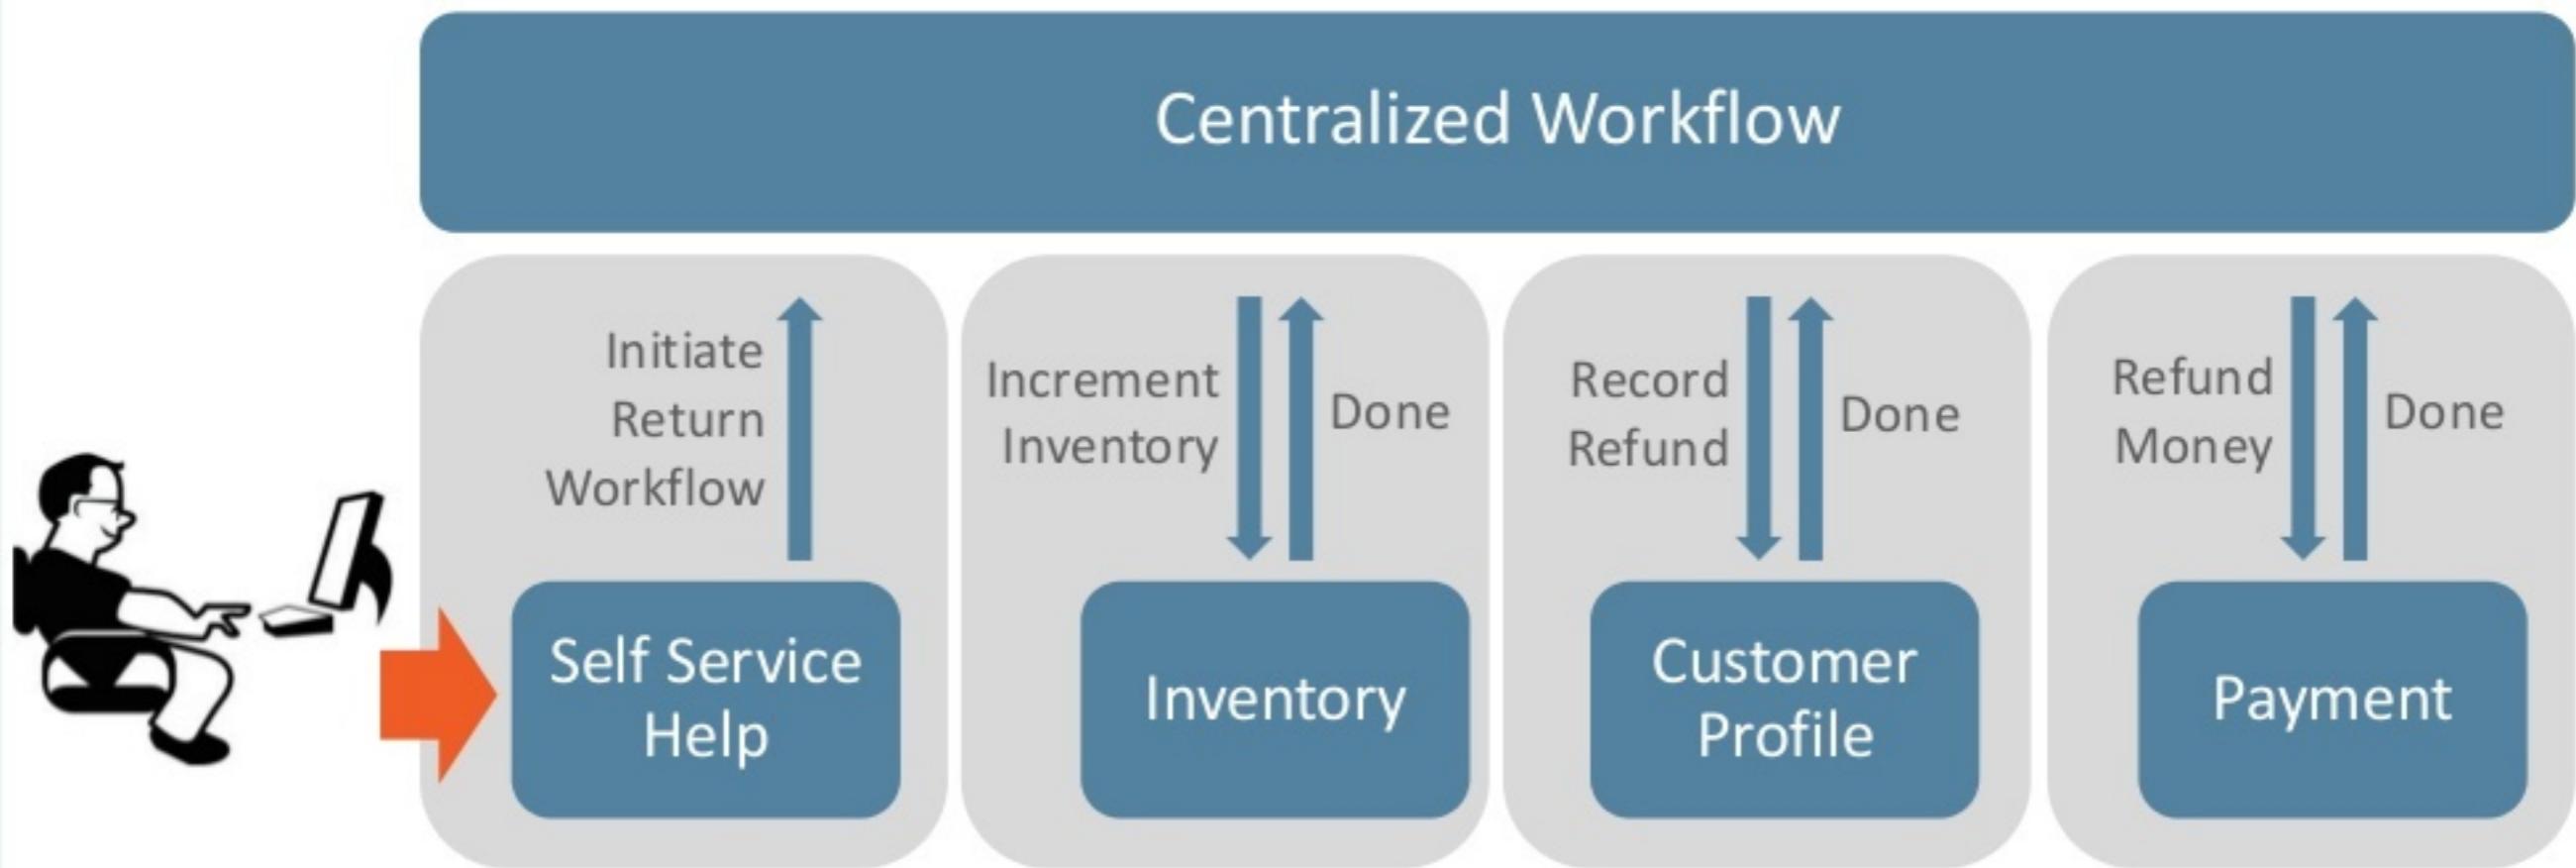
Orchestration vs. Choreography





Example of Orchestration

Scenario: eCommerce user returns a widget through web-facing .com



Brittle | Centralized | Tightly Coupled



Example of Choreography

Scenario: Inventory microservice

Events This Microservice Cares About

- New Inventory
- Product Sold
- Product Returned
- Product Recalled



Events This Microservice Emits

- Inventory Incremented
- Inventory Decrement
- Inventory Created
- Inventory Deleted

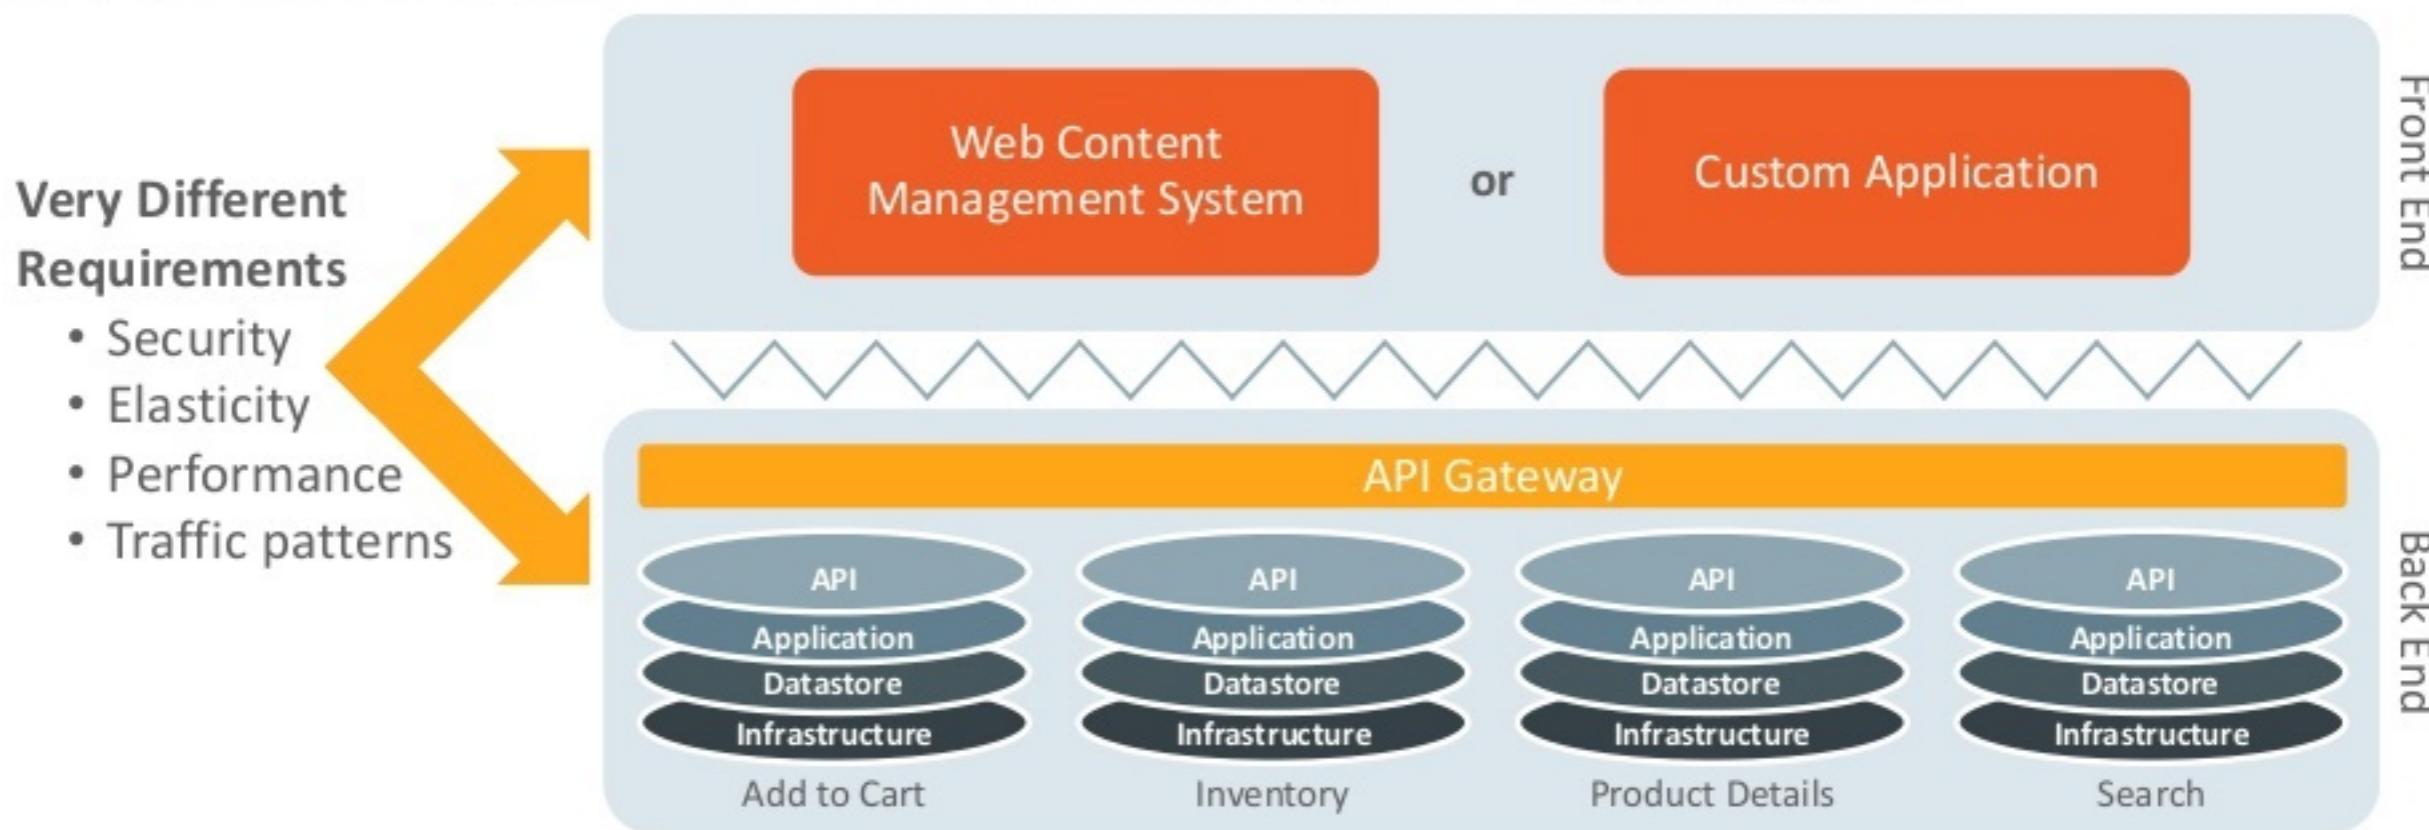
For Anyone Who Cares...

All asynchronous



Best to Ship Your Applications Headless

Put your front and back ends in different clouds, different geographies

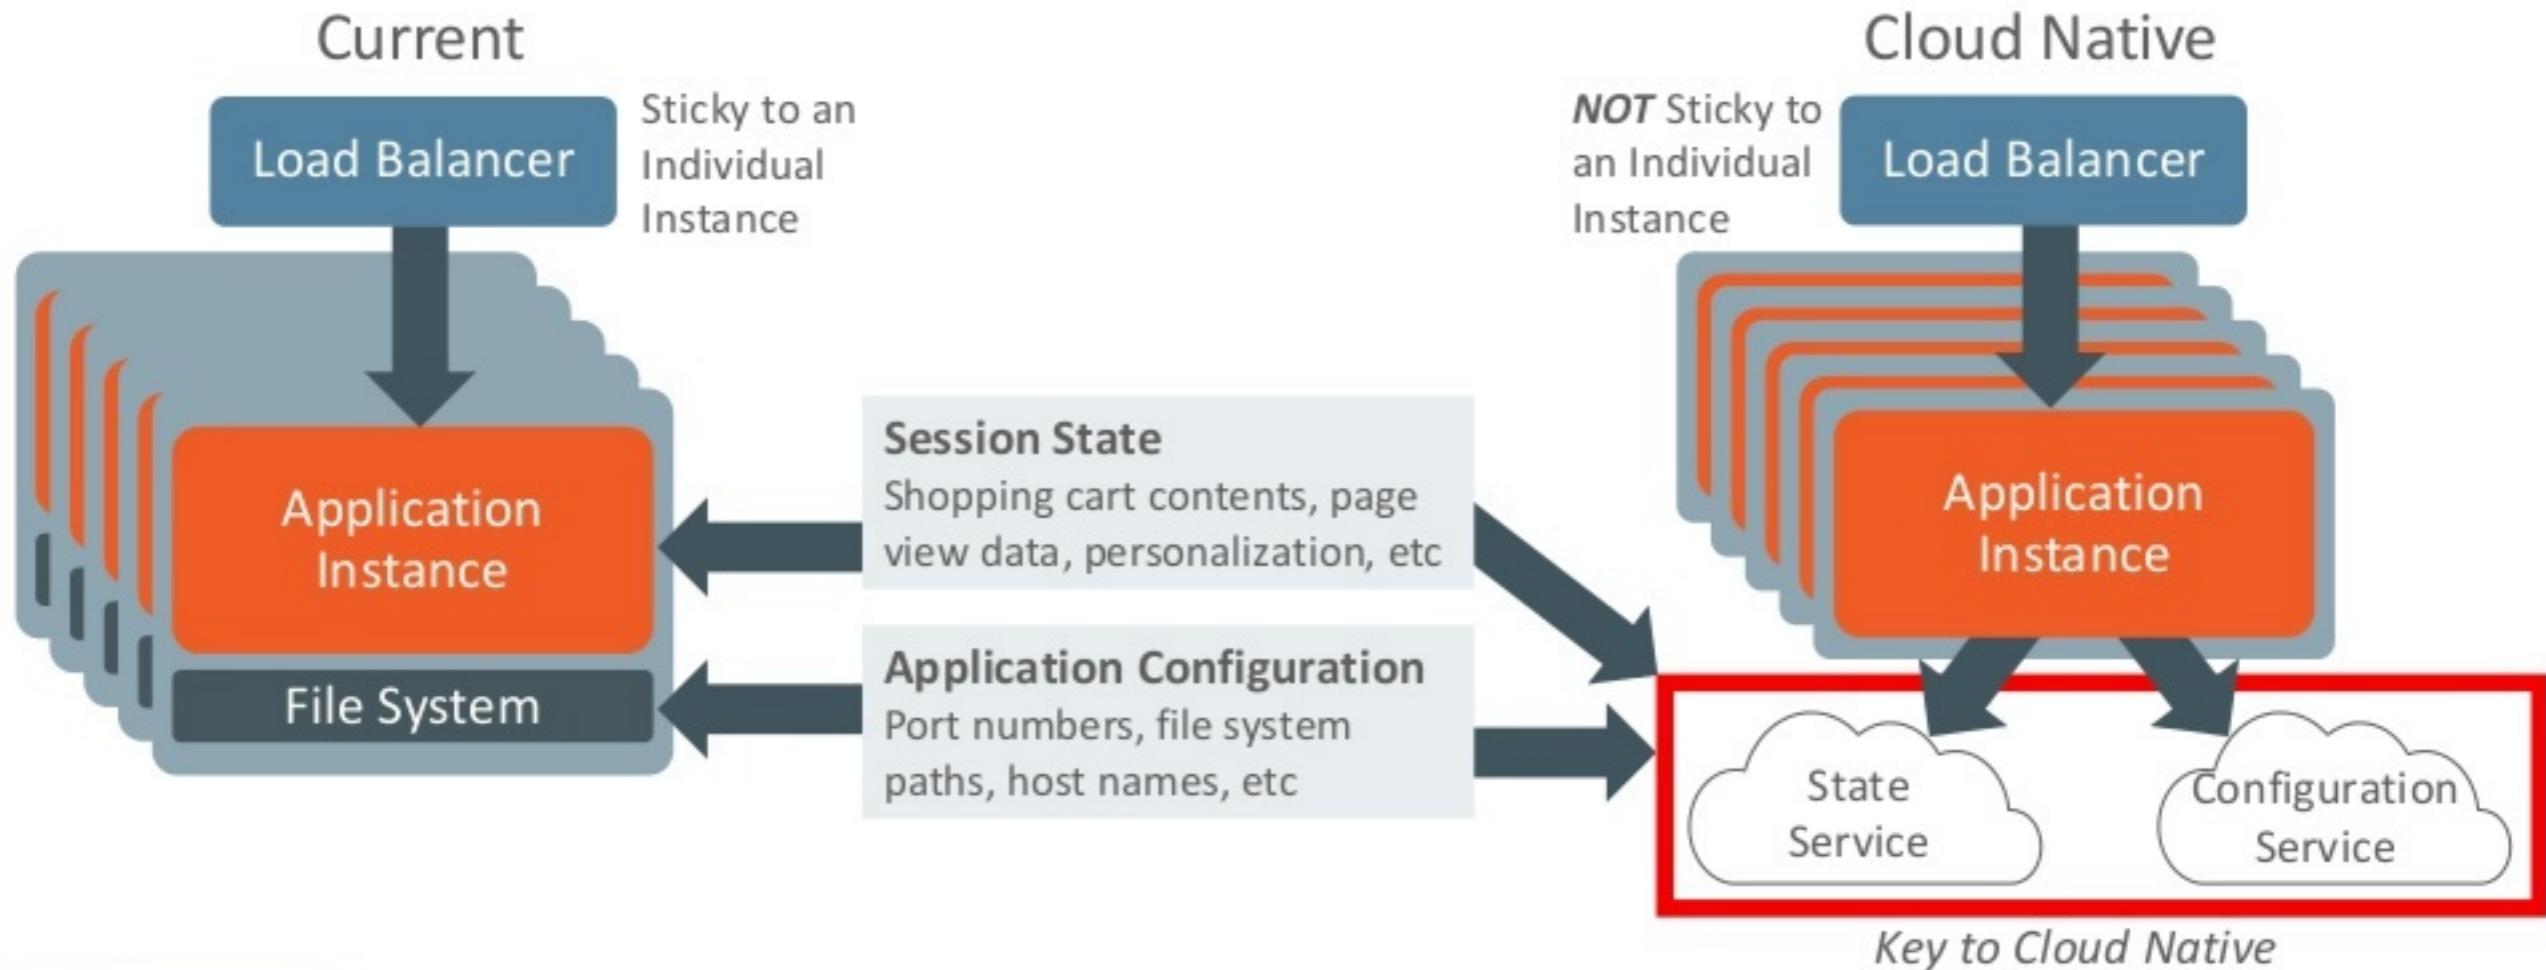


Design, develop, deploy and manage your front and back ends differently



Make Your Middle Tier Stateless If You Can

Push all state and configuration down to highly available cloud services



Remove All Hard-coded IPs, Host Names, etc

Use service discovery, DNS, etc instead. *Everything* should be dynamic



Modify Messaging To Be Less Fragile

Messages should be constructed to be able to be applied over and over

Not Idempotent

```
<credit>
  <amount>100</amount>
  <forAccount>1234</account>
</credit>
```

- Can execute exactly once
- Works fine in a monolith
- Will break with microservices

Best For Microservices

Idempotent

```
<credit>
  <amount>100</amount>
  <forAccount>1234</account>
  <creditMemoID>4567</creditMemoID>
</credit>
```

- Can execute many times
- Works fine in monolith and microservices
- Allows the pipes to be dumb – message just needs to be applied once

* Idempotent = the outcome doesn't change after the first application

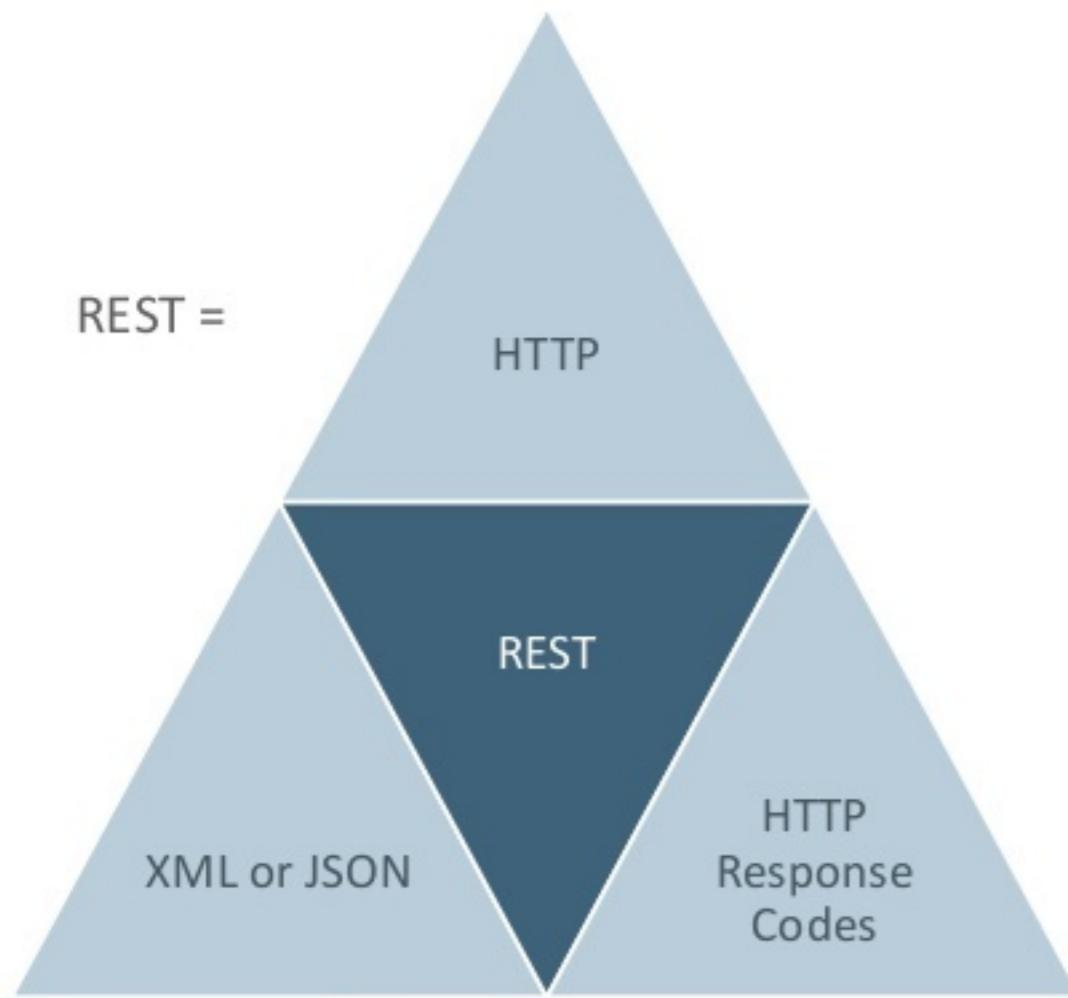


Approaches to Synchronous Network Calls

| | XML/JSON Over HTTP | Binary Over Wire |
|------------------------|---|---|
| Primary Use | Communicating with clients over the public internet | Communicating with other microservices over a private network |
| Pros | <ul style="list-style-type: none">▪ Universally understood format▪ Easy to implement and understand | <ul style="list-style-type: none">▪ Very fast |
| Cons | <ul style="list-style-type: none">▪ Slow since it's text-based | <ul style="list-style-type: none">▪ Can be hard to implement |
| Implementations | <ul style="list-style-type: none">▪ No special software required – natively supported by all major programming languages▪ HTTP is the language of the web! | <ul style="list-style-type: none">▪ Oracle Portable Object Format▪ Google Protocol Buffers▪ Apache Avro▪ Apache Thrift |

REST: Representational State Transfer

Strongly associated with microservices but not a technical requirement



- Much simpler alternative to SOAP
- Uses GET, POST, PUT, DELETE, etc – just like web browsers do
- Synchronous inter-microservice communication often occurs over binary
- Can version APIs - /v1.2/customer
- Can use XML or JSON
 - XML is often better - supports XPath, CSS selectors
- Can't generate strongly typed stubs

Microservices Requires a Higher Level of REST Use

Oracle fully supports level 3



- **Level 3**

- Application itself tells client how to interact with it - similar to hyperlinks in a web page
- <link rel = "delete" uri = "/OrderService/12345/delete"/> under <order id="12345">

- **Level 2**

- Start using HTTP verbs - HTTP GET/POST/PUT/DELETE/etc
- Responses come back using correct HTTP codes – e.g. HTTP 409 for a conflict

- **Level 1**

- Start requesting individual resources
- Interact with /OrderService/12345

- **Level 0**

- Use HTTP as a tunneling mechanism only
- Interact with /OrderService
- Use HTTP GET/POST only

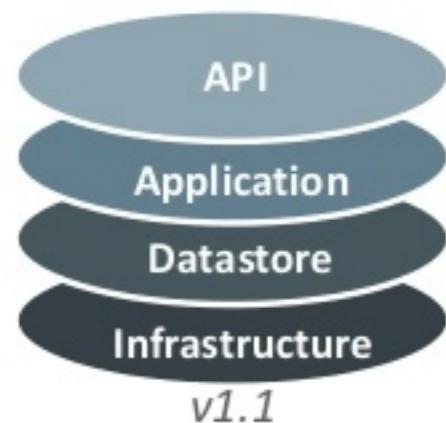
{JSON}

<XML>

Java™
ORACLE

Message Compatibility Across Versions

Design for backwards compatibility

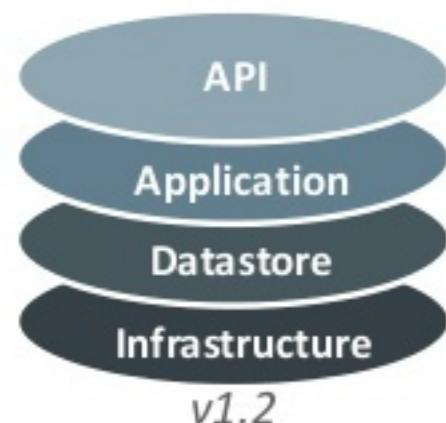


Send update to Product Catalog microservice

```
<product>
  <id>329340224</id>
</product>
```

Don't make properties required!

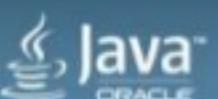
Assume any given microservice has two or more different versions running concurrently. Build for backwards compatibility



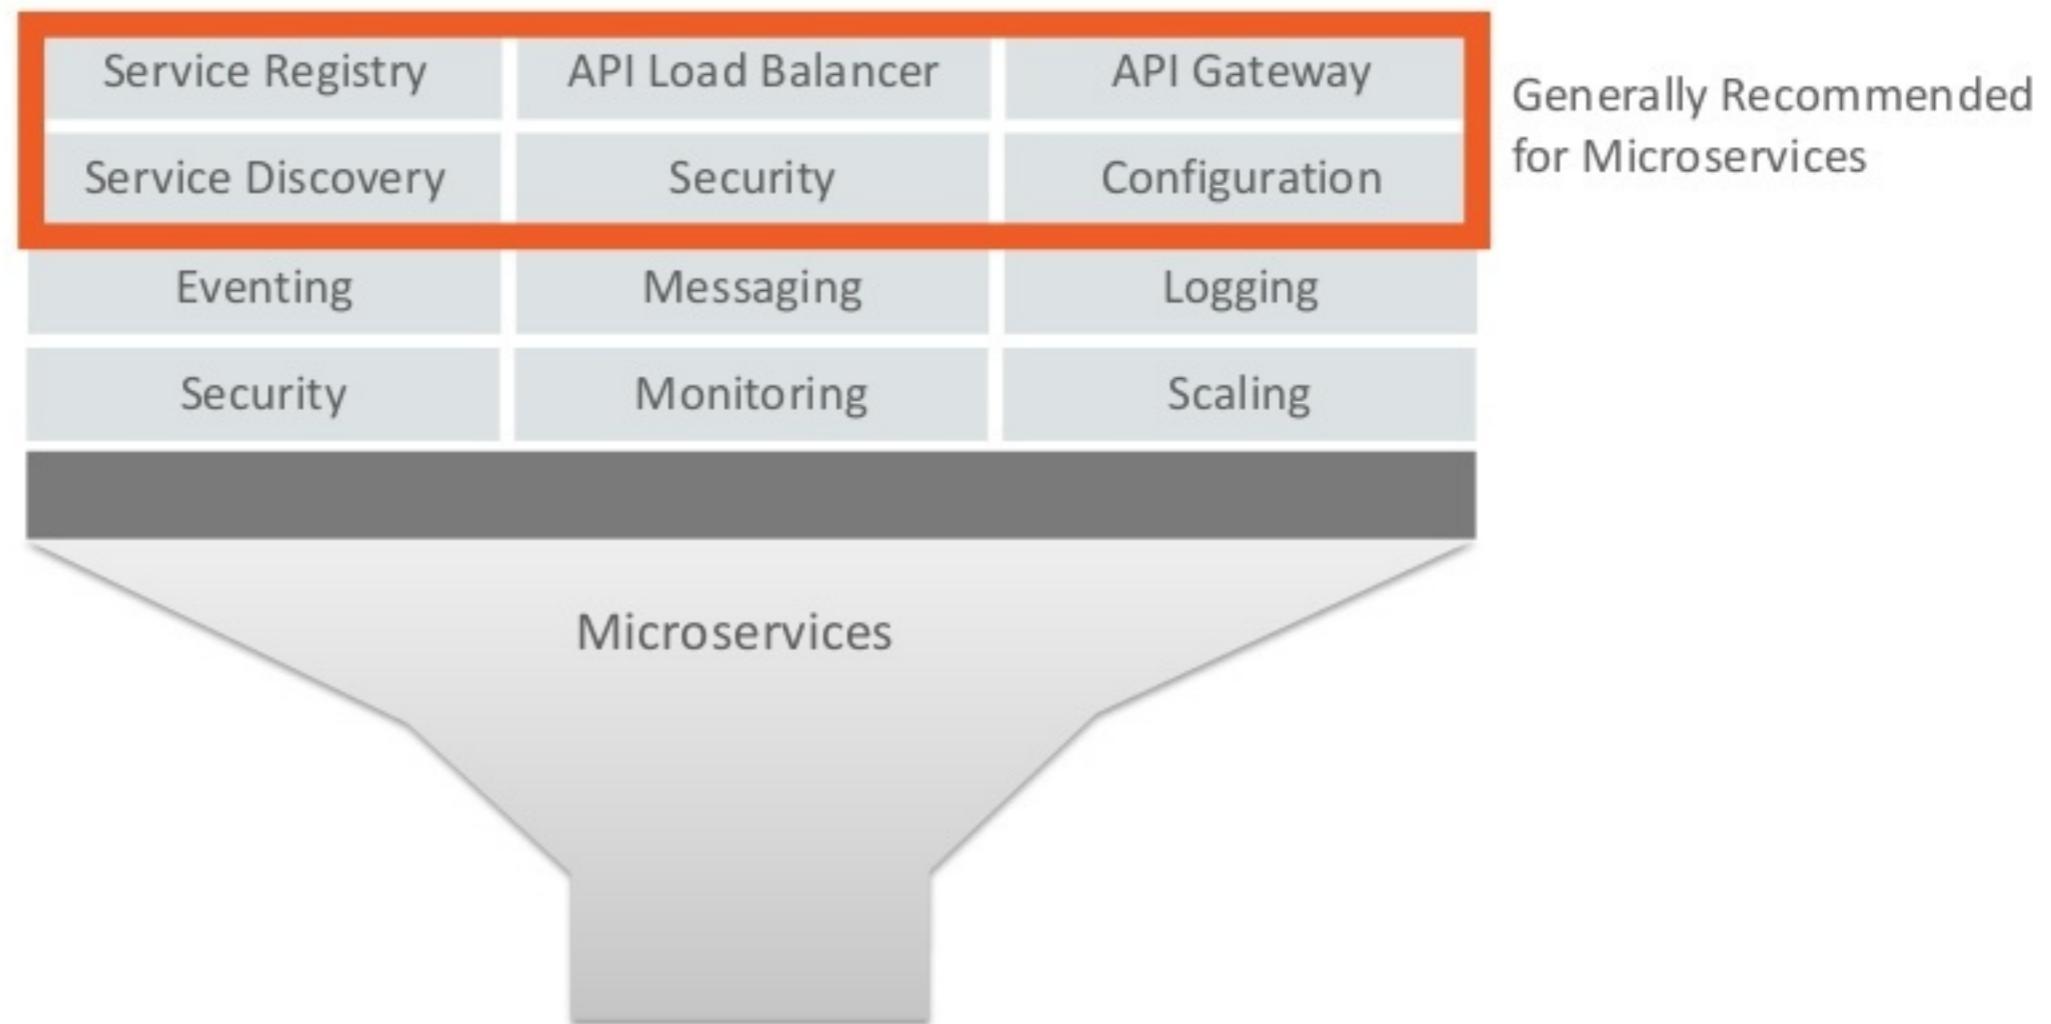
Send update to Product Catalog microservice but v1.2 adds a new required property

```
<product>
  <id>329340224</id>
  <shipToStore>true</shipToStore>
</product>
```

Technical Prerequisites

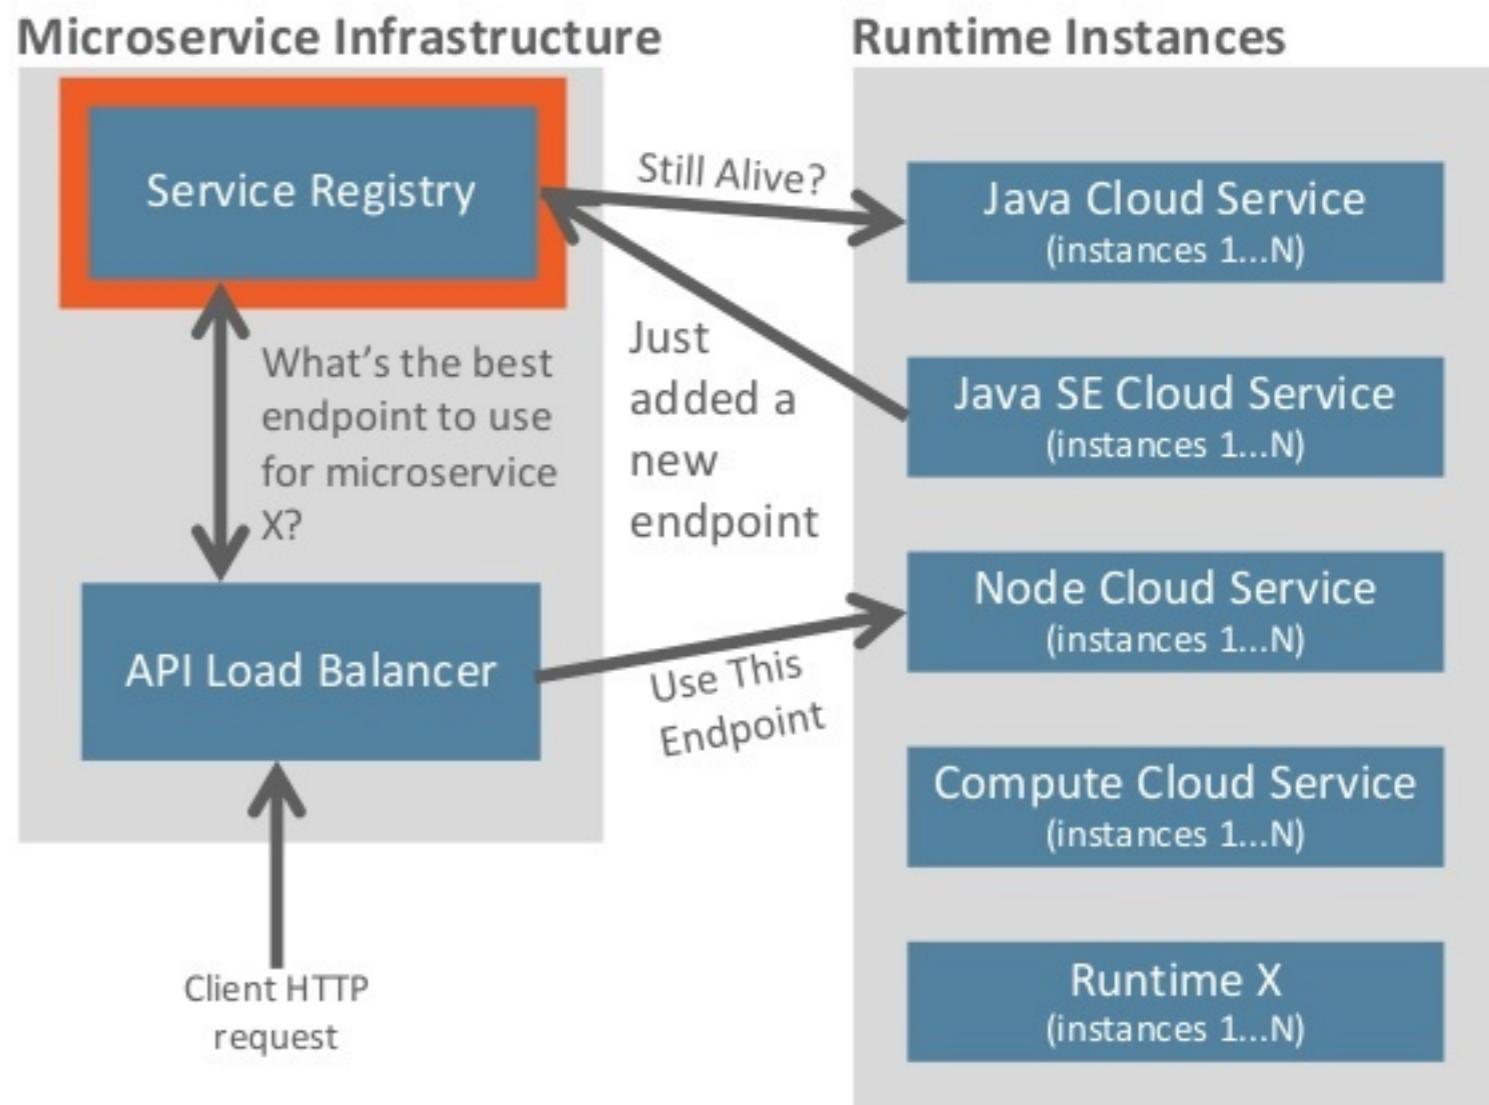


Requirements for Microservices Implementation



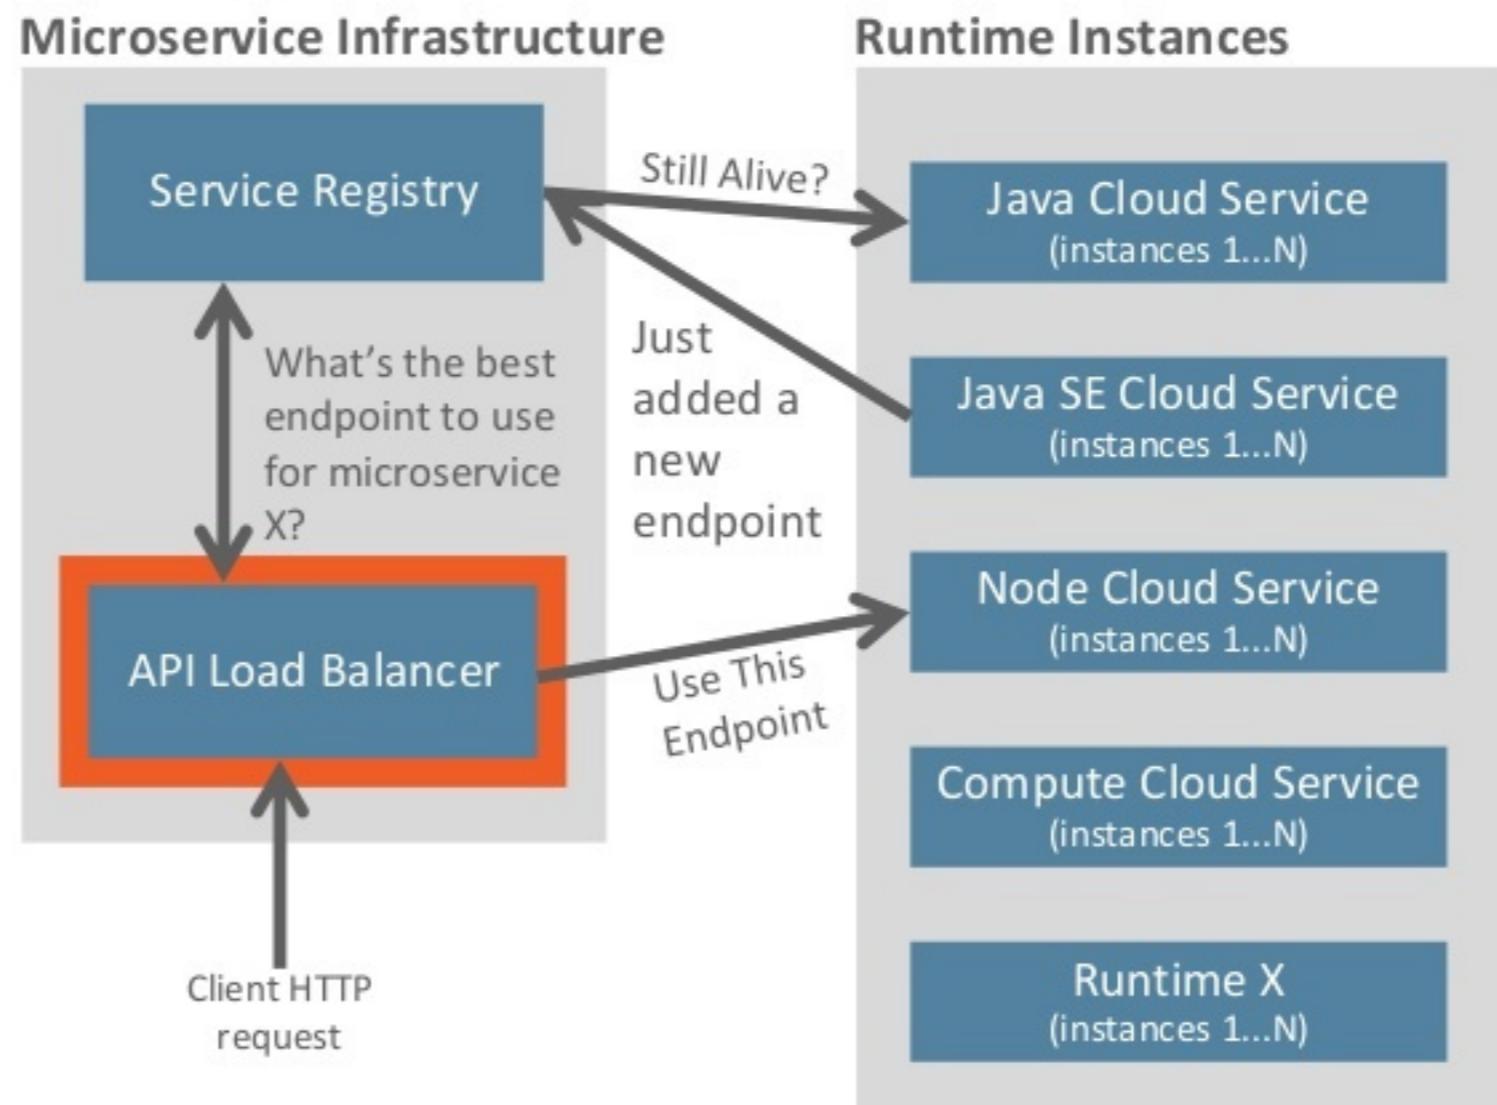
Key Microservices Technology – Service Registry

- Manages the lifecycle of each microservice endpoint
 - Newly-instantiated endpoints register with Service Registry
 - Service Registry continually polls each endpoint's health
- Aware of tenants, microservice versions, and environments
 - Health checking and selecting the most appropriate endpoint are very much dependent upon the tenant, version, and environment
 - Can query for an endpoint based on those attributes



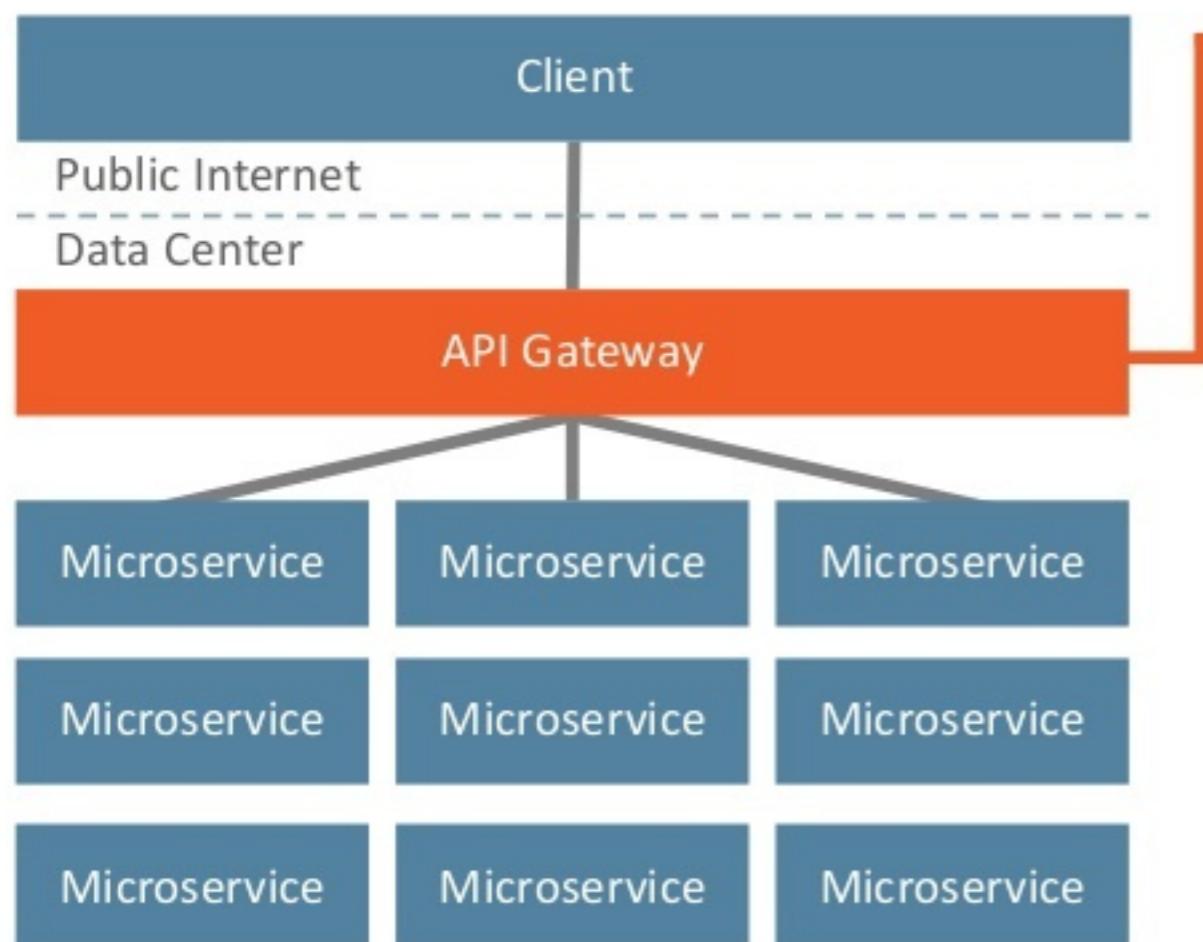
Key Microservices Technology – API Load Balancer

- Matches each request to the best endpoint
- Queries Service Registry to determine the best endpoint
- Ideally stateless – looks up the best endpoint for each HTTP request
- Uses load balancer or web server as the core. Typical implementations are customized Oracle Traffic Director, Nginx, or Apache mod_rewrite
- Handles logging, request rate limiting, authentication



API Gateways Load Balance and Aggregate Responses

API gateways provide a "backend for each frontend"



- Builds a XML or JSON response for each type of client – web, mobile, etc
- Asynchronously calls each of the N microservices required to build a response
- Handles security and hides back-end
- Load balances
- Applies *limited* business logic
- Meters APIs
- Logs centrally
- Common solutions: Netty, Vertex, Nginx, Kong, Apigee

API Load Balancer vs. API Gateway

API Load Balancer

- ✓ Matches each request to the best endpoint
- ✓ Queries Service Registry to determine the best endpoint
- ✓ Handles logging, request rate limiting, authentication

API Load Balancers are missing aggregation
but...

It's best to build a microservice that does aggregation. Keep the API Load Balancer or Gateway as free of business logic as possible

API Gateway

- ✓ Matches each request to the best endpoint
- ✓ Queries Service Registry to determine the best endpoint
- ✓ Handles logging, request rate limiting, authentication
- ✓ Aggregates the responses of many microservices

Service Discovery – Generally Not a Problem With Monoliths

Requirements

- Must resolve:
 - Host/port
 - Version for each microservice
- Must be discoverable

Because

- Many microservices in an environment
- Many environments
- Many versions of each microservice
- Hosts/ports can change quickly
- It's not practical to manage by hand

Solutions Include

- Any DNS provider
- Consul
- Zookeeper
- Custom coding to search instance tags

Approaches Include

- Plain DNS - limited to IP only
- SRV DNS records - gives IP + port
- Hierarchical namespaces
- Instance tagging in the cloud

Microservices Requires Robust Messaging

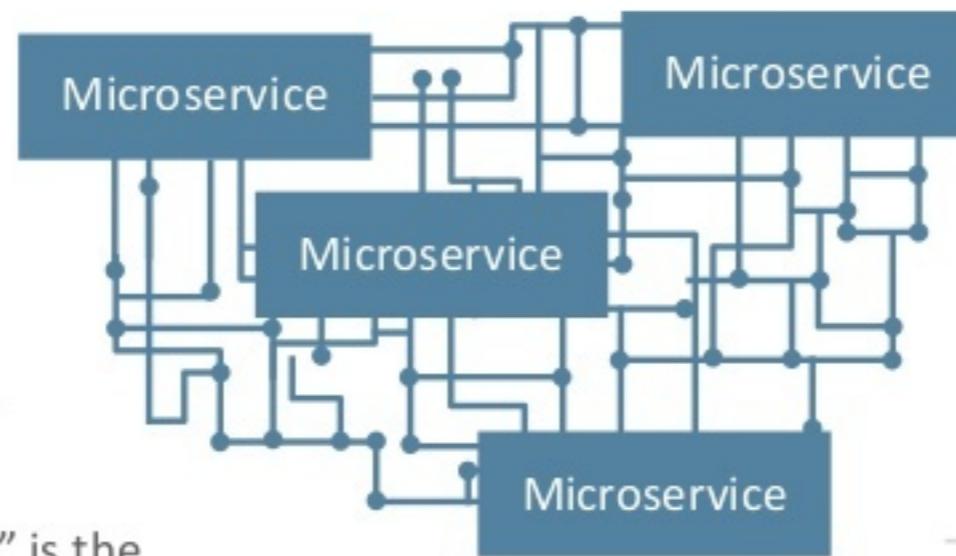
Both traditional durable messaging and non-durable eventing

Why Use Messaging?

- Message broker can buffer messages until the consumer is able to process them – prevents synchronous coupling which leads to outages

Requirements for Messaging

- “Smart endpoints, dumb pipes” is the philosophy of microservices – messaging should just pass messages. Not manipulate them
- Should support a variety of communication patterns including one-way requests and publish-subscribe



Types of Messaging

Normal Messaging

- Durable, ordered
- Relatively low throughput
- Usually brokered
- Often used to keep the data across different microservices in sync, as each microservice has its own data store

Eventing

- Non-durable, un-ordered
- Very high throughput
- Usually non-brokered
- Often used to distribute notification events – scale up, scale down, etc

Circuit Breakers Prevent Cascading Failures

Cascading failures are more common with microservices

- Rule #1 of microservices – avoid coupling!
 - Synchronous = two systems are coupled
 - Asynchronous = no coupling
- Cascading failures happen when request-handling threads are waiting on a response from a remote system
- Circuit breakers make synchronous calls from another thread pool to avoid binding up request-handling threads
- Hystrix (Java-based) is well-known and solves this problem



Emit Logs as Event Streams

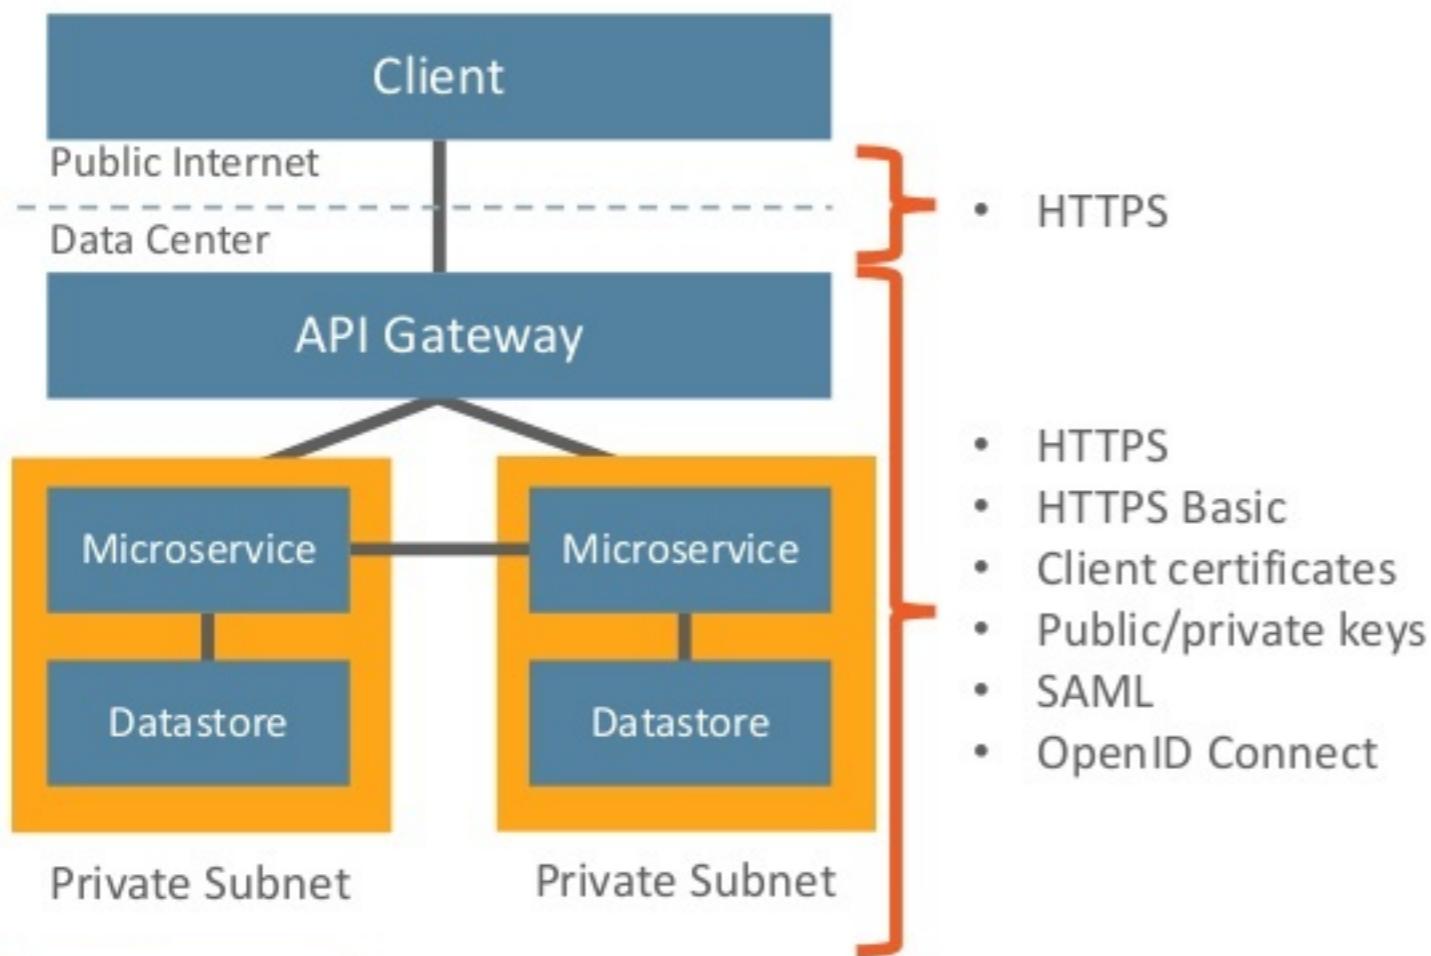
Can't do anything with log files sitting on a container's local storage volume

Capture, Aggregate, Search, Troubleshoot, etc



Security: Requires New Paradigms to Properly Secure

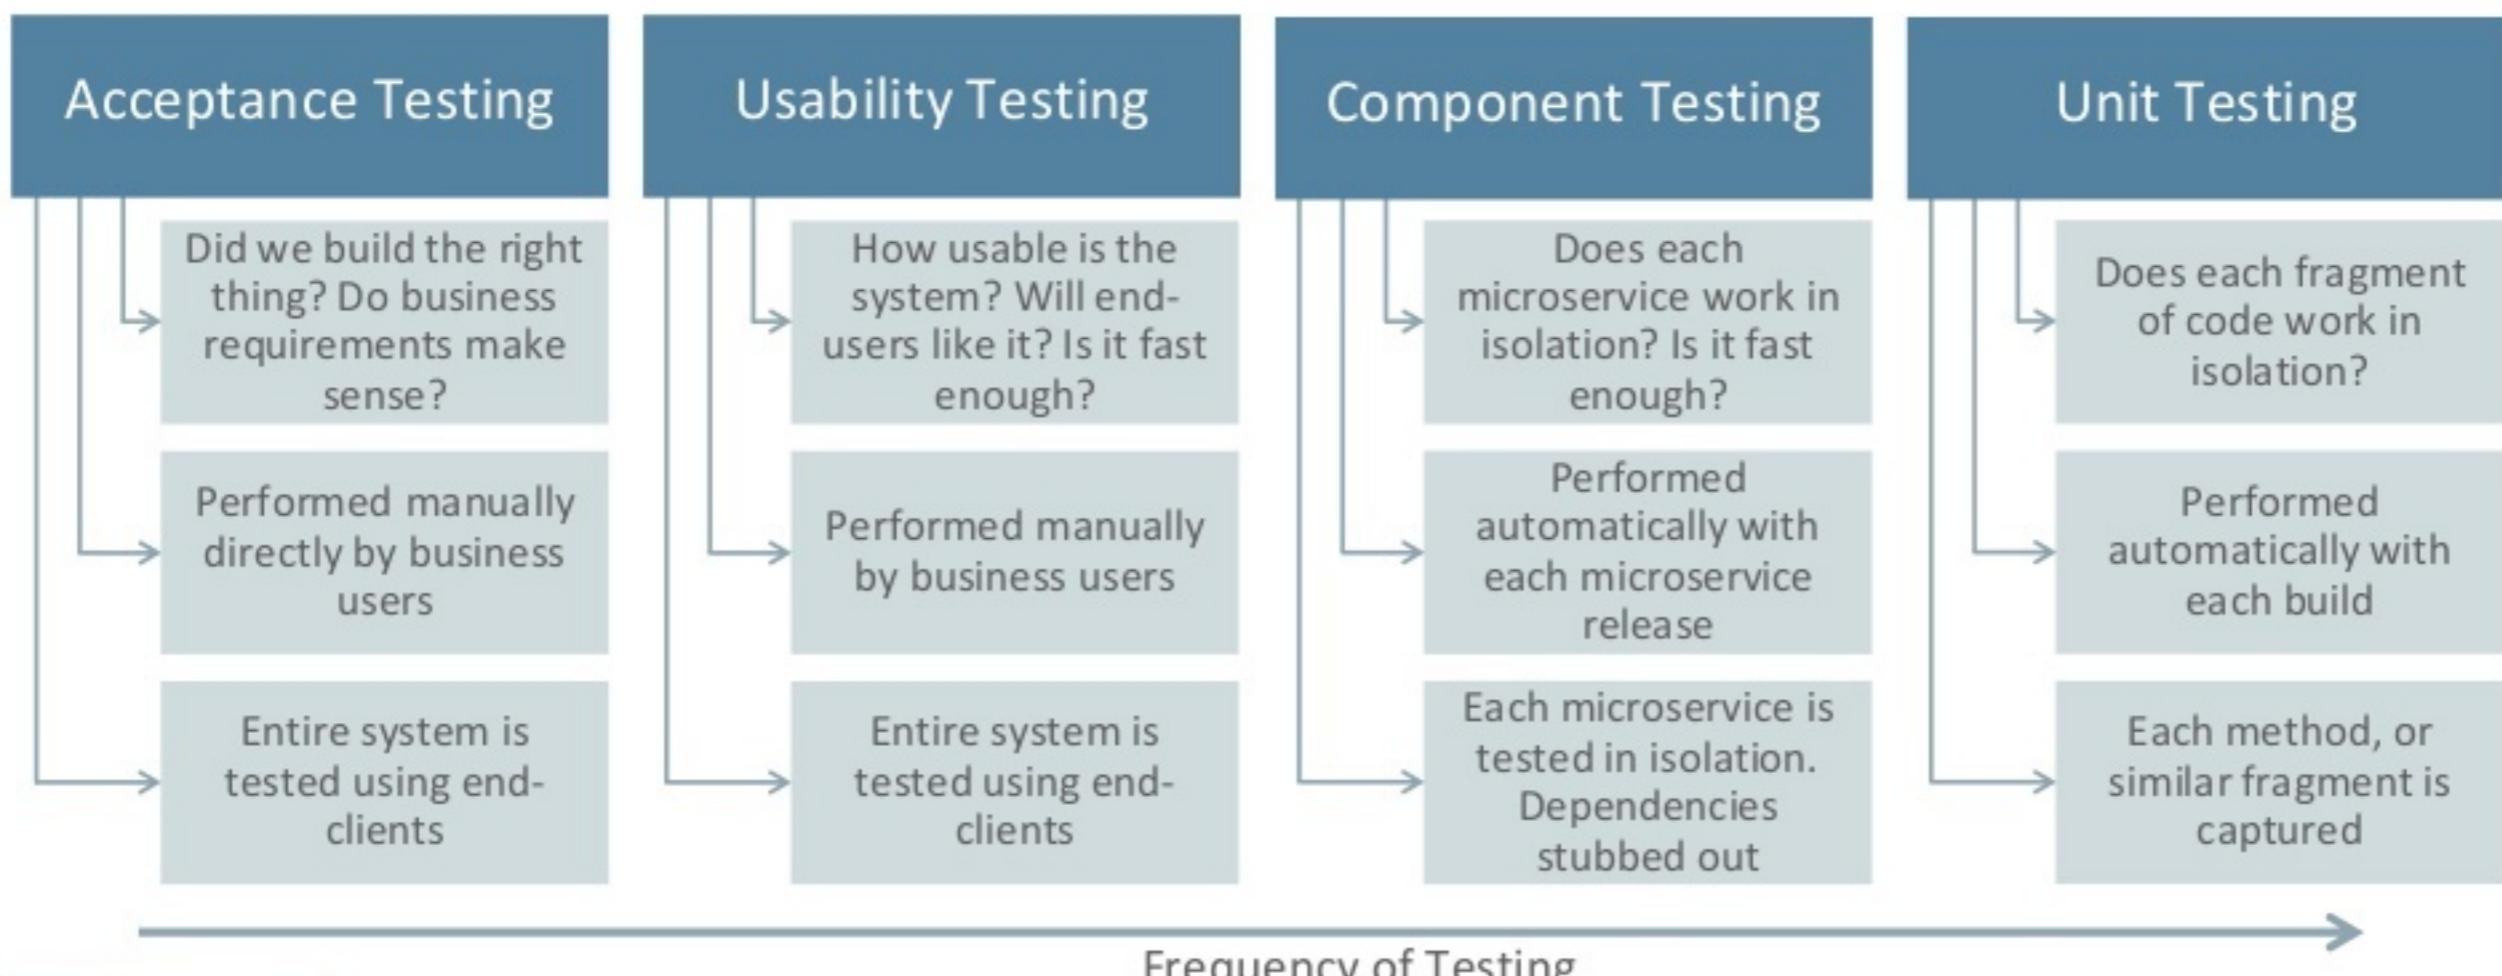
More microservices, more principals, more technologies, more everything



- Must *authorize* and *authenticate* every single principal – best to use common approach
- Monolithic apps typically do authentication and authorization on their own
- Secure every single remote network call
- Use network segmentation to isolate microservices from each other

Microservices Complicates Testing

More code, more microservices to test



What/How To Monitor

Monitoring a monolith is relatively easy – one app. Microservices = many apps



Requirements for Monitoring Microservices

1. Monitor throughput, performance, and business metrics
2. Trace each end-request through every microservice – end-to-end
3. Track the health of downstream dependencies
4. Monitor each process, OS, host, etc

Popular Tooling

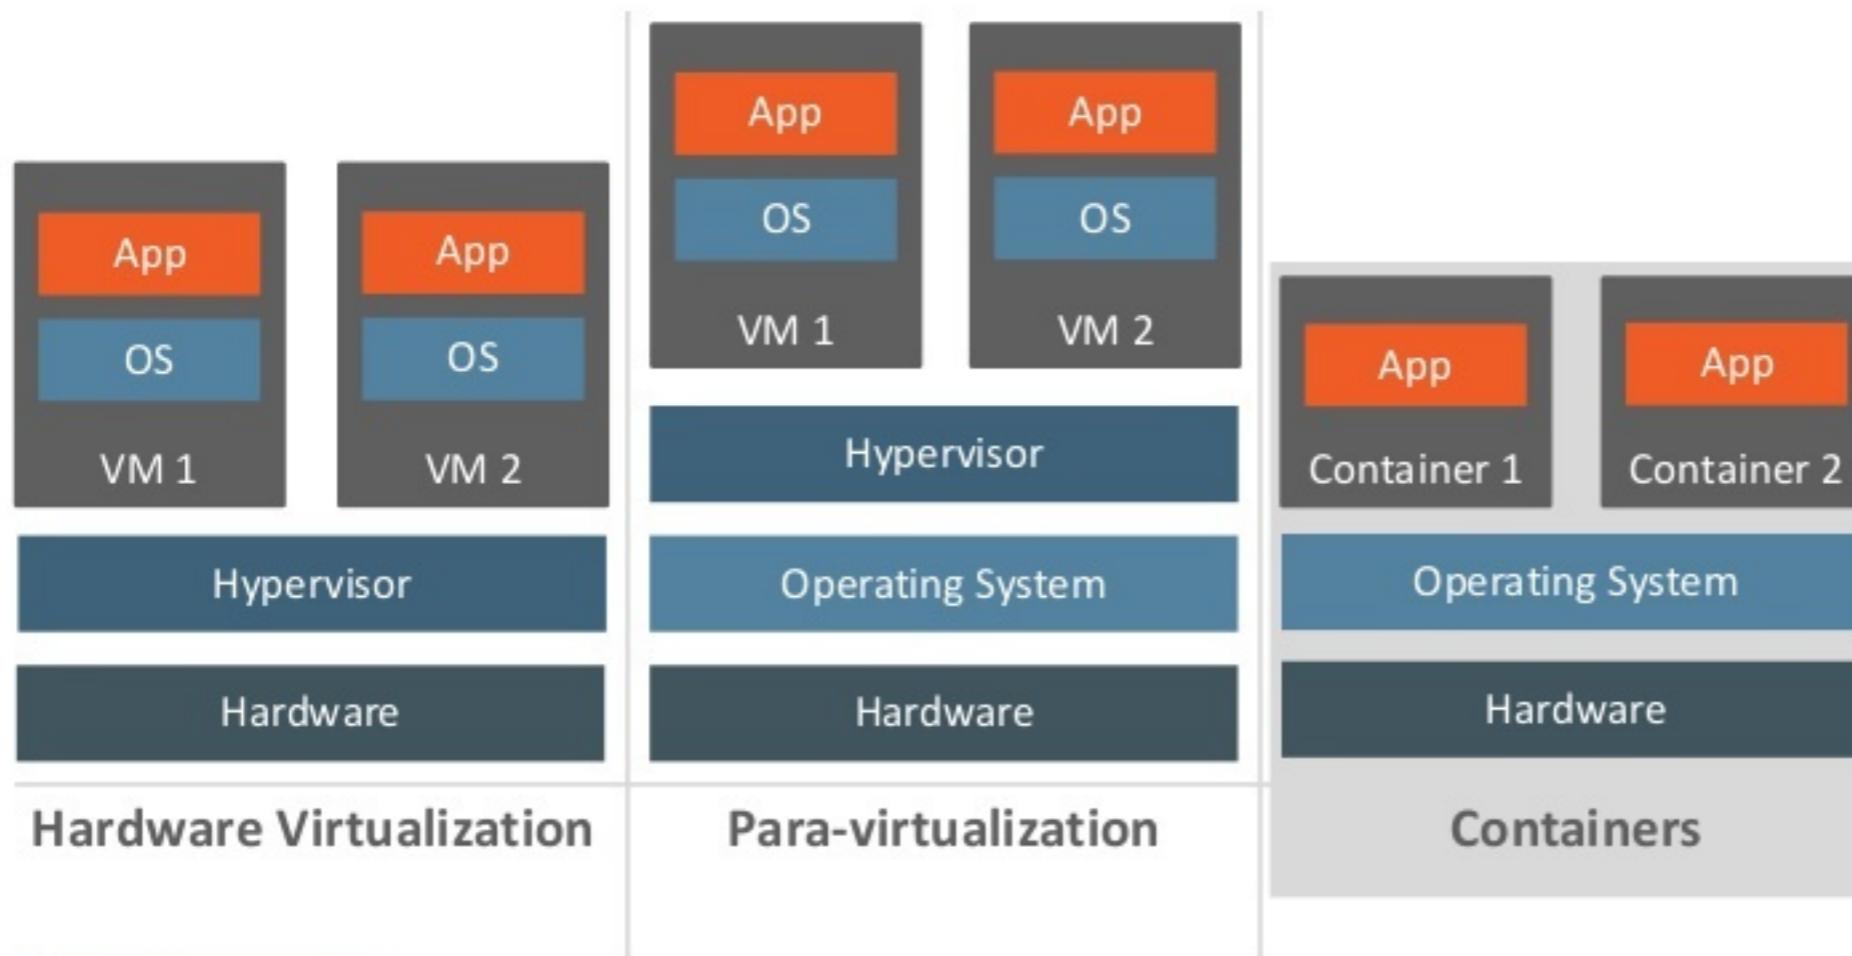


Dropwizard Metrics



Containers Make Microservices Easier

Helpful to microservices but not a requirement



- #1 value – app packaging
- Microservices doesn't rely on containers but they do help:
 - Higher density
 - Easy to start/stop
 - Portability
- Containers are lightweight, just like microservices themselves

Many See Containers As the Standard

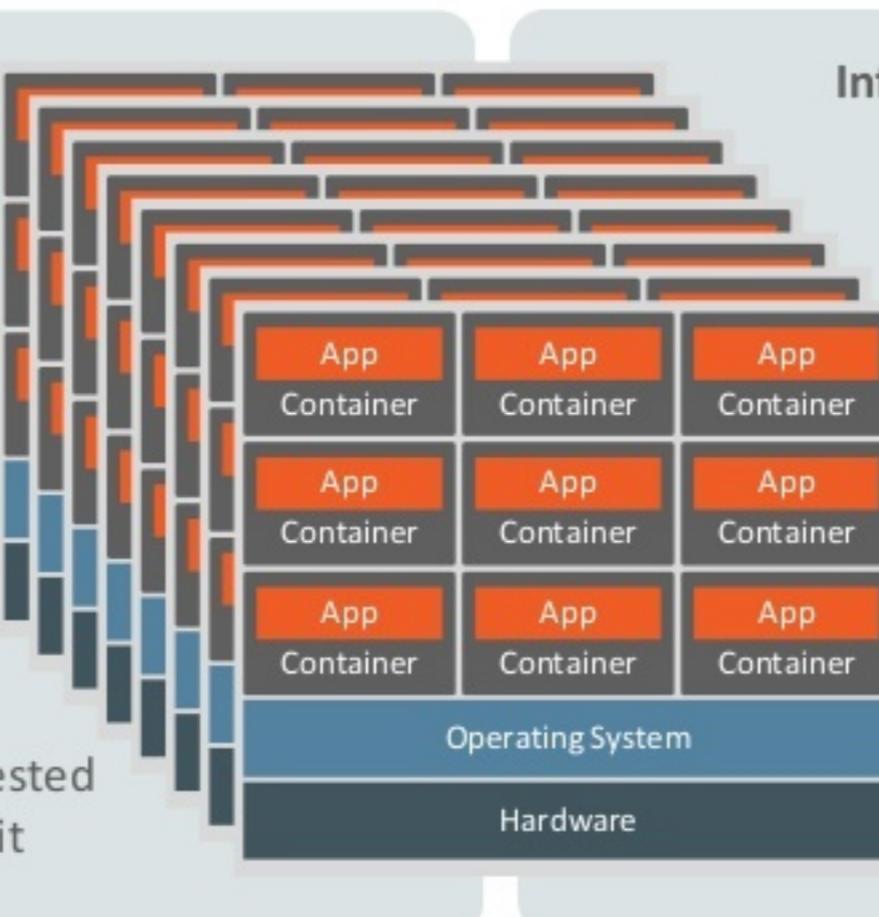
Four main use cases

Application Packaging

Neatly package applications and supporting environment in immutable, portable containers

Continuous Integration

All changes to an app are contained in one immutable container image. Container is tested and deployed as one atomic unit



Infrastructure Consolidation

Get infrastructure utilization up to 100% (vs 5-10% with VMs) due to over-subscription of resources and near bare metal performance.

DIY PaaS

Build a simple PaaS by wiring up containers to a load balancer. New code, patches, etc pushed as new immutable containers.

Containers Should Be Immutable



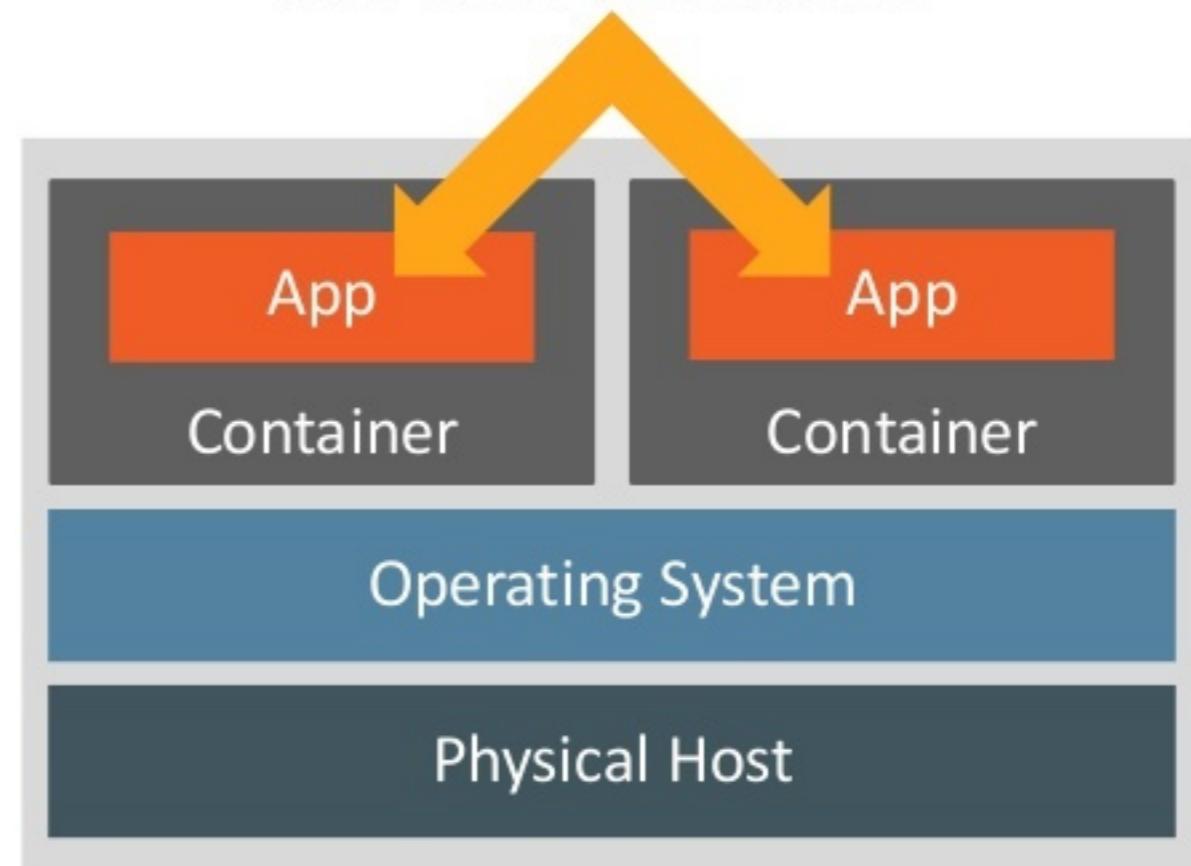
Build and deploy a new container

Never touch a container that's already been built

One Instance Per Container is Typical

- Best to run one instance (unique host/port combination) per container
- Running multiple instances of the same application or different applications will make scheduling very difficult
- Expose one port per container

Just One Per Container



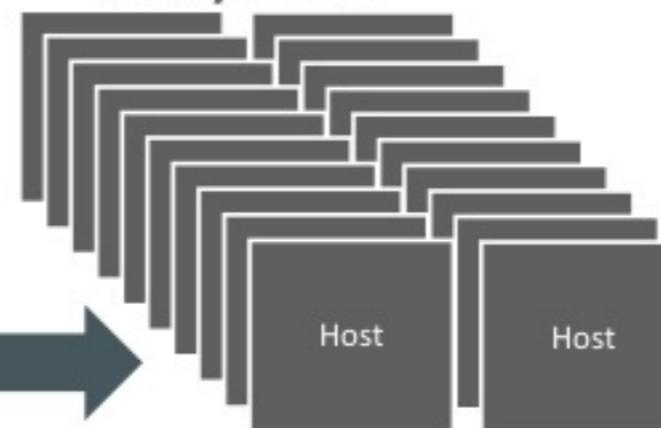
How Do You Deploy Containers to Physical Hosts?

The emerging space of container orchestration

Many Containers



Many Hosts



Docker Swarm



MESOS



RANCHER



kubernetes

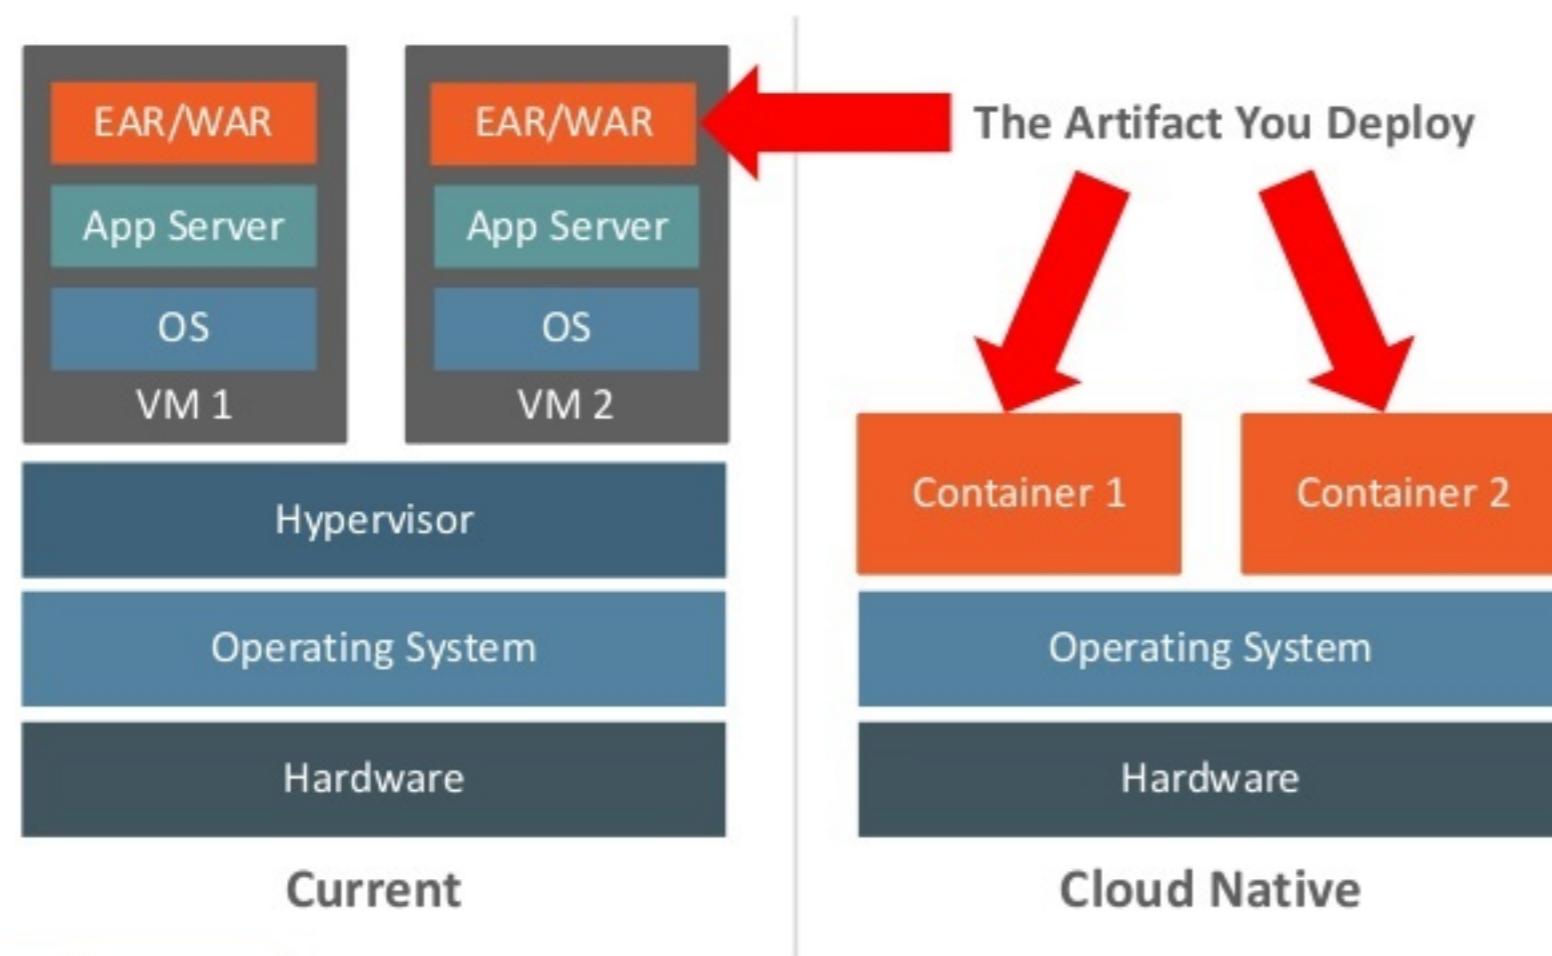
Emerging space. Solutions are very early and lack any real notion of an application. Still very much infrastructure-focused

What Do Container Orchestration Solutions Do?

- Map containers back to physical hosts, taking into account user-defined placement rules, the utilization of each host, and the needs of each container. Can be very complex
- Set up overlay networking, firewalls, ensure network QoS, etc
- Auto-scaling
- Local and external load balancers
- Service registry / discovery

Artifacts Are Now Immutable Containers – Not EARs, WARs

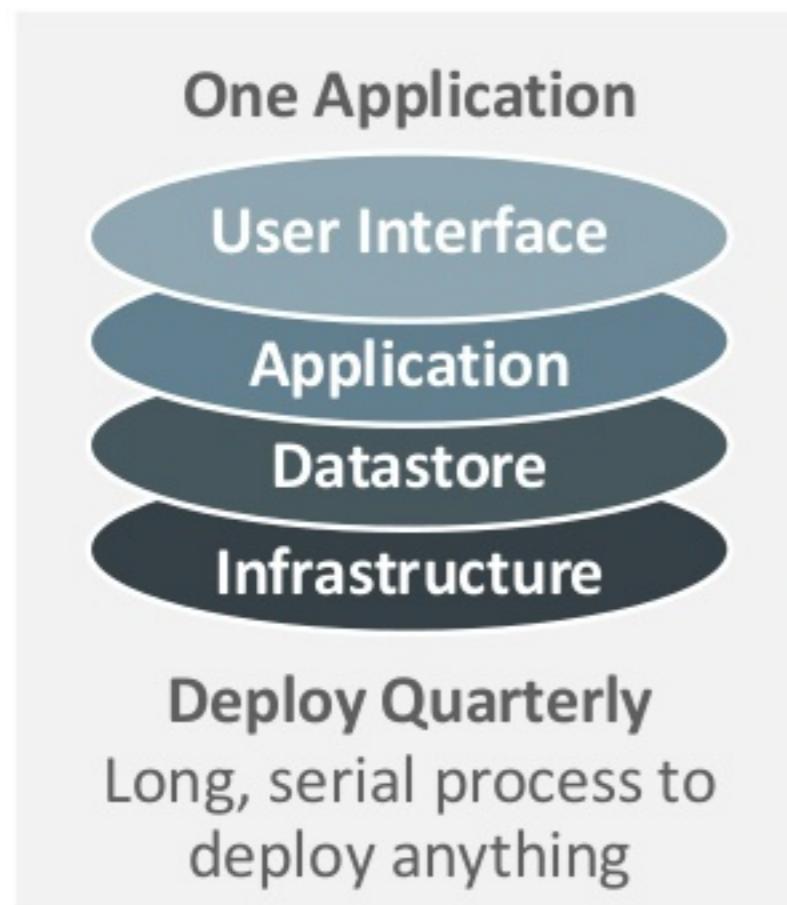
Containers can have anything in them and are highly portable



- No more installing a JVM, app server, and then deploying the artifacts to them
- Build the container once, deploy it anywhere. Can include complex environment variables, scripts, etc
- Containers should be free of state and configuration
- Containers should not assume they are able to write to a persistent local file system

Constantly Deploy Each Microservice

No need to batch together many changes to single app



Run Many Versions of the Same Microservice Concurrently



Run only *one* version of the same application in the same environment

Run *many* versions of each microservice in the same environment

Application Containers – Use What You Have Today

Characteristics of Existing Application Server

- Run a very large, complicated monolithic application
- Many advanced features
- Integrated dependency management
- Few instances per host
- Stateful

Characteristics of Microservices-friendly Application Server

- Run a smaller, microservice that does one thing really well
- Few features
- External dependency management
- Many instances per host
- Stateless

Run whatever works best – few firm requirements



How Oracle Products Support Microservices

Oracle's Microservices Strategy

Oracle's Microservices Roadmap

Today With Oracle Products

| | | |
|------------------------------------|---|--|
| State | Planned - Oracle Cloud State Service | Oracle Coherence or Oracle WebLogic |
| Configuration | Planned - Oracle Cloud Config Service | Oracle Coherence |
| Runtime | Planned - Runtime - Jersey + Grizzly | Oracle WebLogic, Node, Java SE, etc |
| Eventing | Planned - Oracle Cloud Eventing Service | Oracle Coherence |
| Messaging | Oracle Messaging Cloud Service | Oracle Messaging Cloud Service |
| Management/Logging/Alerting | Oracle Management Cloud Service | Oracle Management Cloud Service |
| Datastore | Oracle Database or NoSQL Cloud Service | Oracle Database or NoSQL Cloud Service |
| Central Source of Truth | Planned - Oracle Microservices Platform | Oracle Coherence |
| Service Discovery | Planned - Oracle Microservices Platform | Oracle App Container Cloud Service |
| API Gateway/Load Balancer | Planned - Oracle Microservices Platform | Oracle App Container Cloud Service |
| Container Orchestration | Planned - Oracle Microservices Platform | Oracle App Container Cloud Service |
| Infrastructure | Oracle Cloud | Oracle Cloud |
| Build/Deploy | Oracle Developer Cloud Service | Oracle Developer Cloud Service |



Coming Soon!

An All-New Microservices Platform From Oracle

All New



Light/fast
Jersey-based
Easy Packaging



ORACLE
Microservices
Platform

Rolling Upgrades
Stateless Support
Stateful Support

Rollbacks
Service Discovery
Blue/Green Releases

Canary Testing
Health Monitoring
Language-specific Tooling

Infrastructure Agnostic
Placement Constraints

High Availability
High Density

Hybrid Cloud
Internal Load Balancing

Consume Platform as PaaS on Oracle Cloud

Install Platform On Any IaaS, On Or Off Premise (future)



Oracle Public Cloud



Oracle Public Cloud
Machine



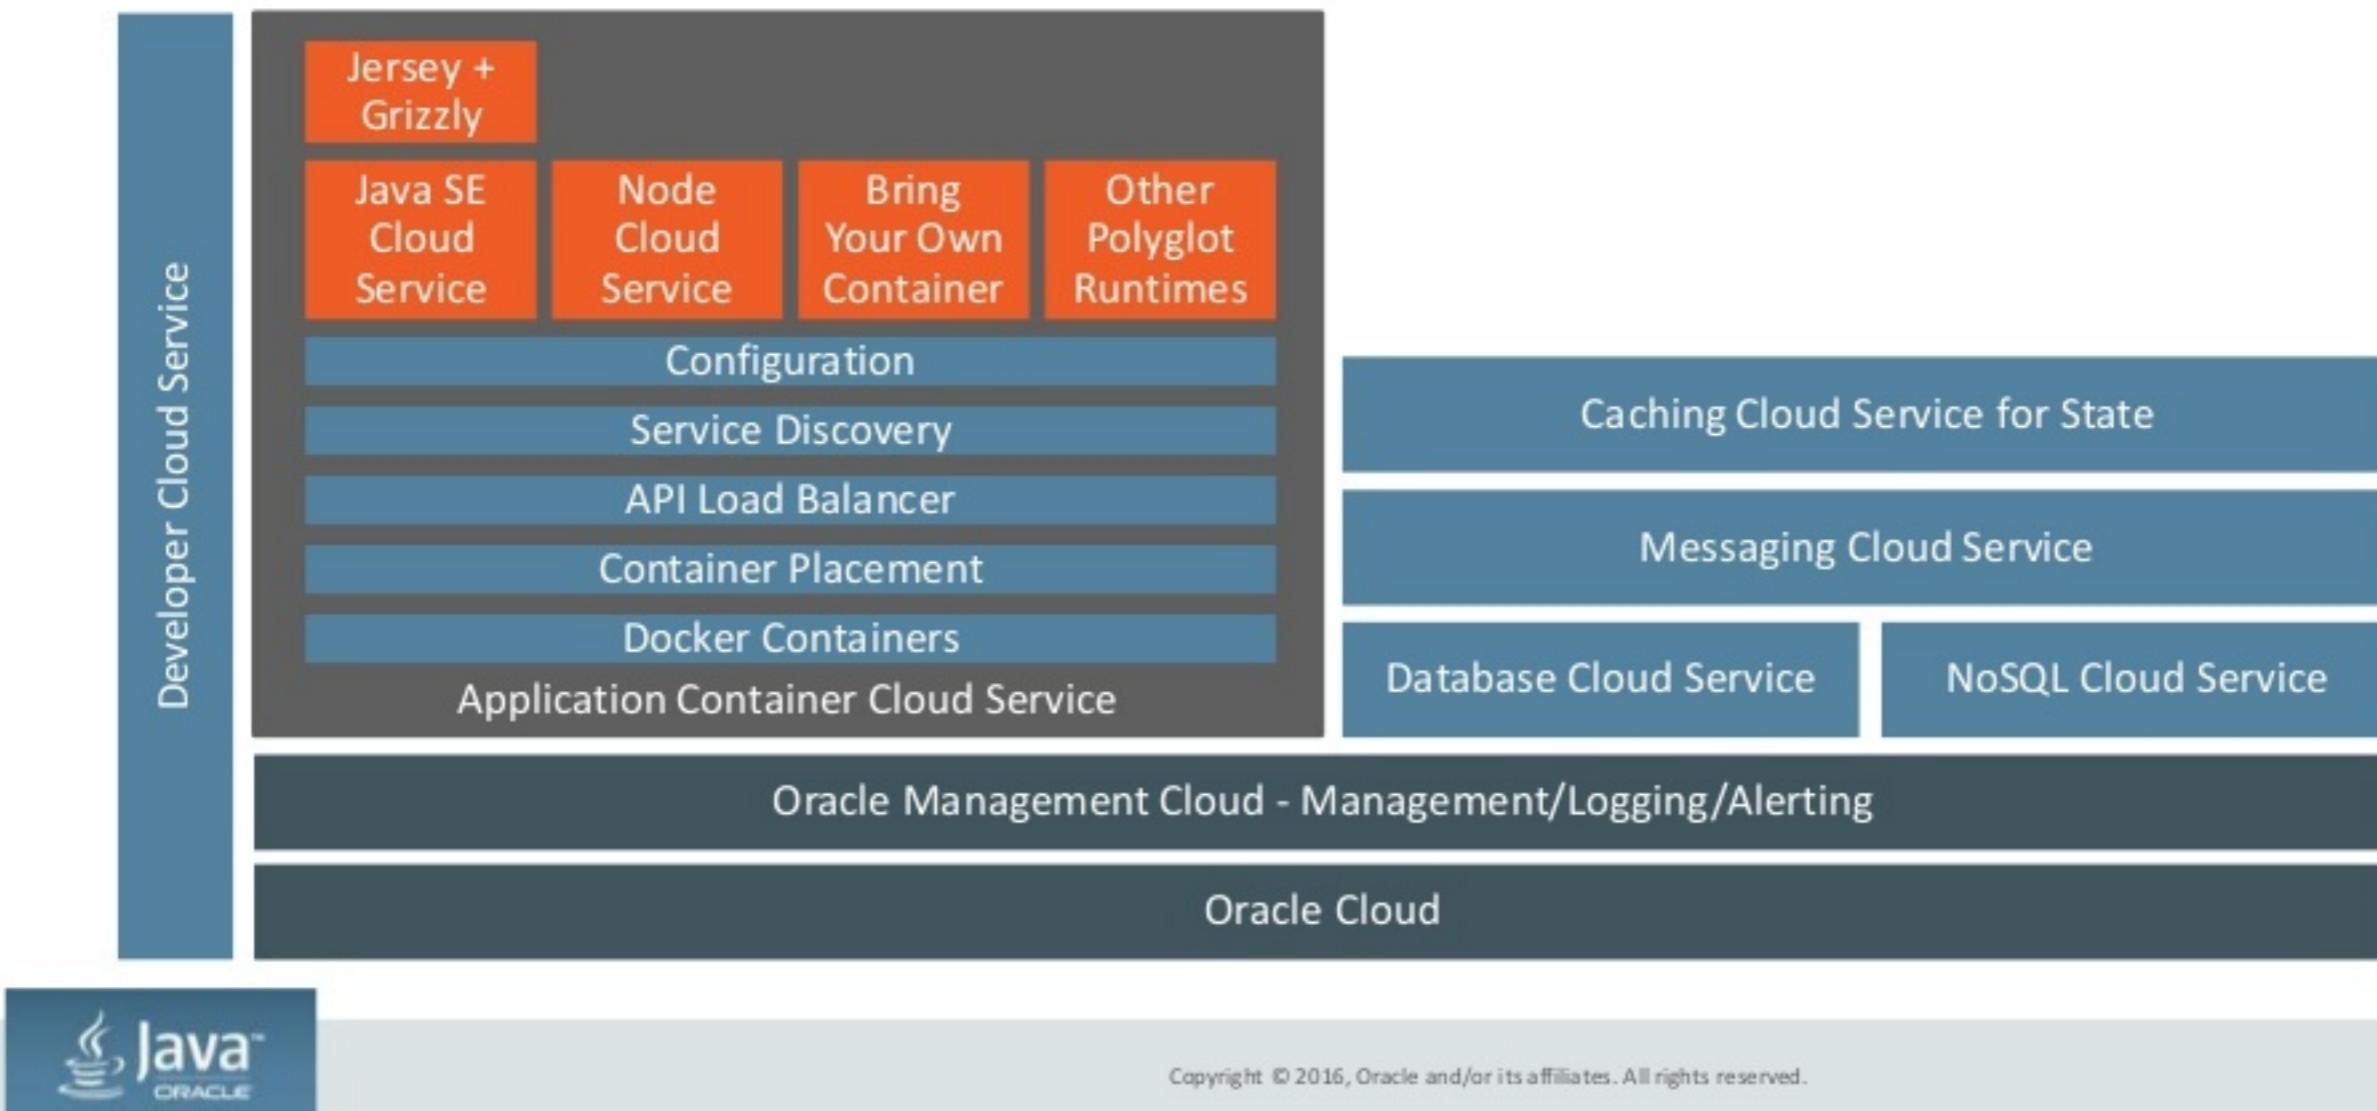
- Config
- State
- Eventing
- Messaging
- Logging
- Monitoring
- Database
- NoSQL
- Caching
- Integration
- Big Data
- Mobile
- Process
- Developer

Services Available
for Consumption
Inside the Platform



Using Application Container Cloud Service For Microservices

A modern platform for lightweight application development



Introducing Oracle App Container Cloud Service



Key Features

- Useful for any Java, Node.js, or polyglot runtime
- IDE Choice - JDeveloper, Eclipse, NetBeans - and API access
- Continuous integration with Oracle Developer Cloud Service
- Cloud tooling for lifecycle management
- Integrated load balancer, support for service discovery
- Bring-your-own-container model supported shortly

Benefits

- Self-service application platform with advanced cloud tools
- Secure, Highly Available with Clustering
- Fully automated provisioning, patching, backup, and recovery

JAX-RS: The Java API for RESTful Web Services

Oracle is a spec lead

Server-side Code

```
@Path("/atm/{cardId}")
public class AtmService {
    @GET @Path("/balance")
    @Produces("text/plain")
    public String balance(
        @PathParam("cardId") String card,
        @QueryParam("pin") String pin) {
```

Client-side Code

```
Client client = ClientFactory.newClient();
String balance =
    client.target("http://xxxx/atm/{cardId}/
                  balance")
    .pathParam("cardId", "1234567890123456")
    .queryParam("pin", "1111")
    .request("text/plain")
    .get(String.class);
```

Simply annotate Java code to expose as REST

Originally defined in JSR 311 - 1.0
Part of Java EE 6 Spec

Use REST without having to parse text

Updated as JSR 339 in 2013 - 2.0
Part of Java EE 7 Spec



Jersey: The Reference Implementation of JAX-RS

Oracle continues to lead the spec



JERSEY

- Oracle sponsored open source
- Implements the JSR 311 specification
- Contains
 - Standalone server
 - Client
 - JAXB/JSON support

Grizzly: High Performance I/O

Great for inter-process communication

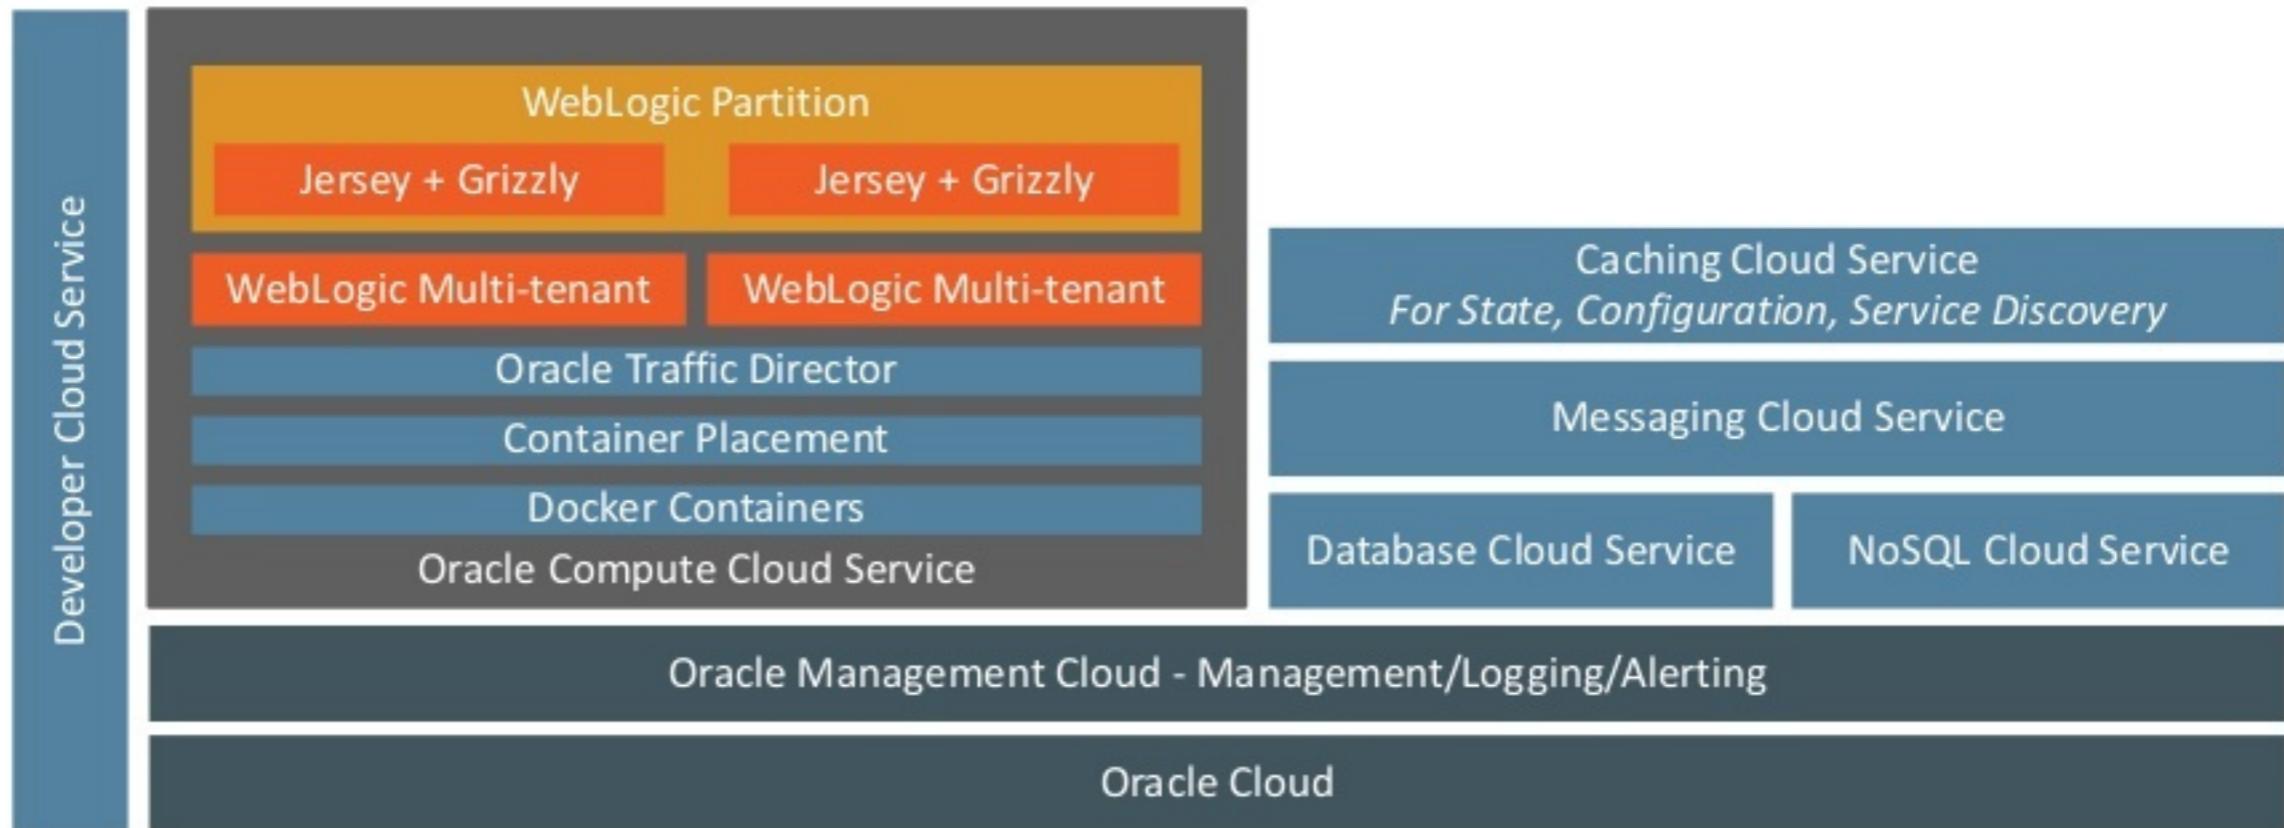
- Oracle sponsored open source
- Allows developers to take advantage of the Java NIO to provide very fast inter-process communication
- Brings non-blocking sockets to the protocol processing layer
 - Support for non-blocking HTTP processing
- WebSocket Support
- APIs make non-blocking interactions simple



Project **Grizzly**

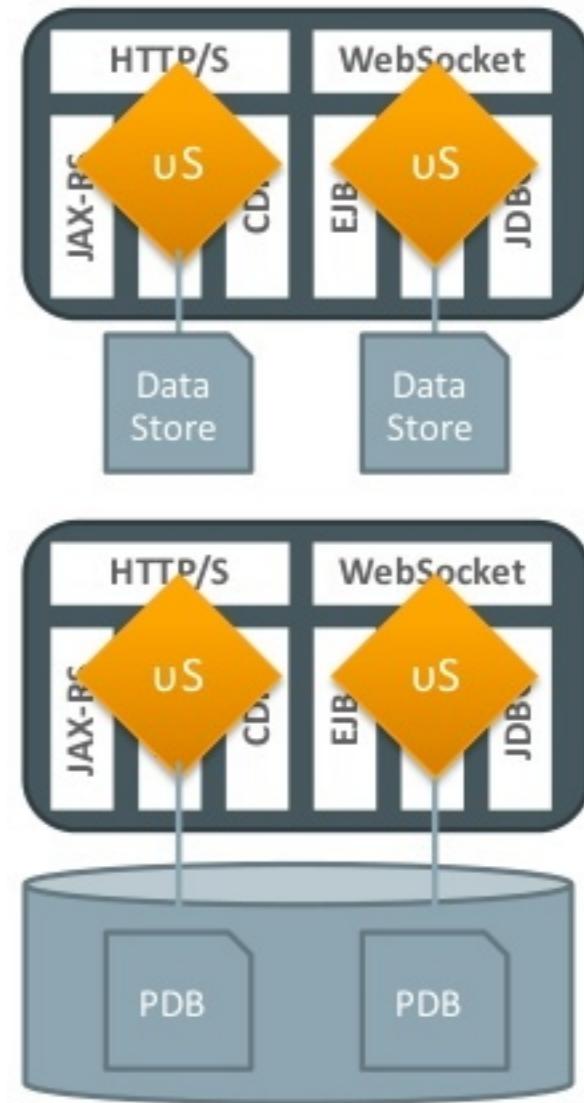
Using WebLogic + Java EE For Microservices

A proven, enterprise-grade platform for application development



Java EE Platform Supports Microservices

- Standards based infrastructure
 - Prescribed, validated set of APIs and services
 - WAR file deployment for service boundary
- Enables development of thin Microservices
 - Assemble application code and resources only
 - APIs and implementation libraries supplied by platform
- Freedom to choose best service topology
 - One service per server, multiple services per server
 - One server per container/host
 - Multiple containers/servers per host



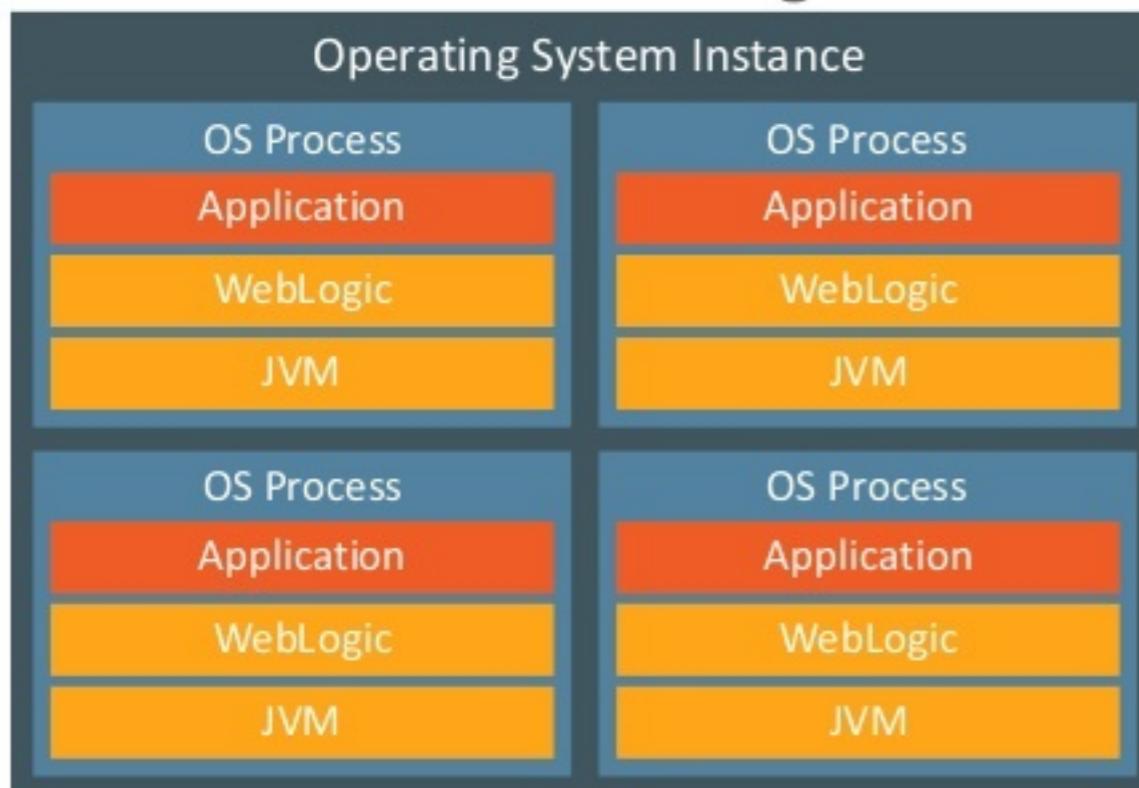
WebLogic Supports Microservices (and then some!)

- Java EE Platform Support
 - With latest Java language support
- Fully automatable infrastructure
 - Scriptable provisioning, configuration, deployment supporting code-as-infrastructure
 - REST API, WLST, JMX
 - Callable from Maven, Gradle, Arquillian
 - Easily incorporated into Continuous Integration and Delivery workflows
- Flexible placement models
 - Single server, multiple servers, clusters of servers
- Diverse datastore options
 - Traditional multi-vendor relational database support
 - Integrated Oracle Database 12c pluggable database support as independent data stores for services
 - Simple configuration for Document/Graph databases
 - Container, application, embedded application scoped database resources

WebLogic Multi Tenant

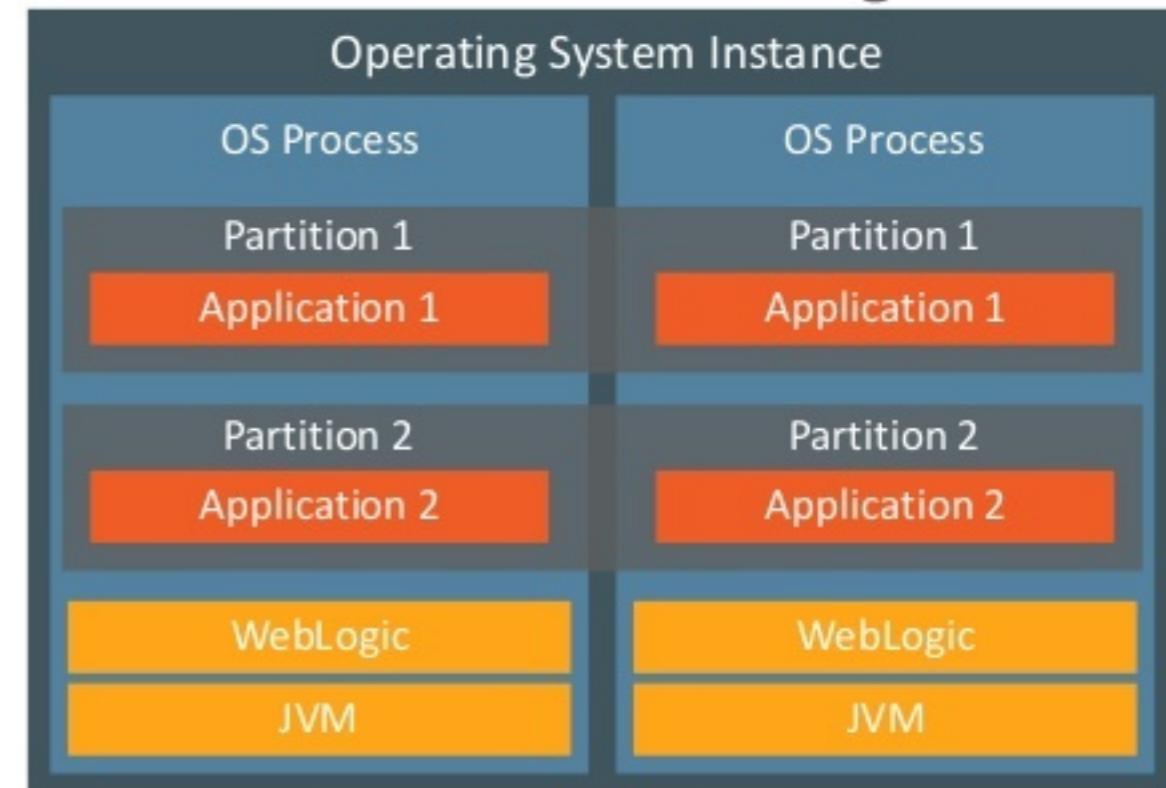
Partition is now a first-class object

Standard WebLogic



Single-tenant

Multi Tenant WebLogic

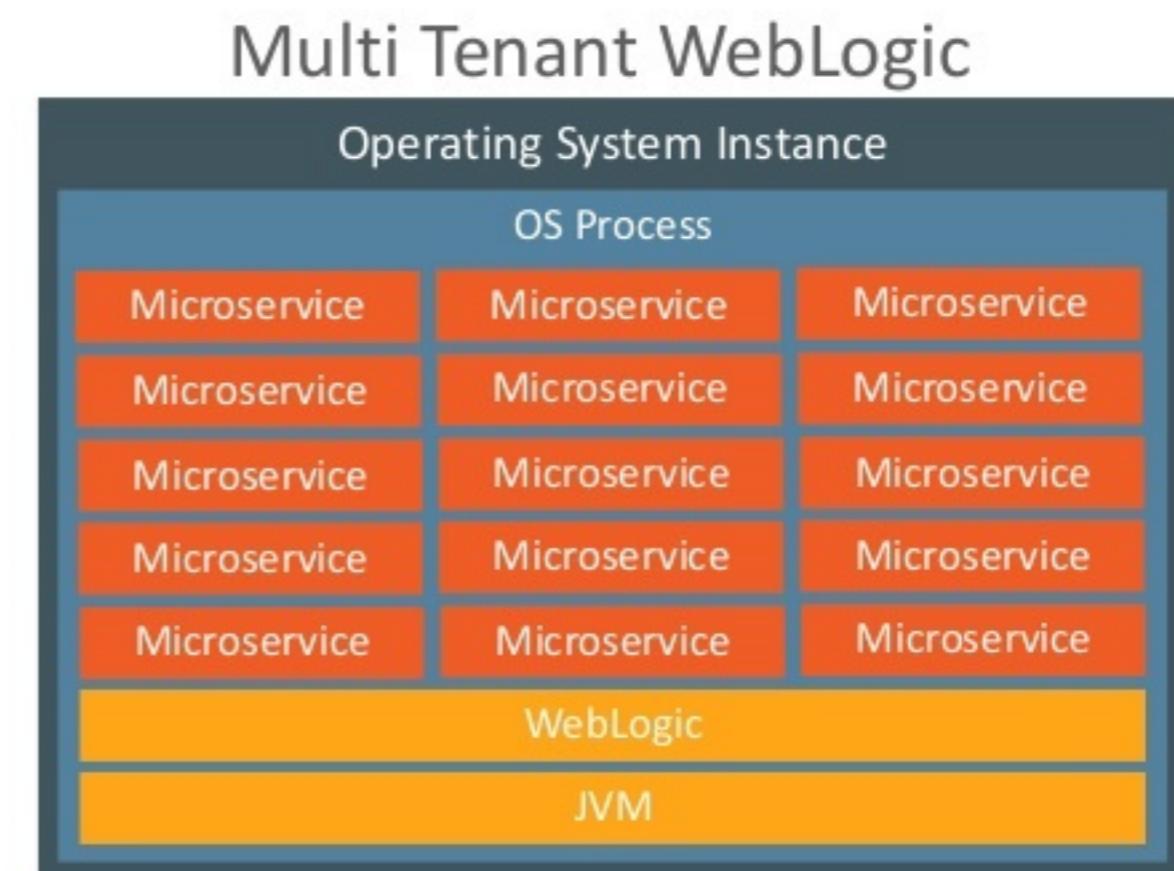


Multi-tenant – strong isolation between tenants

WebLogic Multi Tenant is Perfect for Microservices

Similar to Oracle Database pluggable/container databases

- Each microservice instance can have its own light-weight WebLogic container-like partition
- Easily move partitions between WebLogic hosts
- Each partition is exceptionally light
- Each WebLogic host can support hundreds of partitions



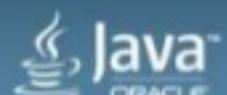


WebLogic Runs Great on Docker

```
sbutton@dhcp-au-adelaide-10-187-112-202] ~ $ docker images
REPOSITORY      TAG        IMAGE ID      CREATED       VIRTUAL SIZE
oracle/weblogic  12.1.3-dev  623062d74b4c  6 hours ago  2.94 GB
oraclelinux      7.0        5f1be1559ccf  4 months ago  265.2 MB
[...]
[1] sbutton@dhcp-au-adelaide-10-187-112-202] ~ $ docker run -d -p 17001:7001 oracle/weblogic:12.1.3-dev
867f223847507c01946f5f1f06797bccec71203c9e6ef0b29400137515478e928
[2] sbutton@dhcp-au-adelaide-10-187-112-202] ~ $ docker run -d -p 27001:7001 oracle/weblogic:12.1.3-dev
4a14cee634879629a0196528ae3cc7a775ee68430f1b70cc91d90e5dbf23b568
[3] sbutton@dhcp-au-adelaide-10-187-112-202] ~ $ docker run -d -p 37001:7001 oracle/weblogic:12.1.3-dev
00985092bd3ff3f17580dbe3b4db87246f8727423e5a7d1344ad109af7a338533
[...]
[1] sbutton@dhcp-au-adelaide-10-187-112-202] ~ $ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              NAMES
[...]
[2] 00985092bd3f    oracle/weblogic:12.1.3-dev   "/u01/oracle/weblogi"  28 seconds ago     Up 24 seconds   sick_babbage
4a14cee63487    oracle/weblogic:12.1.3-dev   "/u01/oracle/weblogi"  36 seconds ago     Up 33 seconds   angry_meitner
[3] 867f22384750    oracle/weblogic:12.1.3-dev   "/u01/oracle/weblogi"  44 seconds ago     Up 40 seconds   lonely_davinci
[...]
[1] sbutton@dhcp-au-adelaide-10-187-112-202] ~ $ docker top 00985092bd3f
PID      USER          COMMAND
1478      docker        {startWebLogic.s} /bin/sh /u01/oracle/weblogic/user_projects/domains/ba
se_domain/bin/startWebLogic.sh
1541      docker        /usr/java/jdk1.8.0_25/bin/java -server -Xms256m -Xmx512m -XX:MaxPermSiz
e=256m -Dweblogic.Name=AdminServer -Djava.security.policy=/u01/oracle/wls12130/wlserver/server/lib/weblogic.pol
icy -Dweblogic.ProductionModeEnabled=true -Djava.security.egd=file:/dev/.urandom -Dweblogic.rjvm.enableproto
colswitch=true -Djava.endorsed.dirs=/usr/java/jdk1.8.0_25/jre/lib/endorsed:/u01/oracle/wls12130/oracle_common/mod
ules/endorsed -da -Dwls.home=/u01/oracle/wls12130/wlserver/server -Dweblogic.home=/u01/oracle/wls12130/wlserver
/server -Dweblogic.utils.cmm.lowertier.ServiceDisabled=true weblogic.Server
[1] sbutton@dhcp-au-adelaide-10-187-112-202] ~ $
```

- Build oracle/weblogic:12.1.3-dev image
- Start 3 WebLogic Server Containers in seconds
- View status
\$ docker run -p 17001:7001 oracle/weblogic:12.1.3-dev

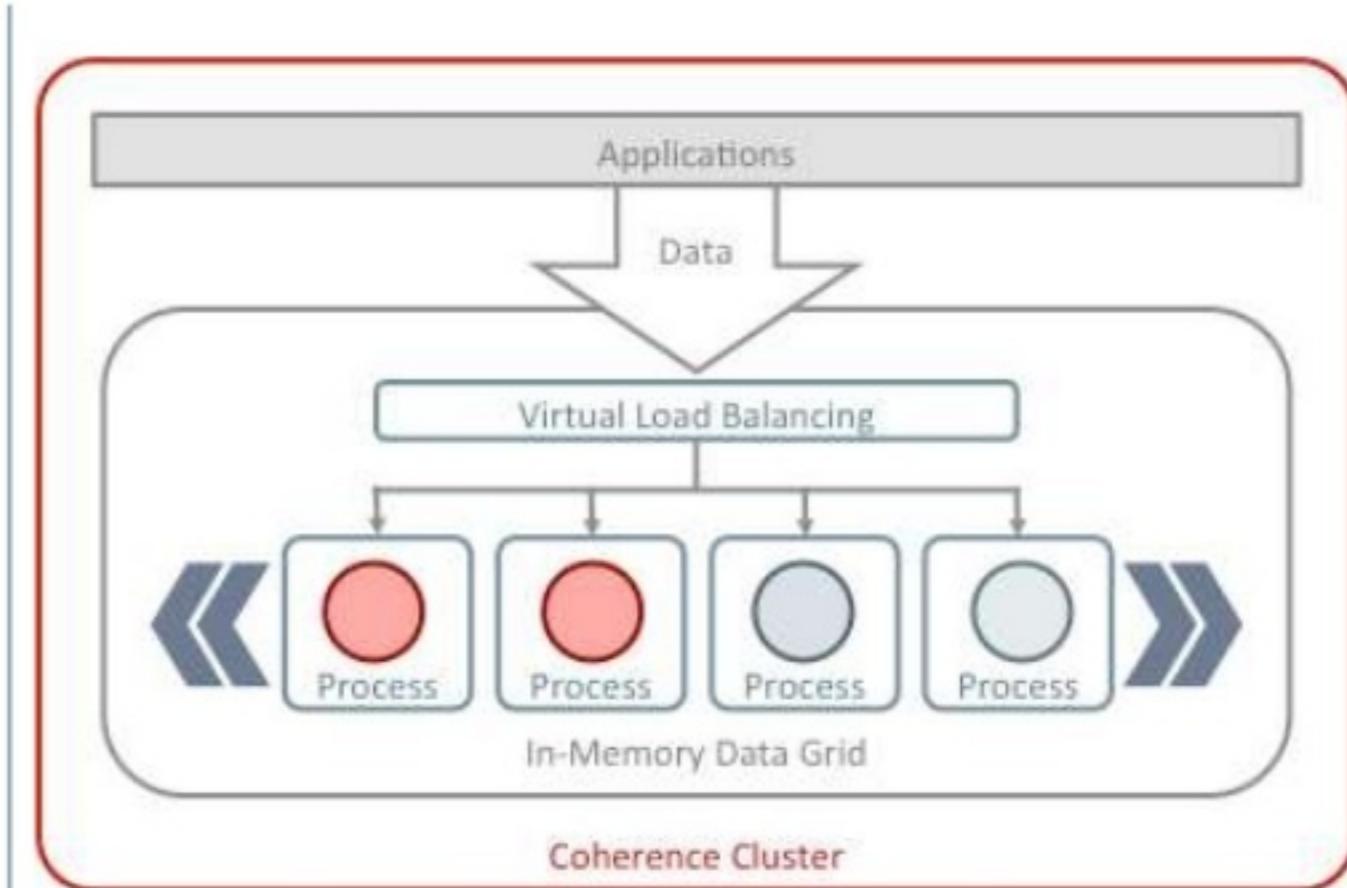
Other Supporting Products from Oracle



Oracle Coherence – Oracle's Distributed Data Grid

Architected for microservice state management

- Reliable in-memory key-value store
 - Dynamically scalable
 - Scale processing with data
 - Flexible topology support
- Java, .NET, C++, REST, Memcached, Jcache client support
- In-place distributed processing
- Queries & continuous queries
- Map-reduce aggregation
- Event notification / event programming model
- Distributed Lambdas and Streams



Oracle Messaging Cloud Is Perfect for Microservices

Implements AMQP standard

|  |  |  |  |
|---|--|---|---|
| Standardized Interfaces | Versatile | Delivery Choices | Reliability Mechanisms |
| <i>REST</i> | <i>Oracle Cloud</i> | <i>Pull</i> | <i>Transactions</i> |
| <i>JMS</i> | <i>On Premise</i> | <i>Push</i> | <i>Acknowledgements</i> |
| <i>Message push over HTTP</i> | <i>Hybrid</i> | <i>Filter</i> | <i>Durable subscriptions</i> |

Oracle Management Cloud Services – Initial Offerings



Application Performance Monitoring

Improve End-User Experience and System Performance; Diagnose Performance Issues Faster



Log Analytics

Extract Value from Logs by Collecting, Correlating, and Searching Any Kind of Log Data; Quickly Discover Anomalies



IT Analytics

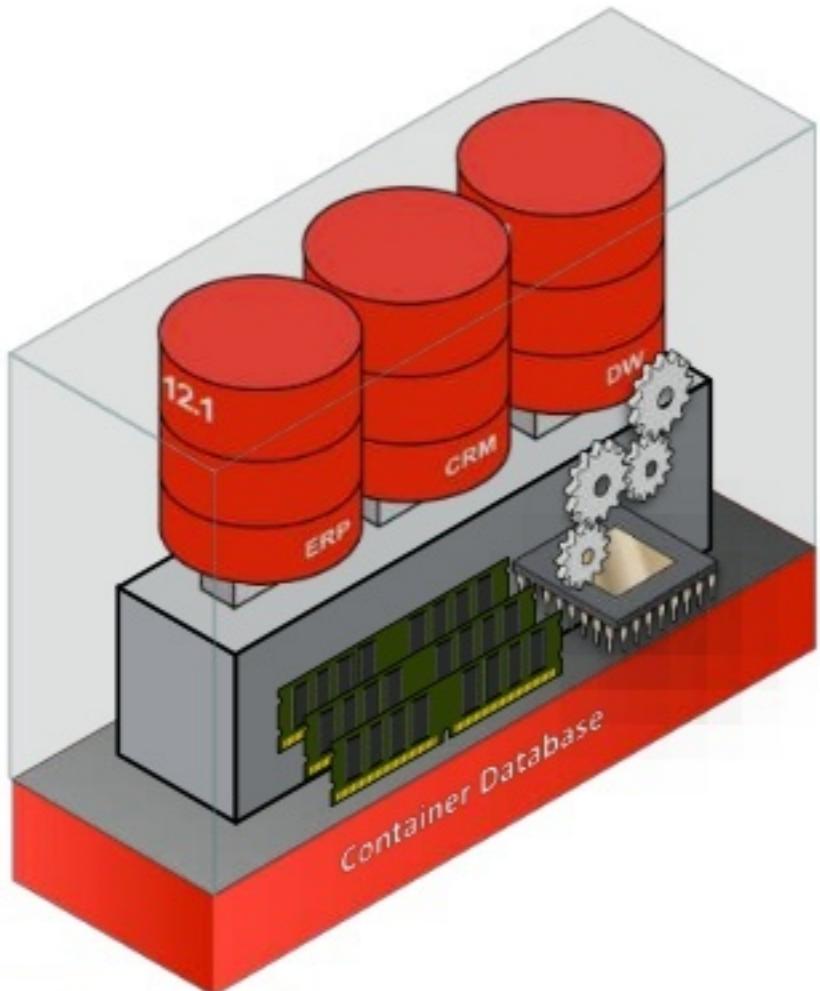
Make Critical Decisions About Your IT Estate; Plan For Growth, Run What-If Analyses, Compare Resource Usage

Oracle Datastore Options – Offered On Premise and In Cloud

| CATEGORY | PRODUCTS | DESCRIPTION |
|------------------|------------------------------|---|
| RDBMS | ORACLE® DATABASE | Supports JSON, XML, CLOBs, BLOBs, and multi-media. Accessible over client-specific APIs, REST |
| Key/Value Stores | ORACLE® NoSQL | Distributed key/value pairs, schema-less, nearly ACID compliant, scale out. Berkeley DB behind the scenes |
| Object Data Grid | ORACLE® COHERENCE | Distributed data grid that supports grid-side processing |

Pluggable Databases

One Container Database per application, one Pluggable Database per microservice



- Container Database
 - Multi-tenant database that includes zero, one or many pluggable databases
 - Upgrades, etc are performed against container
- Pluggable Database
 - A full database to the client except that behind the scenes it doesn't have its own controlfiles, redo logs, undo, etc
 - Just a collection of datafiles and tempfiles to handle its own objects, including its own data dictionary
 - Can easily move Pluggable Databases from one container to another

Oracle Database: Simple Oracle Document Access (SODA)

NoSQL backed by Oracle Database

- Enable schemaless development on top of an Oracle Database
 - Provide a simple NoSQL-style API for working with documents
- Make it easy to use Oracle as a NoSQL-style document store
- Supports all common application development environments
- Supports SQL!
- Works with Oracle Database tooling
 - backup, restore, security, etc



SODA For Java

Native library for Java

- Provides:
 - Connection to the database (replaces JDBC!)
 - Collection Management
 - CRUD operations
 - Query-by-Example
 - Utility and control functions
- SODA for Java applications can use a JDBC connection to talk to the database
- SODA for Java is transactional
 - Supports Hybrid model with JDBC and SODA based operations

```
// Create a Connection
OracleRDBMSClient client = new OracleRDBMSClient();
OracleDatabase database = client.getDatabase(conn);

// Now create a collection
OracleCollection collection =
database.getDatabaseAdmin().createCollection("MyCollection");

// Create a document
OracleDocument document = database.createDocumentFromString("{\"name\" : \"Alexander\" }");

// Next, insert it into the collection
OracleDocument insertedDocument = collection.insertAndGet(document);

// Get the key of the inserted document
String key = insertedDocument.getKey();

// Get the version of the inserted document
String version = insertedDocument.getVersion();
```



SODA For REST

Can be invoked from any programming language

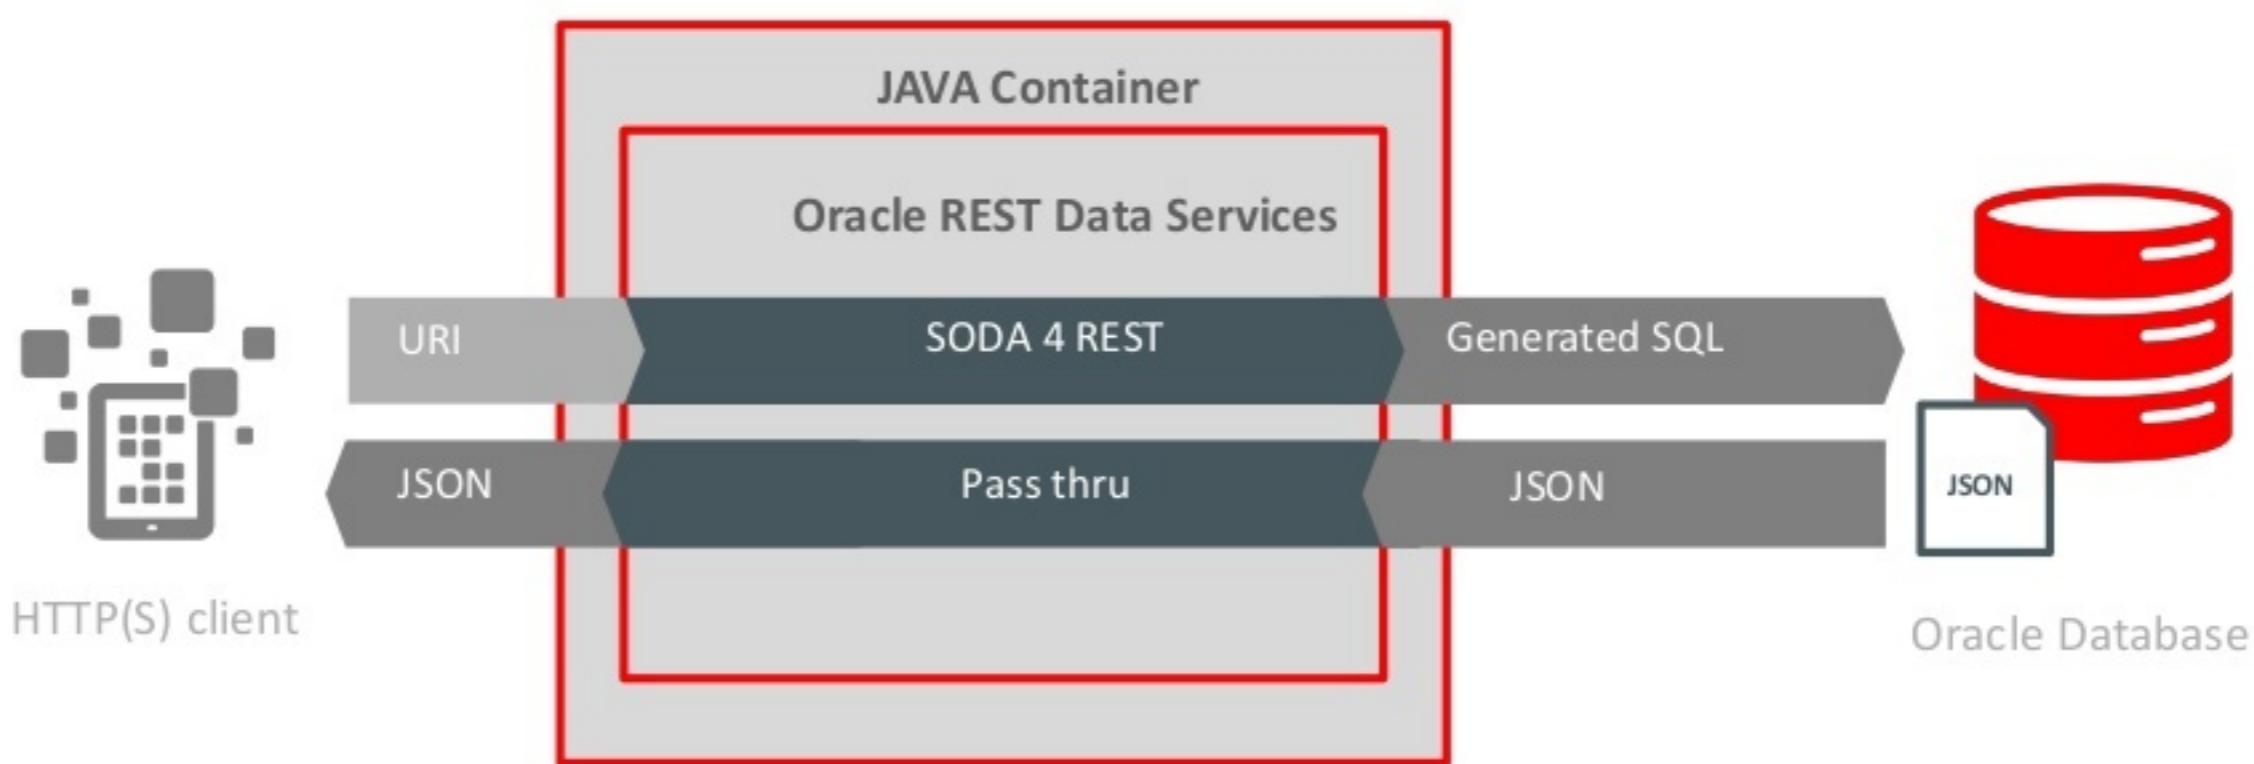
- Standard REST based model
- CRUD operations map to HTTP Verbs
 - Create / Update : PUT / POST
 - Retrieve : GET
 - Delete : DELETE
- Other operations, such as Query by Example, Bulk Insert and Indexing are mapped to variants of POST
- JSON document forms the payload of the HTTP Request or Response
- Stateless model, no transaction support
- Implemented as a Java Servlet

| | |
|----------------------------------|-----------------------------------|
| GET /DBSODA/schema | List all collections in a schema |
| GET /DBSODA/schema/collection | Get all objects in collection |
| GET /DBSODA/schema/collection/id | Get specific object in collection |
| PUT /DBSODA/schema/collection | Create a collection if necessary |
| PUT /DBSODA/schema/collection/id | Update object with id |
| POST /DBSODA/schema/collection | Insert object into collection |
| POST | Find objects matching filter in |



SODA For REST Architecture

Clean architecture



- Data stored in Oracle Database as JSON documents
- App Developer make standard HTTP(S) calls to SODA for REST API

Summary

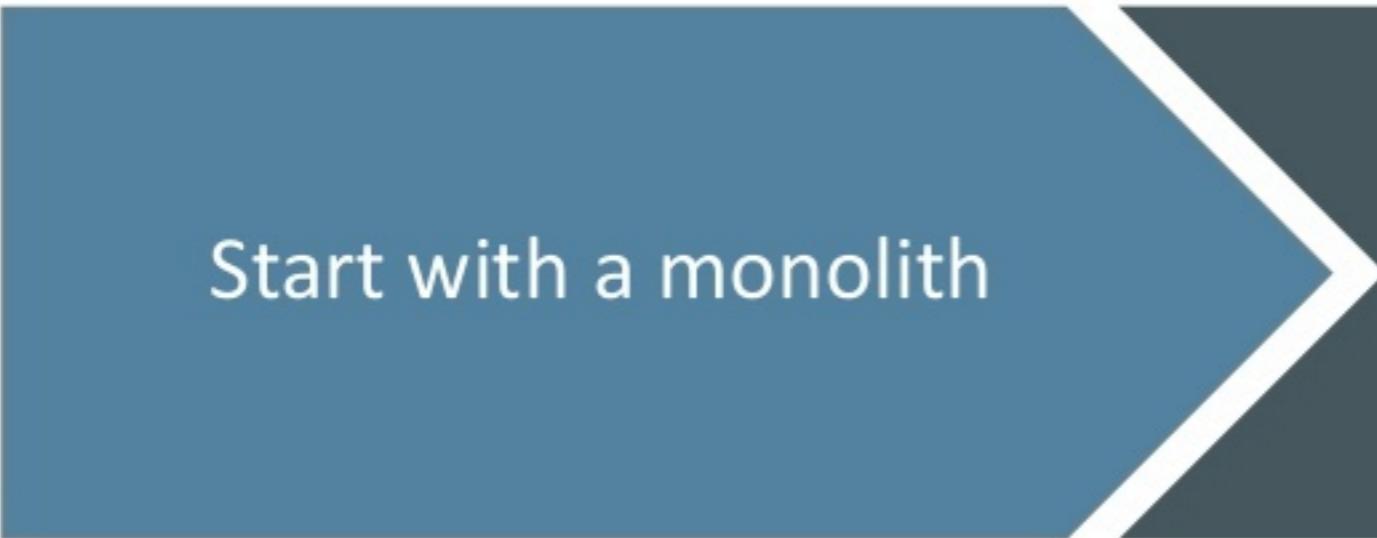


Problems to Solve With Microservices

Many, very sticky problems to solve



Microservices Adoption Patterns



Start with a monolith



Adopt Microservices
(Only if the monolith gets too big and
you've met the other prerequisites)



Microservice Takeaways for Enterprises

The term “microservices” is probably a passing fad, but there is real business value

Constraints:

Time | Competency | Resources

Ability to Execute

Technical Prerequisites

Architectural Prerequisites

Non-technical Prerequisites



Remember: Microservices Are Not a Silver Bullet

It's probably best to stick with monoliths

"There is no single development, in either technology or management technique, which by itself promises even one order of magnitude [tenfold] improvement within a decade in productivity, in reliability, in simplicity."



"We cannot expect ever to see two-fold gains every two years" in software development, like there is in hardware development (Moore's law)

Integrated Cloud Applications & Platform Services





ORACLE®