


Building a Distributed Data Ingestion System with RabbitMQ

Alvaro Videla - RabbitMQ



Alvaro Videla


- Developer Advocate at Pivotal / RabbitMQ
- Co-Author of RabbitMQ in Action
- Creator of the RabbitMQ Simulator
- Blogs about RabbitMQ Internals: <http://videlalvaro.github.io/internals.html>
- @old_sound — alvaro@rabbitmq.com — github.com/videlalvaro

About Me

Co-authored

RabbitMQ in Action

<http://bit.ly/rabbitmq>



About this Talk

- Exploratory Talk
- A 'what could be done' talk instead of 'this is how you do it'

Agenda

- Intro to RabbitMQ
- The Problem
- Solution Proposal
- Improvements

What is RabbitMQ

RabbitMQ

RabbitMQ

RabbitMQ

- Multi Protocol Messaging Server

RabbitMQ

- Multi Protocol Messaging Server
- Open Source (MPL)

RabbitMQ

- Multi Protocol Messaging Server
- Open Source (MPL)
- Polyglot

RabbitMQ

- Multi Protocol Messaging Server
- Open Source (MPL)
- Polyglot
- Written in Erlang/OTP

Multi Protocol



<http://bit.ly/rmq-protocols>

Polyglot

Polyglot

Polyglot

- Java

Polyglot

- Java
- node.js

Polyglot

- Java
- node.js
- Erlang

Polyglot

- Java
- node.js
- Erlang
- PHP

Polyglot

- Java
- node.js
- Erlang
- PHP
- Ruby


Polyglot

- Java
- node.js
- Erlang
- PHP
- Ruby
- .Net

Polyglot

- Java
- node.js
- Erlang
- PHP
- Ruby
- .Net
- Haskell

Polyglot



Even COBOL!!!11

Some users of RabbitMQ

Some users of RabbitMQ

- Instagram

Some users of RabbitMQ

- Instagram
- Indeed.com

Some users of RabbitMQ

- Instagram
- Indeed.com
- Telefonica

Some users of RabbitMQ

- Instagram
- Indeed.com
- Telefonica
- Mercado Libre

Some users of RabbitMQ

- Instagram
- Indeed.com
- Telefonica
- Mercado Libre
- NHS

Some users of RabbitMQ

- Instagram
- Indeed.com
- Telefonica
- Mercado Libre
- NHS
- Mozilla

<http://www.rabbitmq.com/download.html>

Unix - Mac - Windows


Messaging with RabbitMQ

A demo with the RabbitMQ Simulator

<https://github.com/RabbitMQSimulator/RabbitMQSimulator>

<http://tryrabbitmq.com>

RabbitMQ Simulator



The Problem

Distributed Application



Distributed Application



Ad-hoc solution

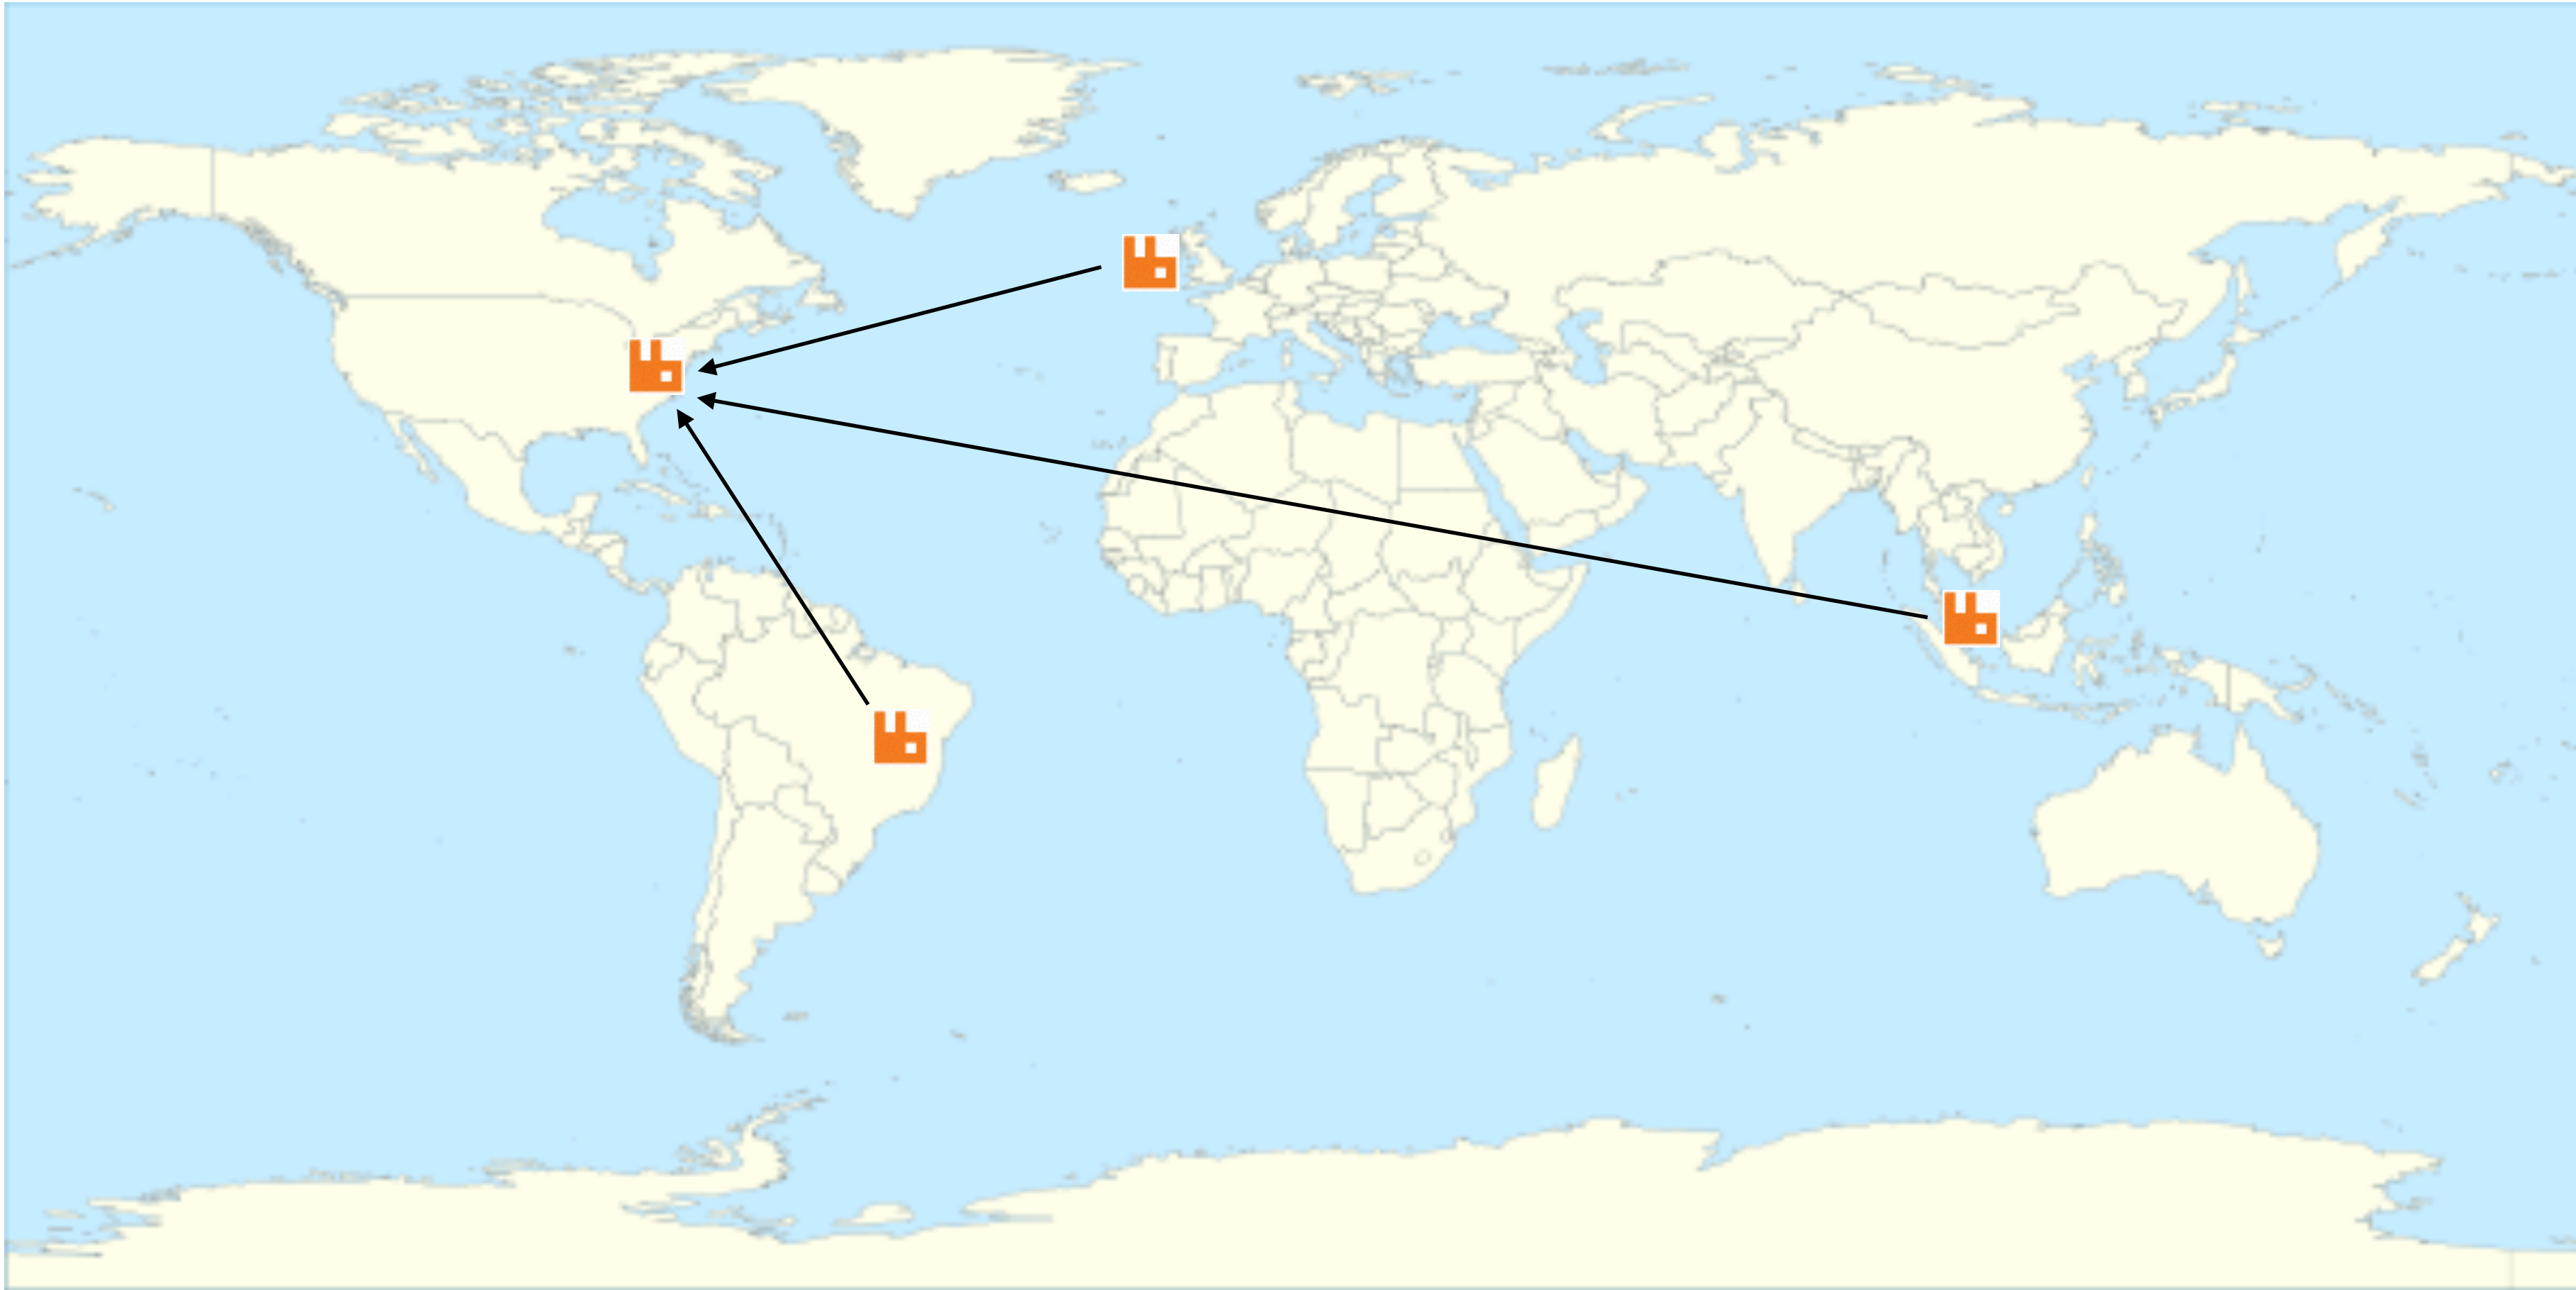
A process that replicates data
to the remote server

Possible issues

- Remote server is offline
 - Prevent unbounded local buffers
 - Prevent message loss
- Prevent unnecessary message replication
 - No need for those messages on remote server
 - Messages that became stale

Can we do better?


RabbitMQ Federation




RabbitMQ Federation

- Supports replication across different administrative domains
- Supports mix of Erlang and RabbitMQ versions
- Supports Network Partitions
- Specificity - not everything has to be federated

RabbitMQ Federation



RabbitMQ Federation



RabbitMQ Federation

RabbitMQ Federation

- It's a RabbitMQ **Plugin**

RabbitMQ Federation

- It's a RabbitMQ **Plugin**
- Internally uses **Queues** and **Exchanges Decorators**

RabbitMQ Federation

- It's a RabbitMQ **Plugin**
- Internally uses **Queues** and **Exchanges Decorators**
- Managed using **Parameters** and **Policies**

Enabling the Plugin

```
rabbitmq-plugins enable rabbitmq_federation
```

Enabling the Plugin

```
rabbitmq-plugins enable rabbitmq_federation
```

```
rabbitmq-plugins enable rabbitmq_federation_management
```

Federating an Exchange

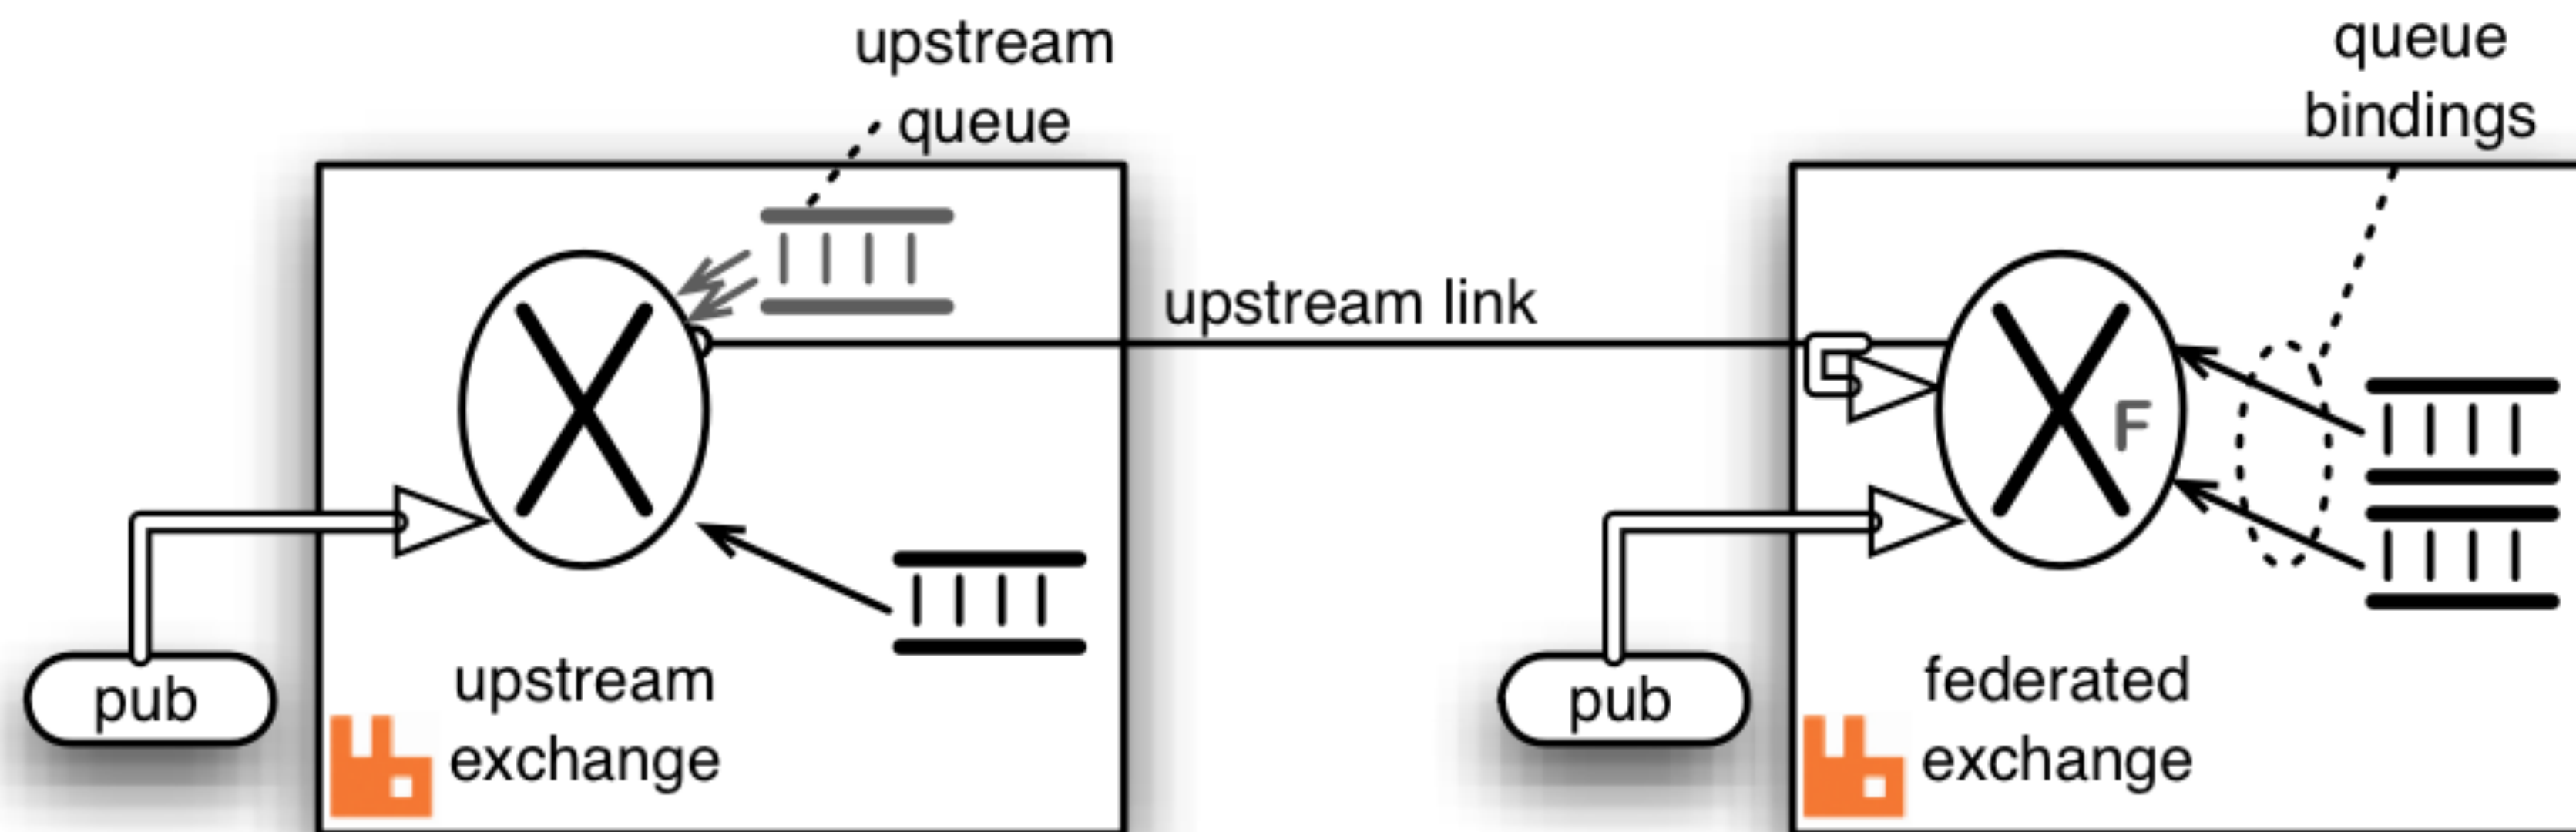
```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```

Federating an Exchange

```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```

```
rabbitmqctl set_policy --apply-to exchanges federate-me "^amqp\." \  
'{"federation-upstream-set": "all"}'
```

Federating an Exchange



Configuring Federation

Config Options

```
rabbitmqctl set_parameter federation-upstream \  
name 'json-object'
```


Config Options

```
rabbitmqctl set_parameter federation-upstream \  
name 'json-object'
```

```
json-object: {  
  'uri': 'amqp://server-name/',  
  'prefetch-count': 1000,  
  'reconnect-delay': 1,  
  'ack-mode': on-confirm  
}
```

<http://www.rabbitmq.com/federation-reference.html>

Prevent unbound buffers

`expires: N // ms.`

`message-ttl: N // ms.`

Prevent message forwarding

`max-hops: N`

Speed vs No Message Loss

`ack-mode: on-confirm`

`ack-mode: on-publish`

`ack-mode: no-ack`

AMQP URI:


`amqp://user:pass@host:10000/vhost`

<http://www.rabbitmq.com/uri-spec.html>

Config can be applied via

- CLI using **rabbitmqctl**
- HTTP API
- RabbitMQ Management Interface

RabbitMQ Federation



Scaling the Setup

The Problem

The Problem

- Queues contents live in the node where the Queue was declared

The Problem

- Queues contents live in the node where the Queue was declared
- A cluster can access the queue from every connected node

The Problem

- Queues contents live in the node where the Queue was declared
- A cluster can access the queue from every connected node
- Queues are an Erlang process (tied to one core)

The Problem

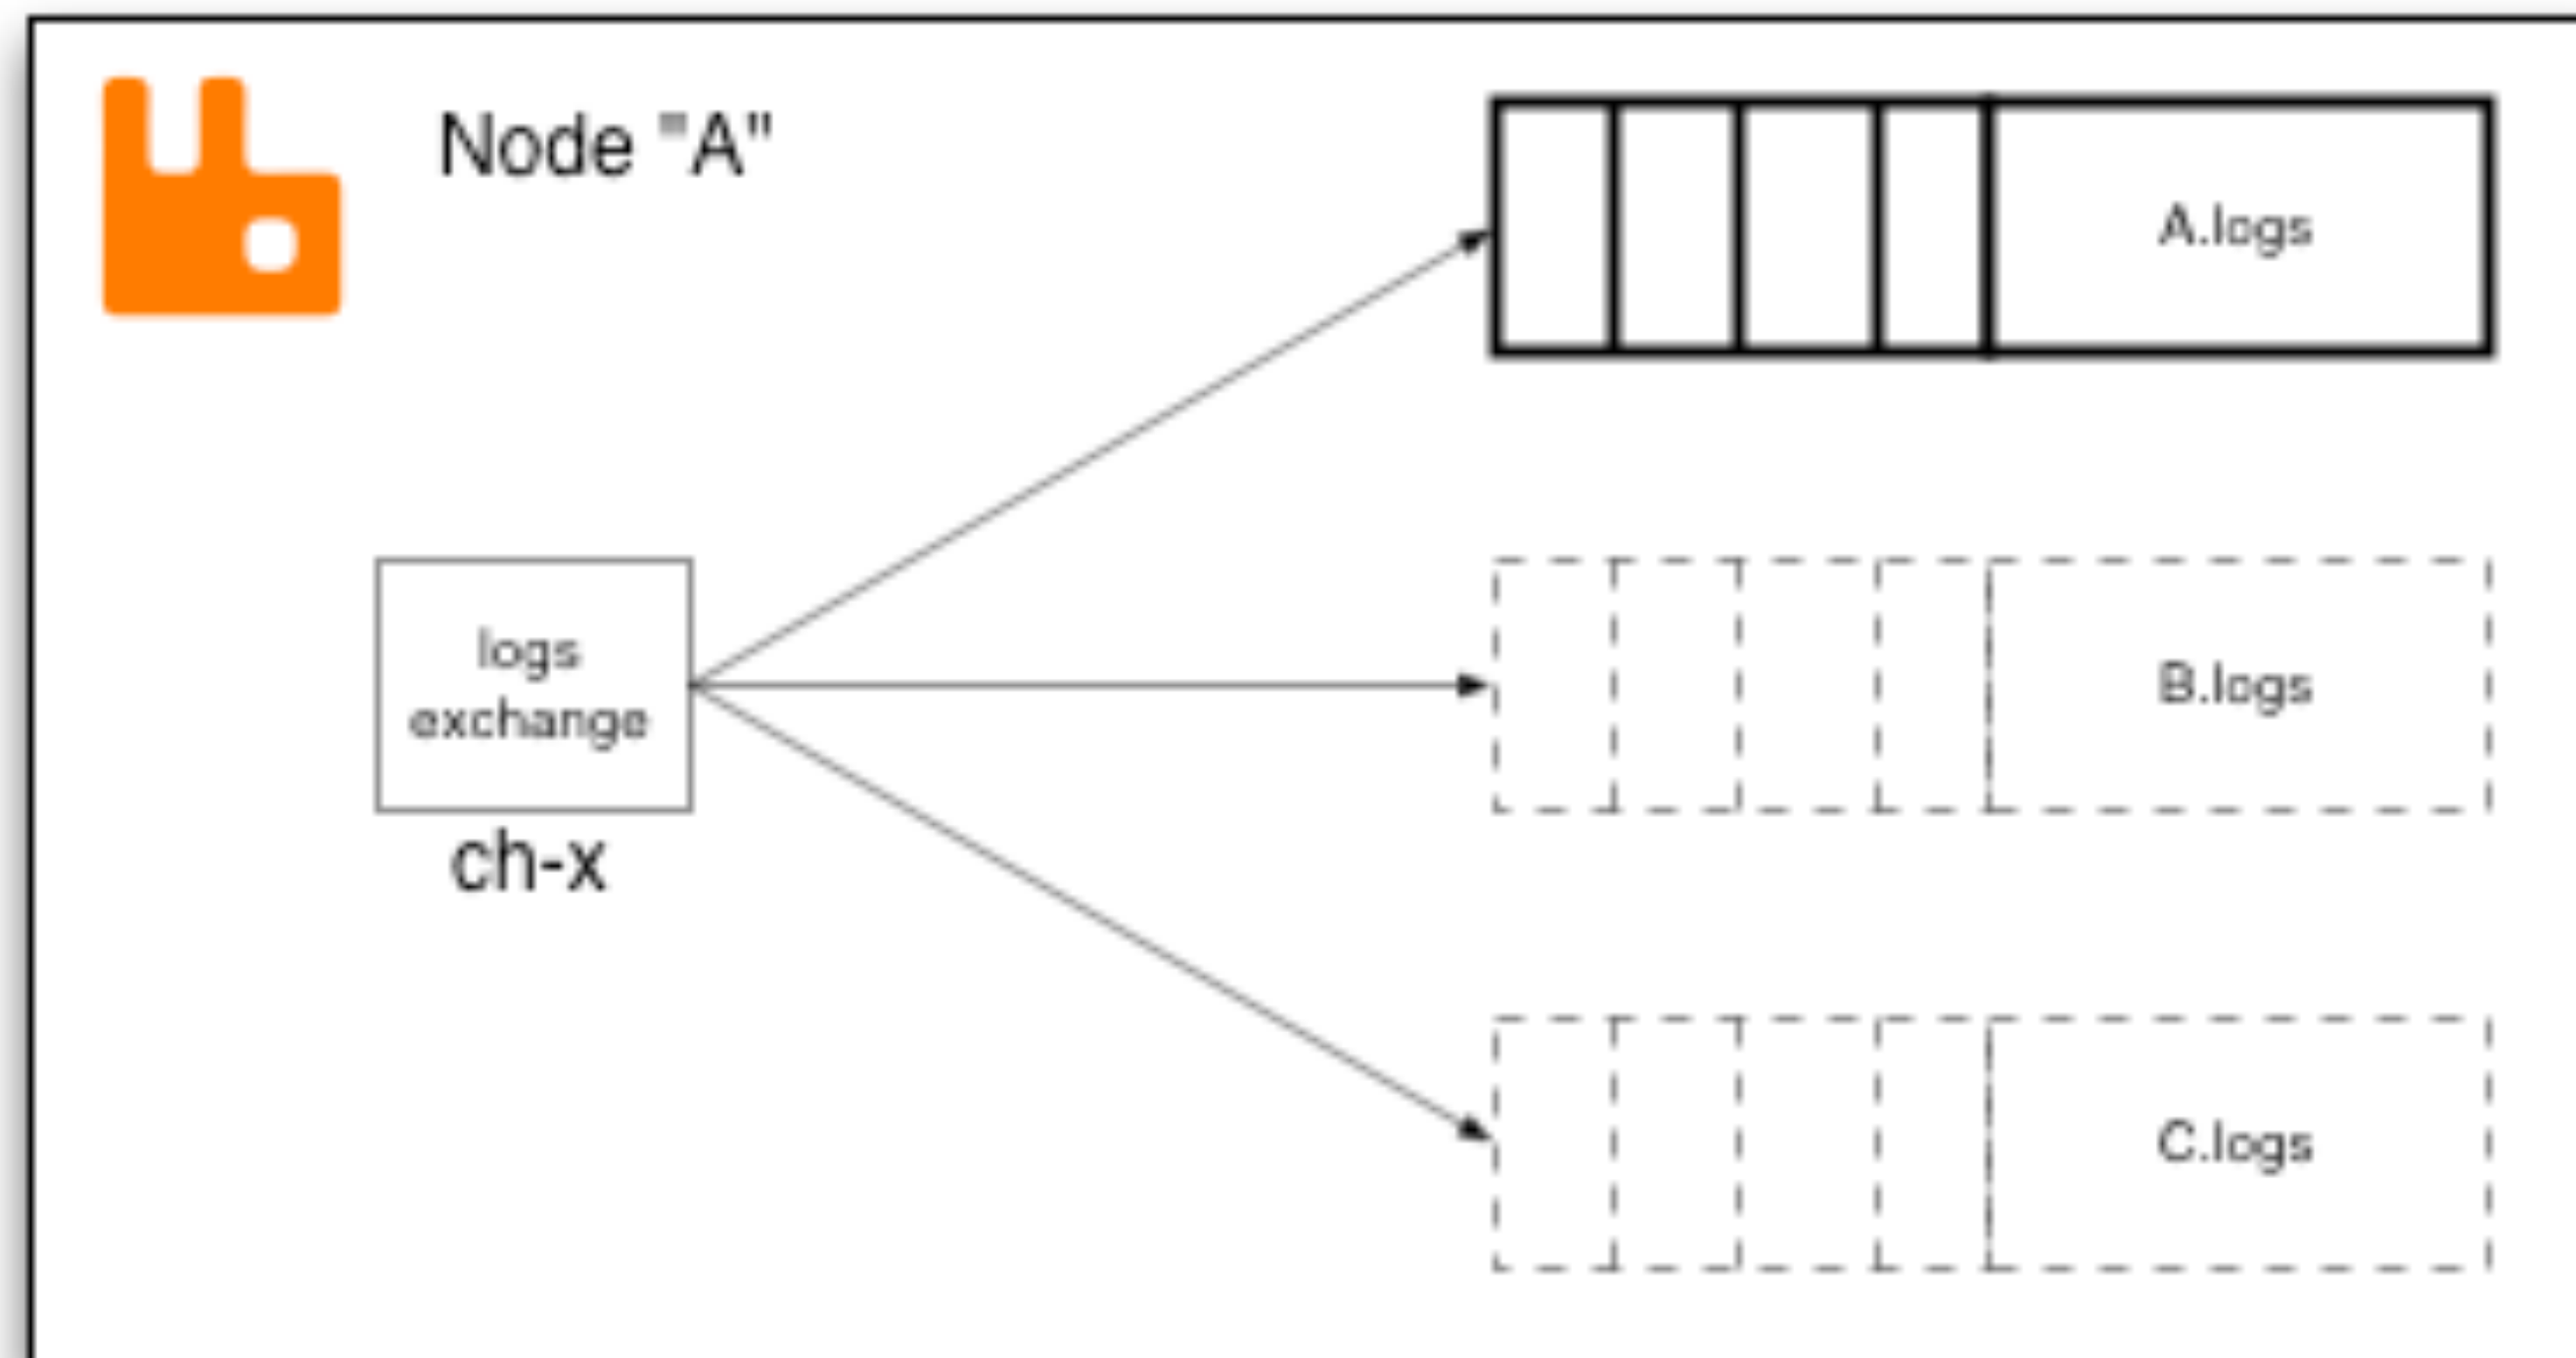
- Queues contents live in the node where the Queue was declared
- A cluster can access the queue from every connected node
- Queues are an Erlang process (tied to one core)
- Adding more nodes doesn't really help

Sharded Queues


Pieces of the Puzzle

- consistent hash exchange
- good ol' queues


Sharded Queues



Sharded Queues



Sharded Queues



Sharded Queues

- Declare Queues with name: `nodename.queueName.index`

Sharded Queues

- Declare Queues with name: `nodename.queueindex`
- Bind the queues to a consistent hash exchange

Sharded Queues

- Declare Queues with name: `nodename.queueindex`
- Bind the queues to a consistent hash exchange
- Get the consumer to randomly build the queue name

We need more scale!

Federated Queues

Federated Queues

- Load-balance messages across federated queues
- Only moves messages when needed

Federating a Queue

```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```

Federating a Queue

```
rabbitmqctl set_parameter federation-upstream my-upstream \  
'{"uri": "amqp://server-name", "expires": 3600000}'
```

```
rabbitmqctl set_policy --apply-to queues federate-me "^images\." \  
'{"federation-upstream-set": "all"}'
```

With RabbitMQ we can

With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP

With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP
- Distribute that data globally using Federation

With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP
- Distribute that data globally using Federation
- Scale up using Sharding

With RabbitMQ we can

- Ingest data using various protocols: AMQP, MQTT and STOMP
- Distribute that data globally using Federation
- Scale up using Sharding
- Load balance consumers with Federated Queues

Credits

world map: [wikipedia.org](https://en.wikipedia.org)

federation diagrams: rabbitmq.com

Questions?

Thanks

Alvaro Videla - @old_sound